

# Practical Black-box Optimization

Luigi Nardi

WASP-AI Assistant Professor Lund University

Researcher Stanford University

@RSS seminar, November 12, 2019



LUND  
UNIVERSITY



# Main Papers Behind This Talk

1. Nardi, et al., "Practical Design Space Exploration", MASCOTS, 2019.
2. Koeplinger, et al., "Spatial: A Language and Compiler for Application Accelerators", PLDI, 2018.
3. Nardi, et al., "Algorithmic performance-accuracy trade-off in 3D vision applications using HyperMapper", iWAPT-IPDPS, 2017.
4. Saeedi, et al., "Application-oriented design space exploration for SLAM algorithms", ICRA, 2017.
5. Bodin, et al., "Integrating algorithmic parameters into benchmarking and design space exploration in 3D scene understanding", PACT, 2016.

# Outline

## 1. Black-box Optimization

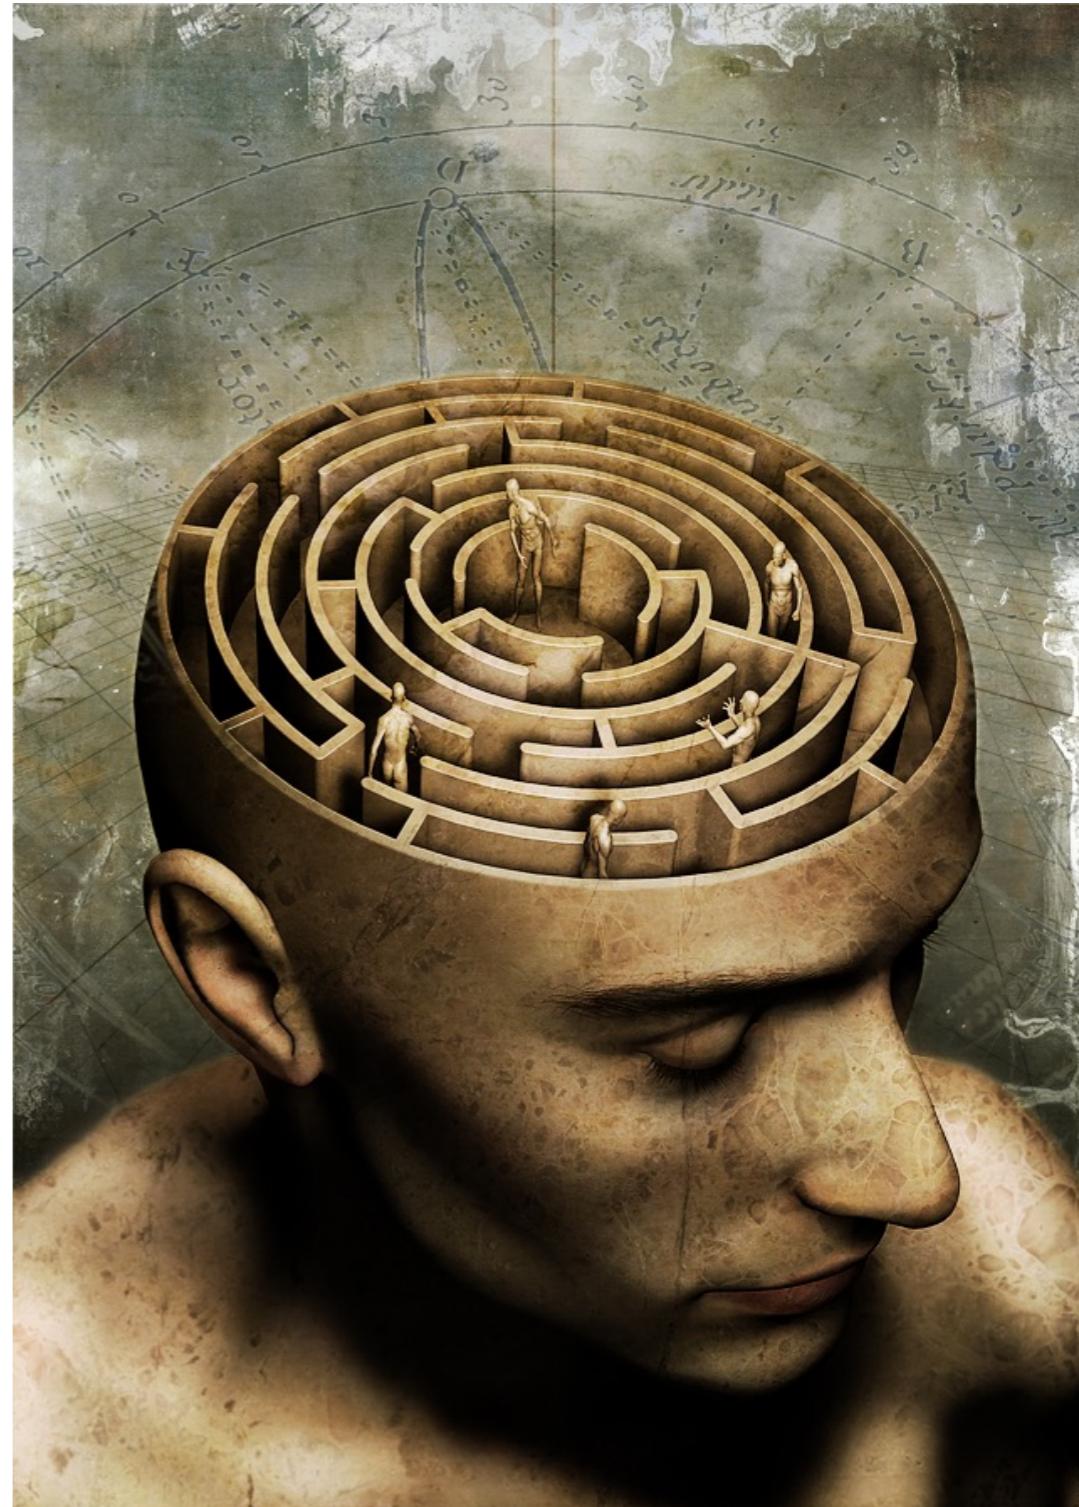
- Problem Setting
- Motivation Examples
- Taxonomy

## 2. The HyperMapper Framework

- Bayesian Optimization
- Pareto and Hypervolume Indicator
- Prior Distribution

## 3. Experimental Results

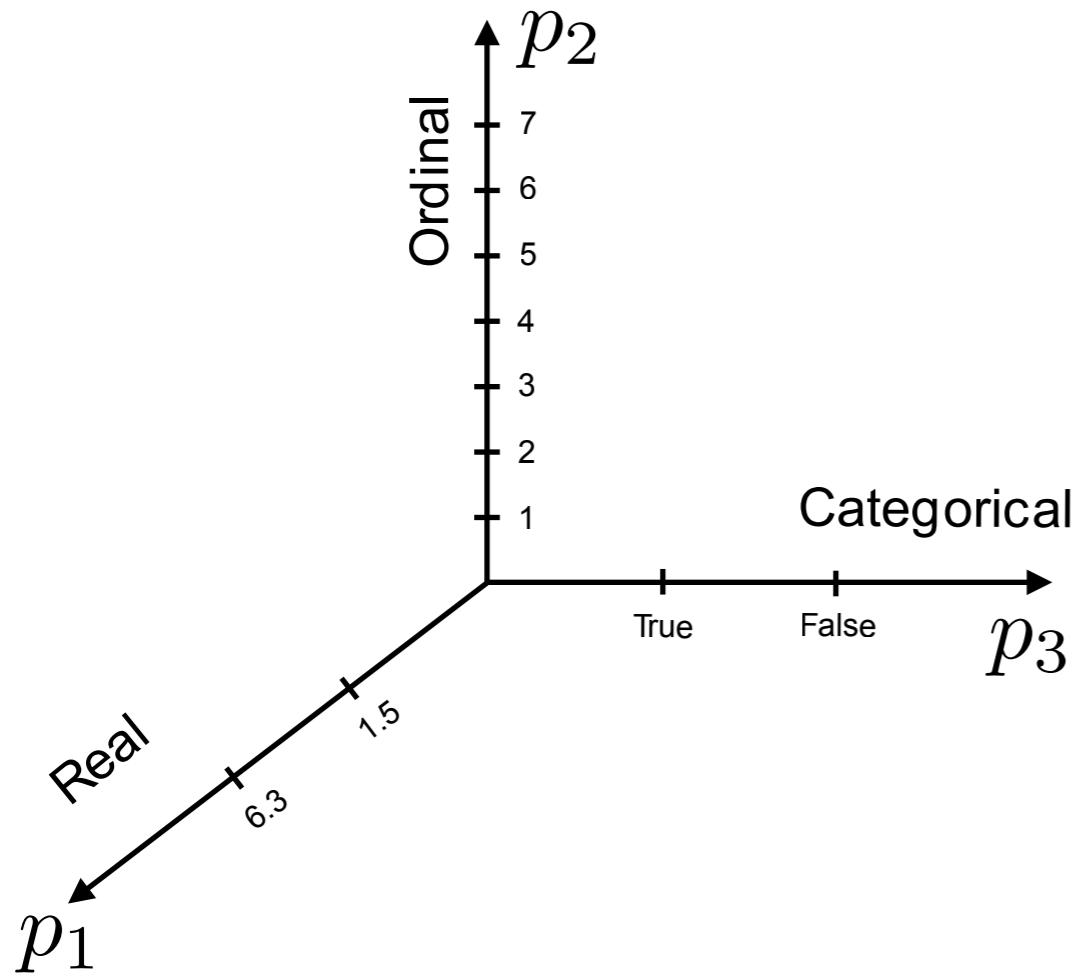
- Hardware Design
- Computer Vision
- Databases



# Black-box Optimization (BBO)

## 3-parameters and 2-objective Pictorial

Input space  
(a.k.a. search or design space)

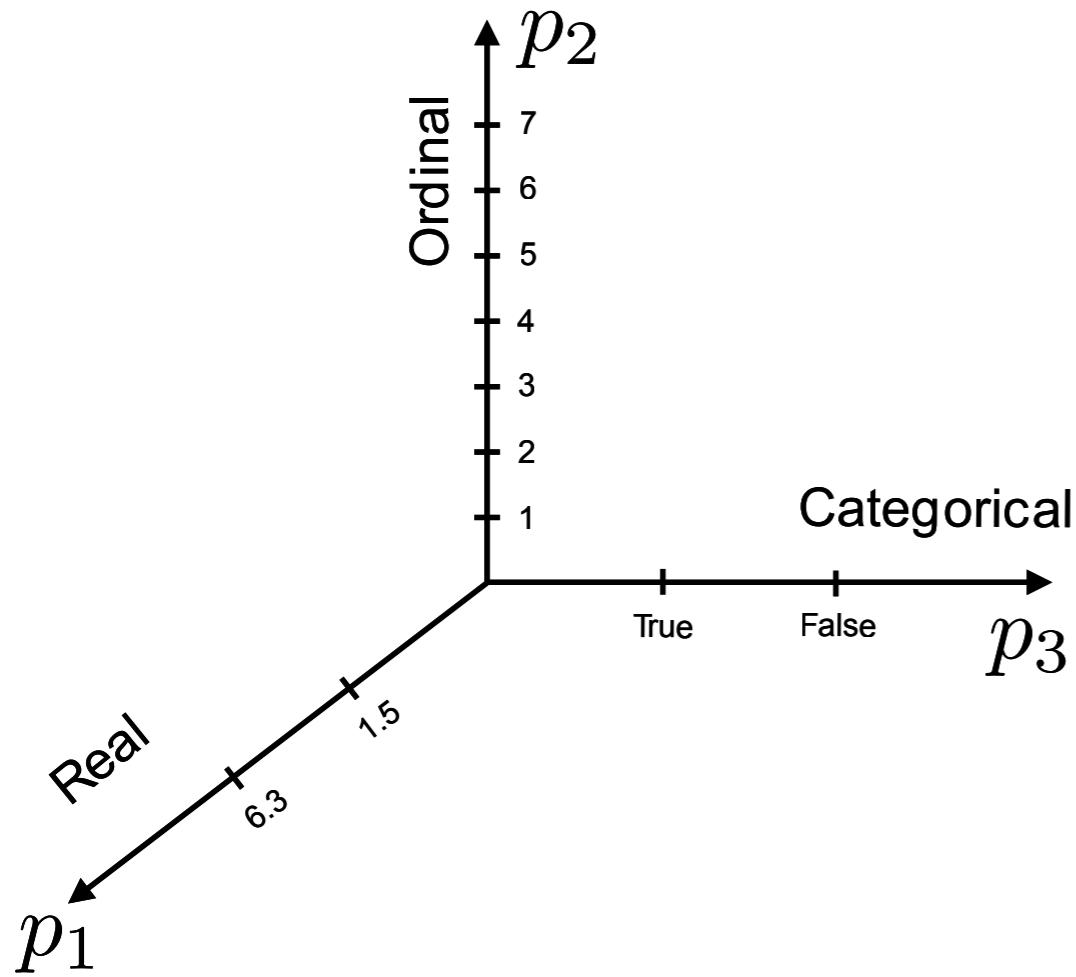


Derivatives are unavailable: e.g., categorical variables

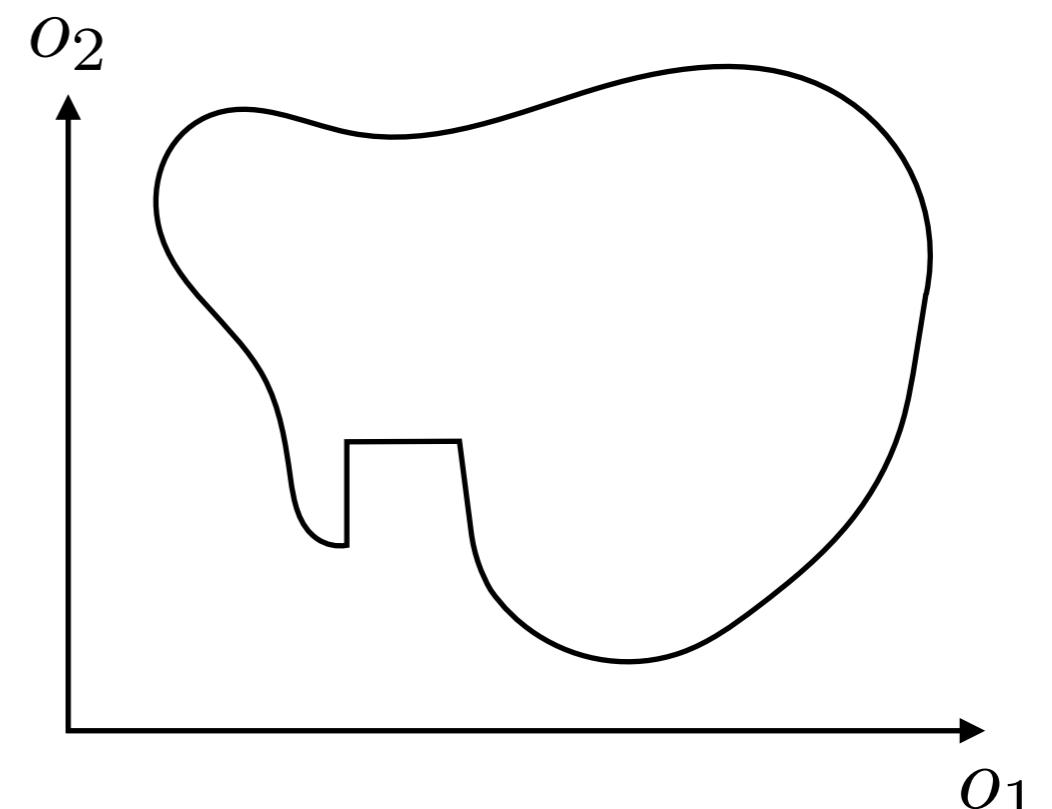
# Black-box Optimization (BBO)

## 3-parameters and 2-objective Pictorial

**Input space**  
(a.k.a. search or design space)



**Output space**  
(a.k.a. objective space)

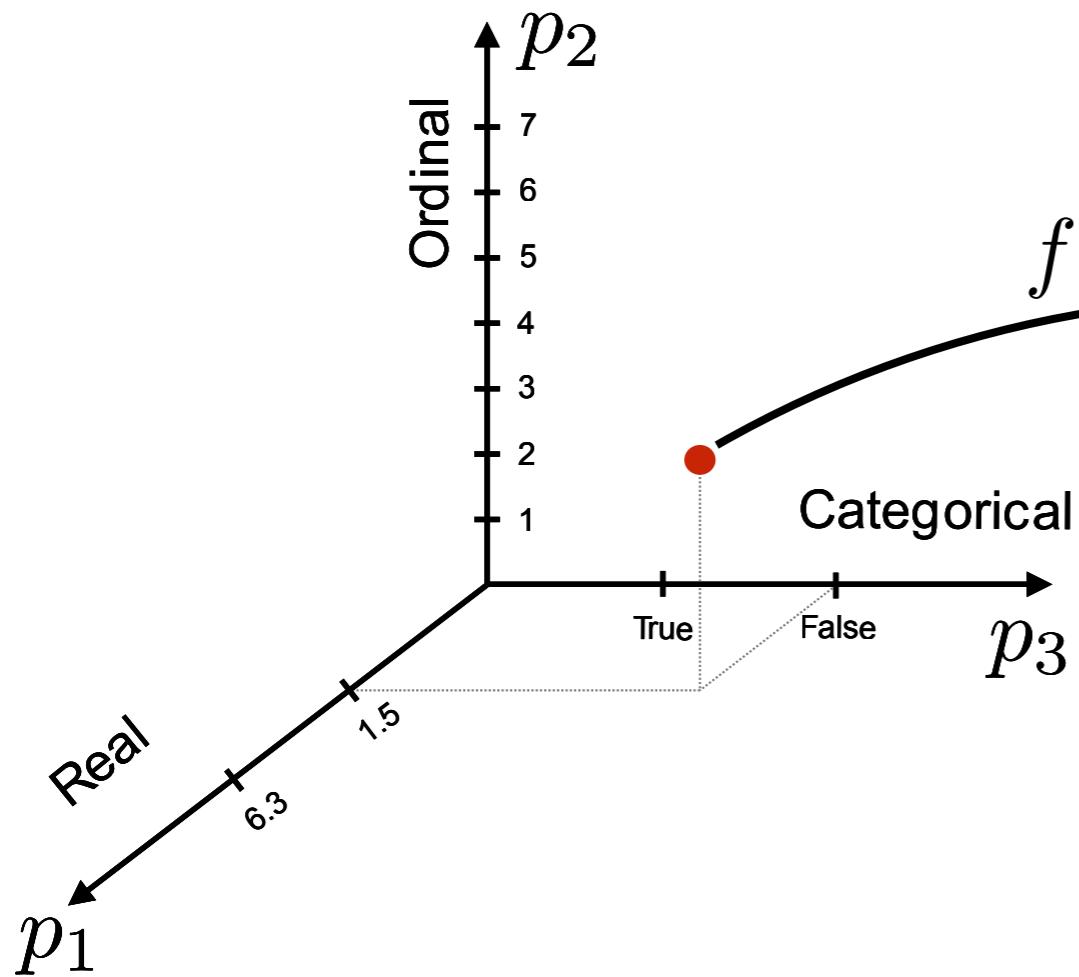


Derivatives are unavailable: e.g., categorical variables

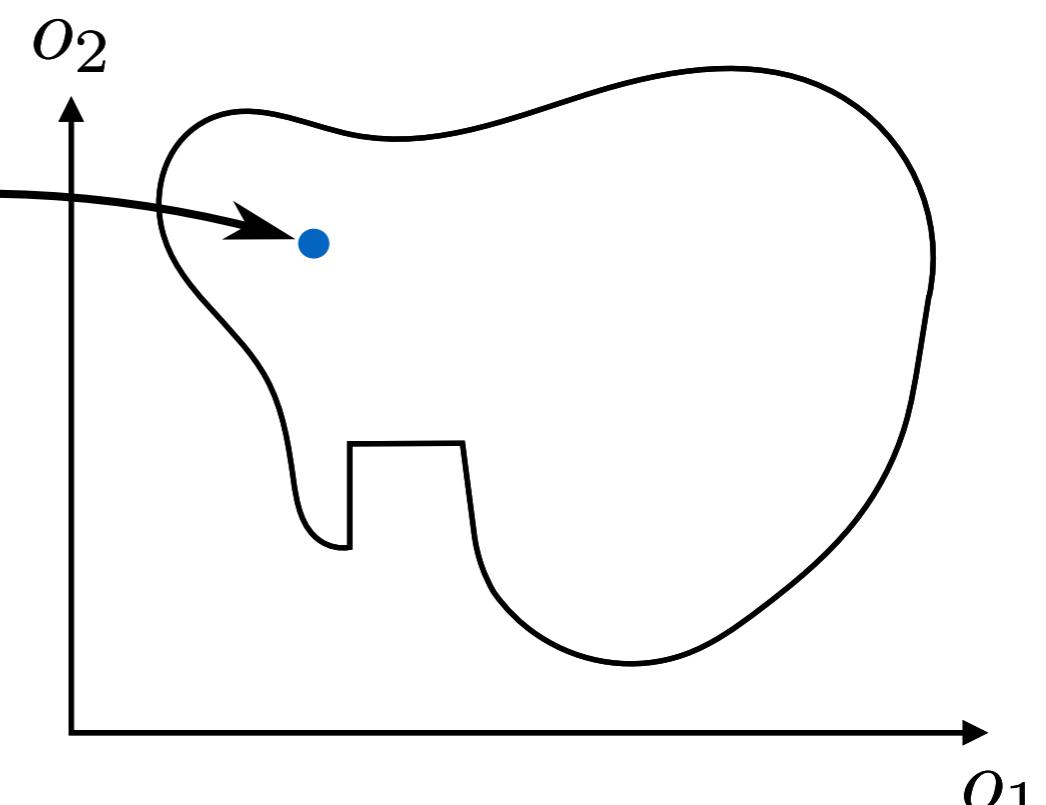
# Black-box Optimization (BBO)

## 3-parameters and 2-objective Pictorial

**Input space**  
(a.k.a. search or design space)



**Output space**  
(a.k.a. objective space)

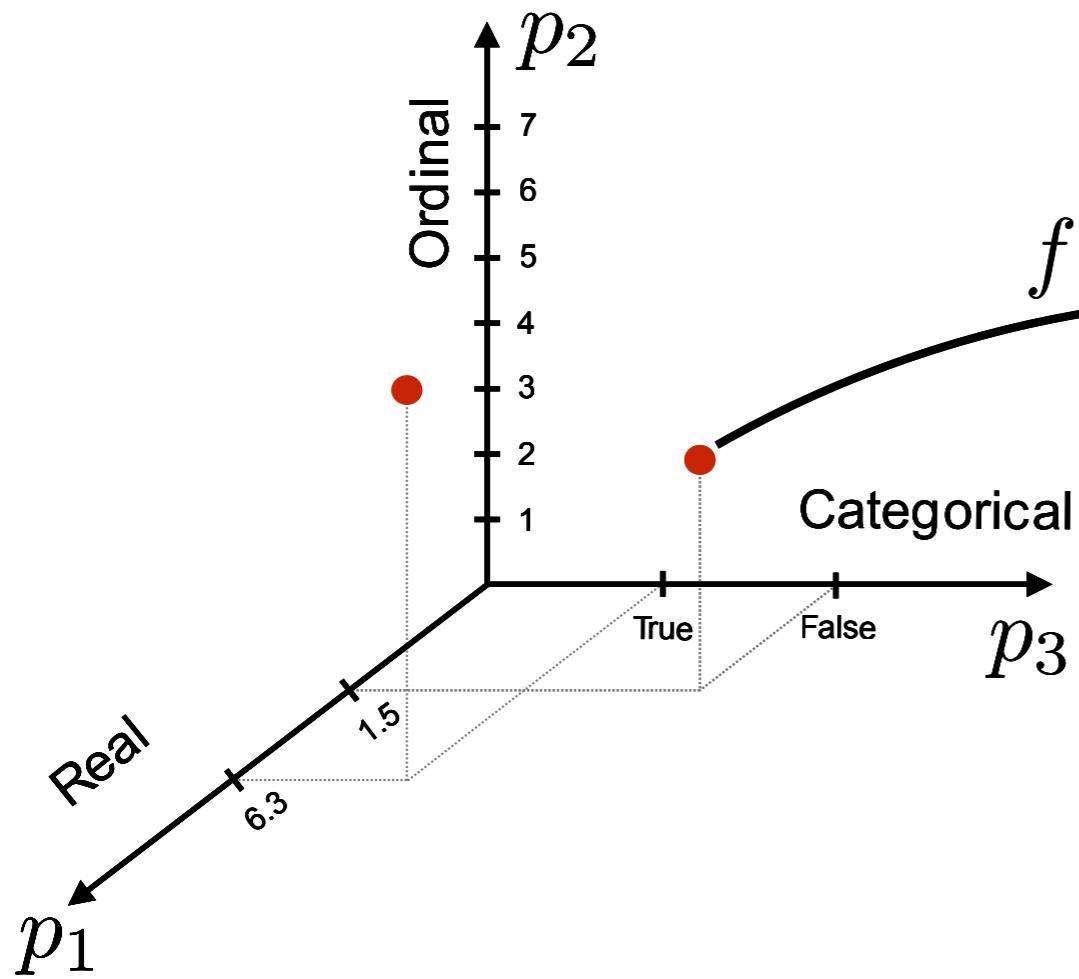


Derivatives are unavailable: e.g., categorical variables

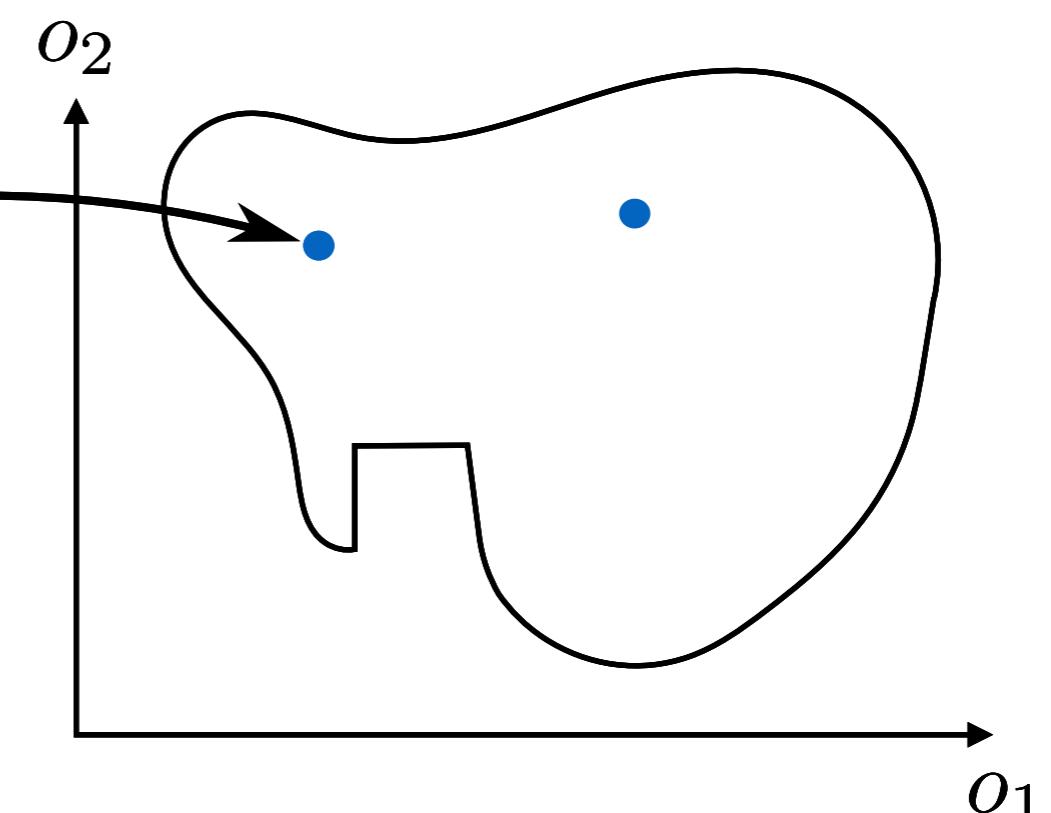
# Black-box Optimization (BBO)

## 3-parameters and 2-objective Pictorial

**Input space**  
(a.k.a. search or design space)



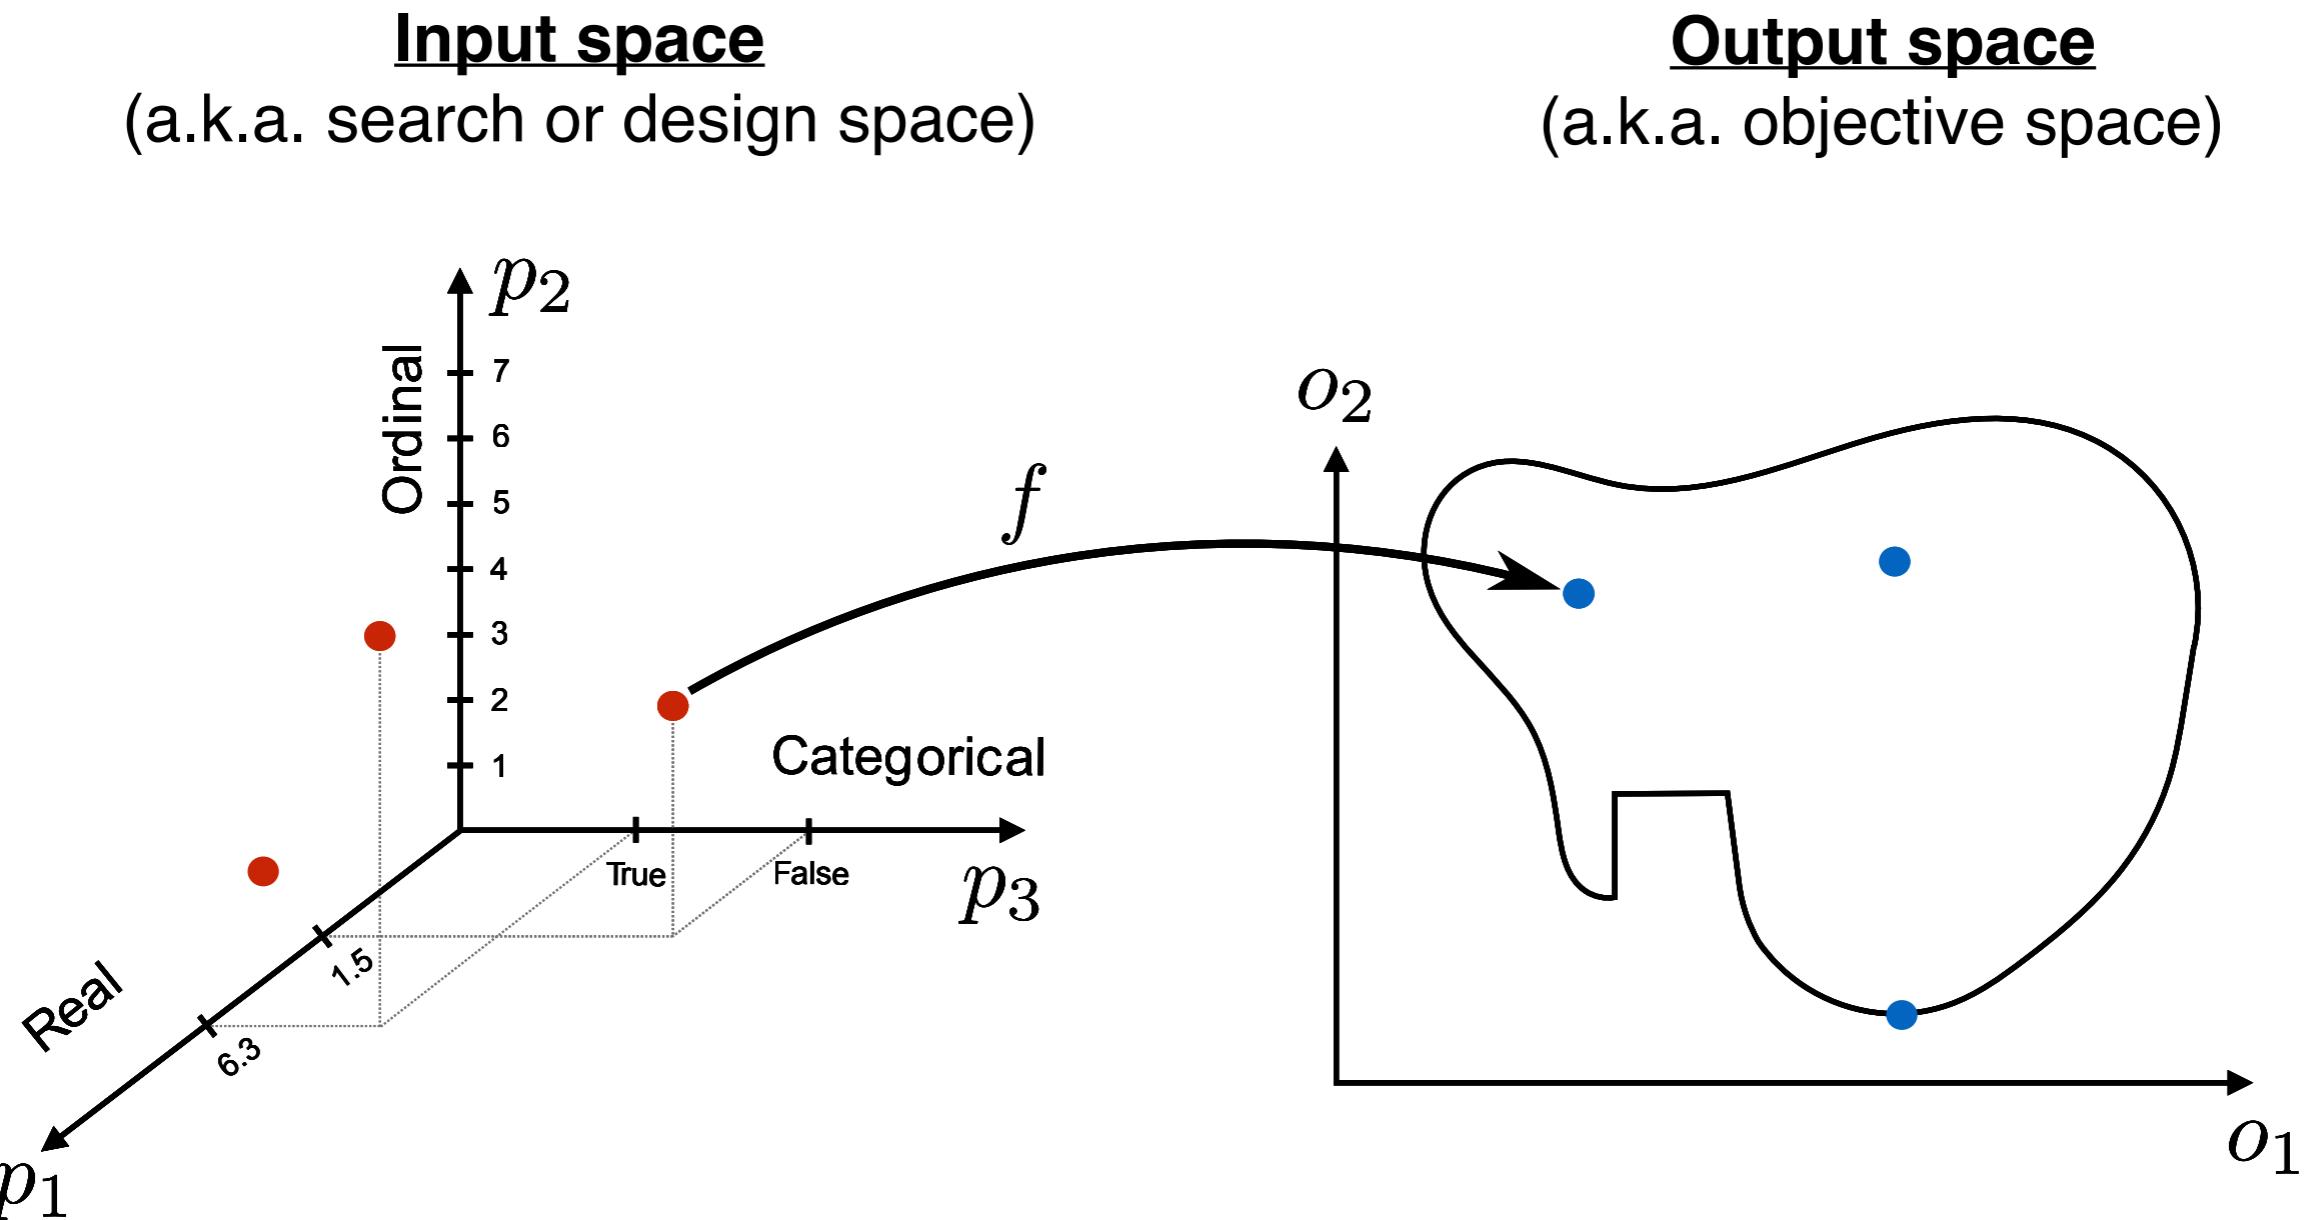
**Output space**  
(a.k.a. objective space)



Derivatives are unavailable: e.g., categorical variables

# Black-box Optimization (BBO)

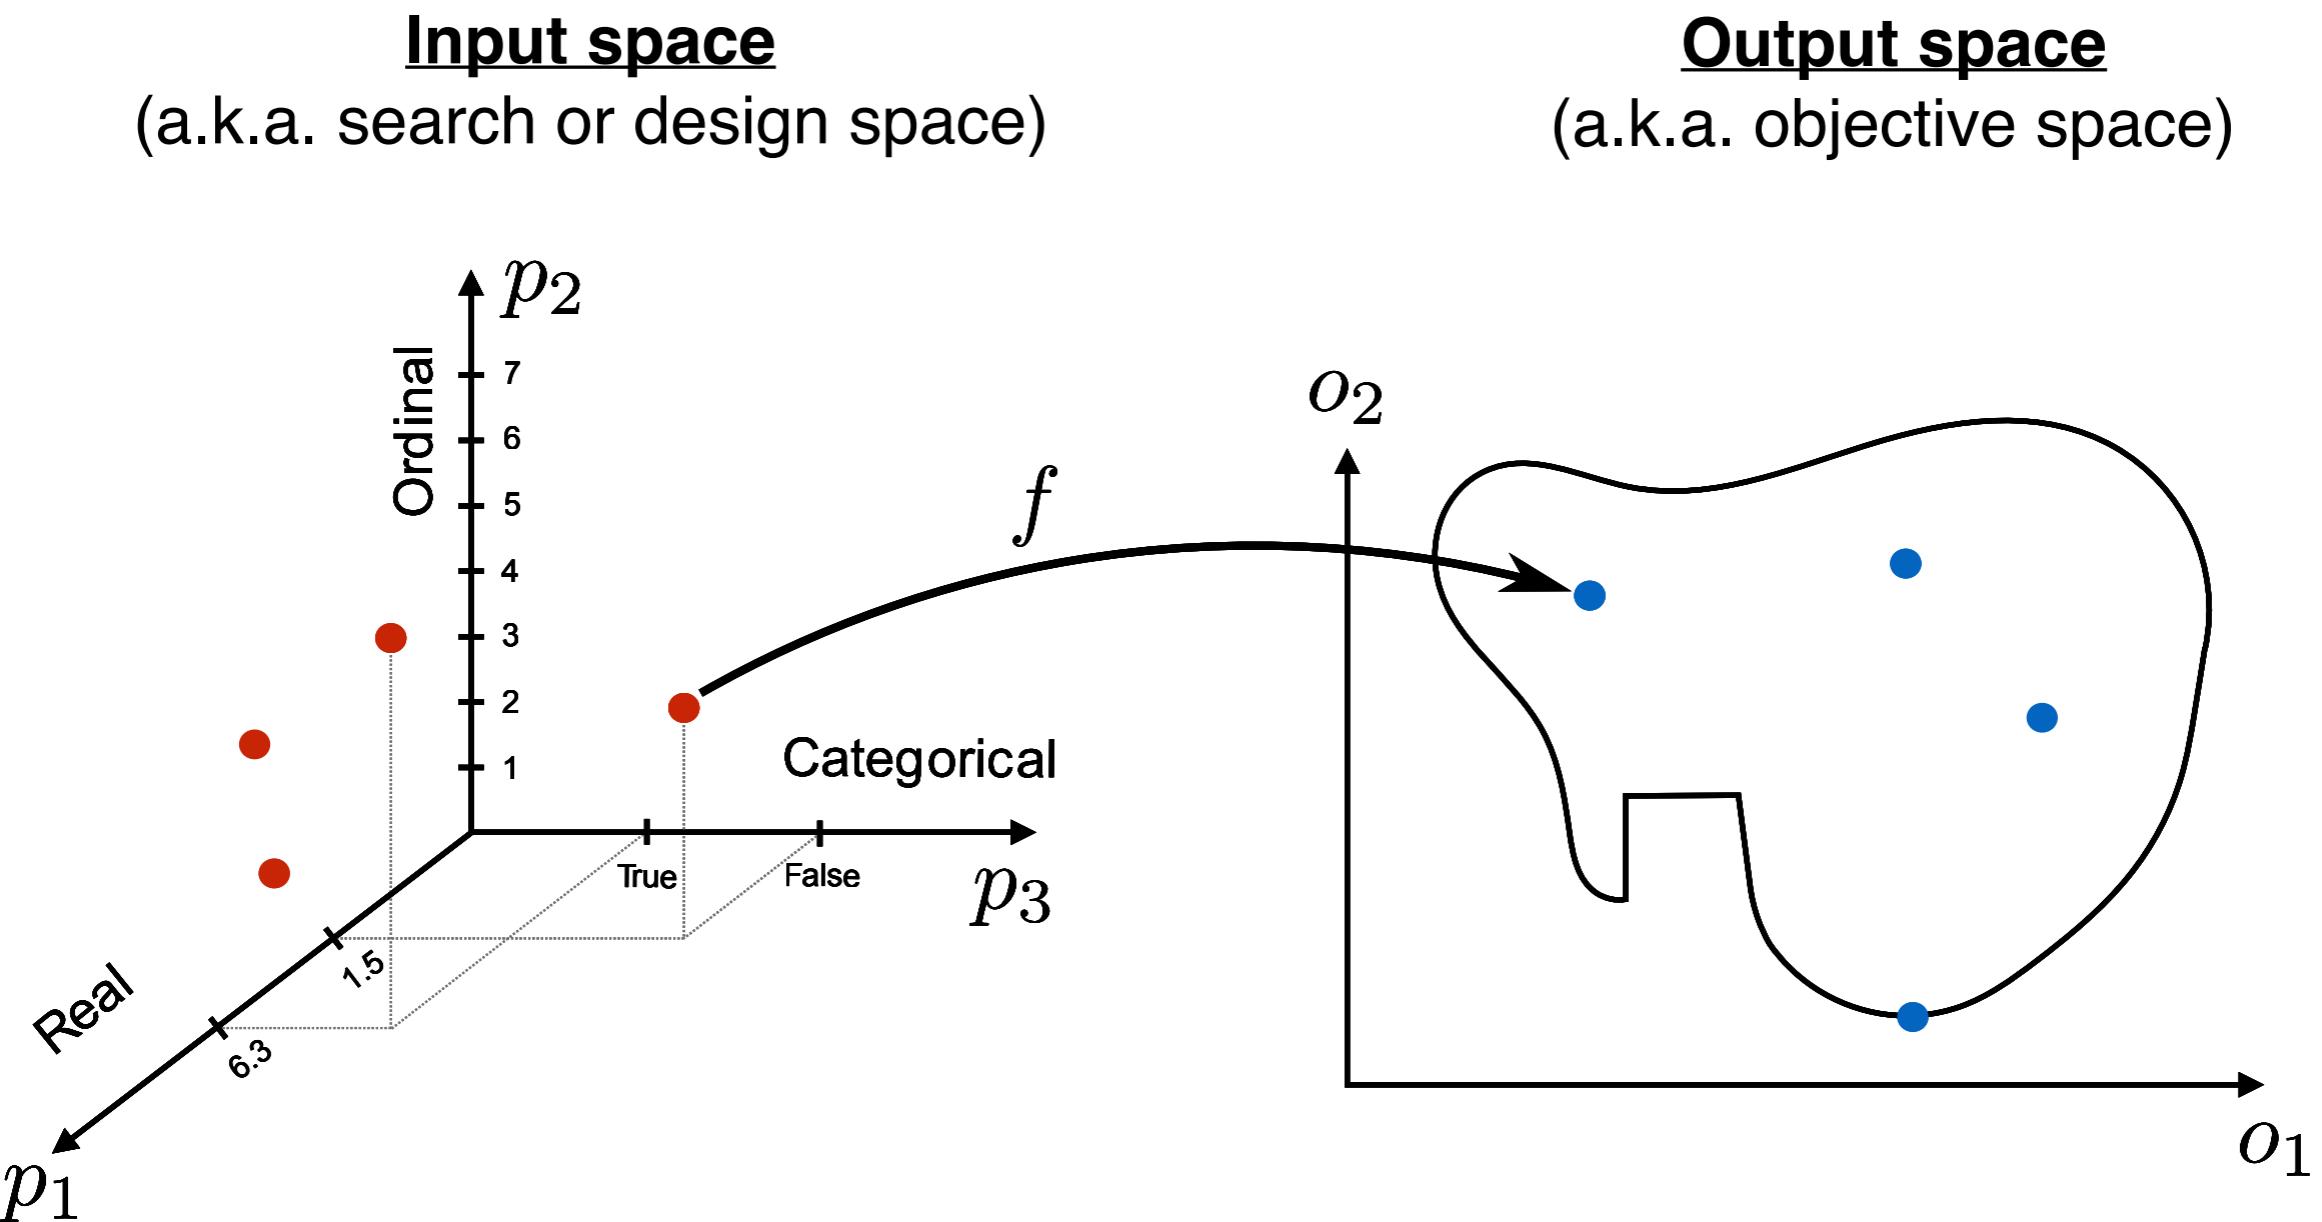
## 3-parameters and 2-objective Pictorial



Derivatives are unavailable: e.g., categorical variables

# Black-box Optimization (BBO)

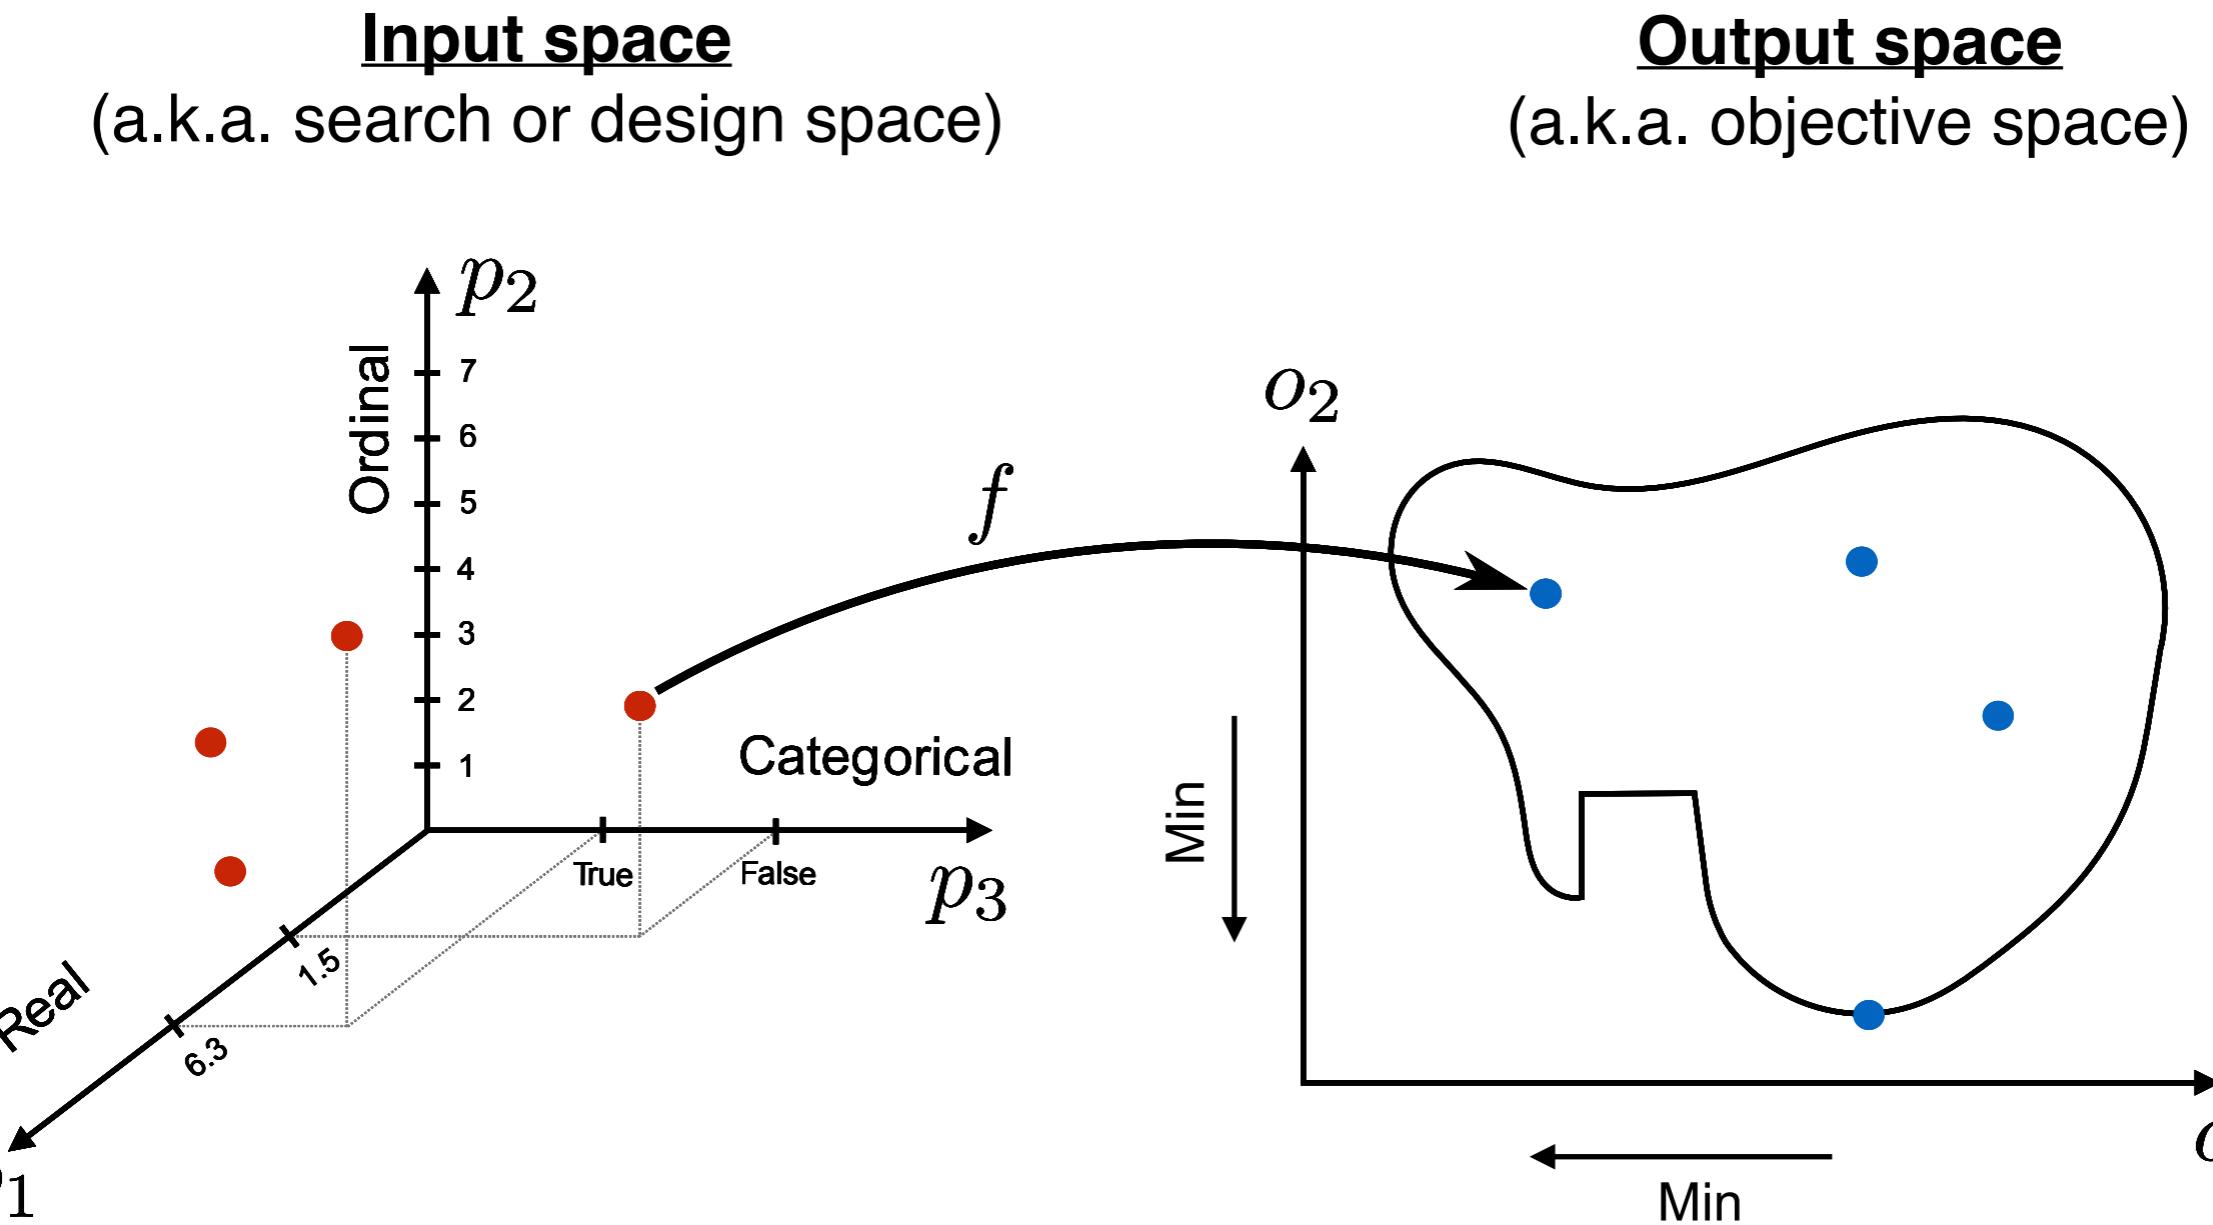
## 3-parameters and 2-objective Pictorial



Derivatives are unavailable: e.g., categorical variables

# Black-box Optimization (BBO)

## 3-parameters and 2-objective Pictorial

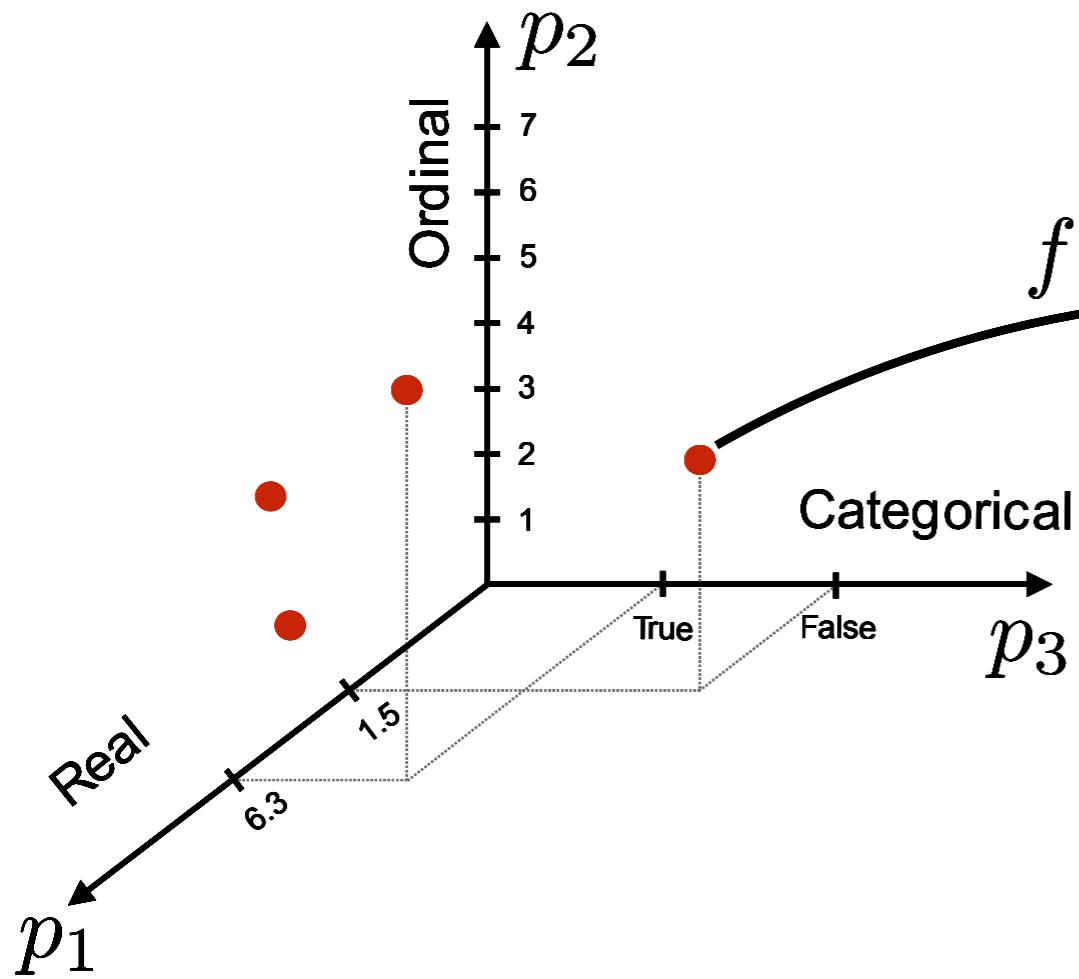


Derivatives are unavailable: e.g., categorical variables

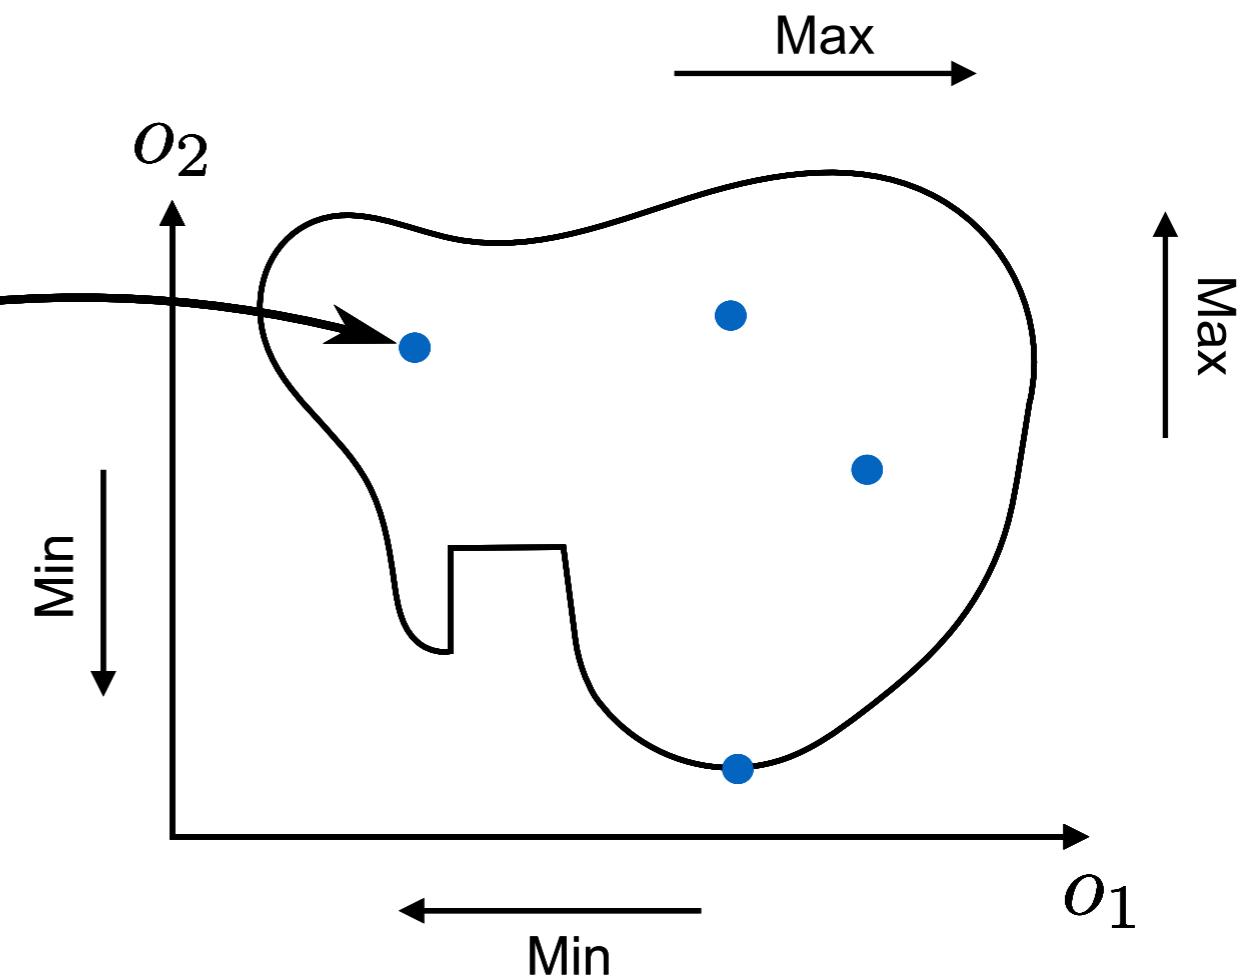
# Black-box Optimization (BBO)

## 3-parameters and 2-objective Pictorial

**Input space**  
(a.k.a. search or design space)



**Output space**  
(a.k.a. objective space)



Derivatives are unavailable: e.g., categorical variables

# Black-box Optimization

- Define an objective function. Examples:
  - Accuracy, runtime, throughput, latency, area, energy, etc.

# Black-box Optimization

- Define an objective function. Examples:
  - Accuracy, runtime, throughput, latency, area, energy, etc.
- Standard search procedures:
  - Grid search
  - Random search (very simple, works surprisingly well)
  - Heuristics: e.g., genetic algorithms
  - Black magic

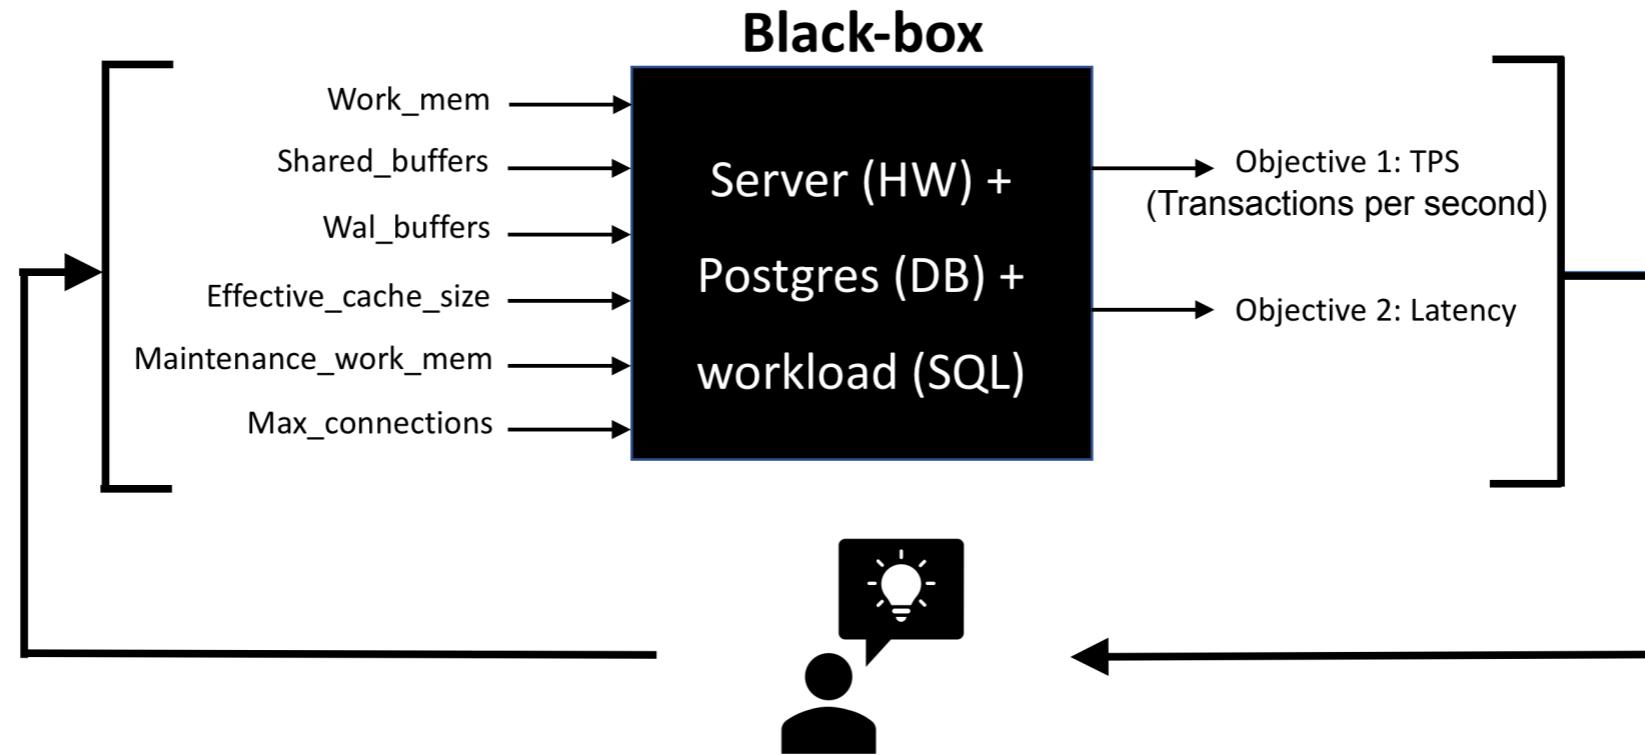
# Black-box Optimization

- Define an objective function. Examples:
  - Accuracy, runtime, throughput, latency, area, energy, etc.
- Standard search procedures:
  - Grid search
  - Random search (very simple, works surprisingly well)
  - Heuristics: e.g., genetic algorithms
  - Black magic
- Painful:
  - Search may be very expensive (e.g., time or money)
  - Usually limited sampling budget
  - Possibly noisy, unknown feasibility constraints, multi-objective

# Black-box Optimization

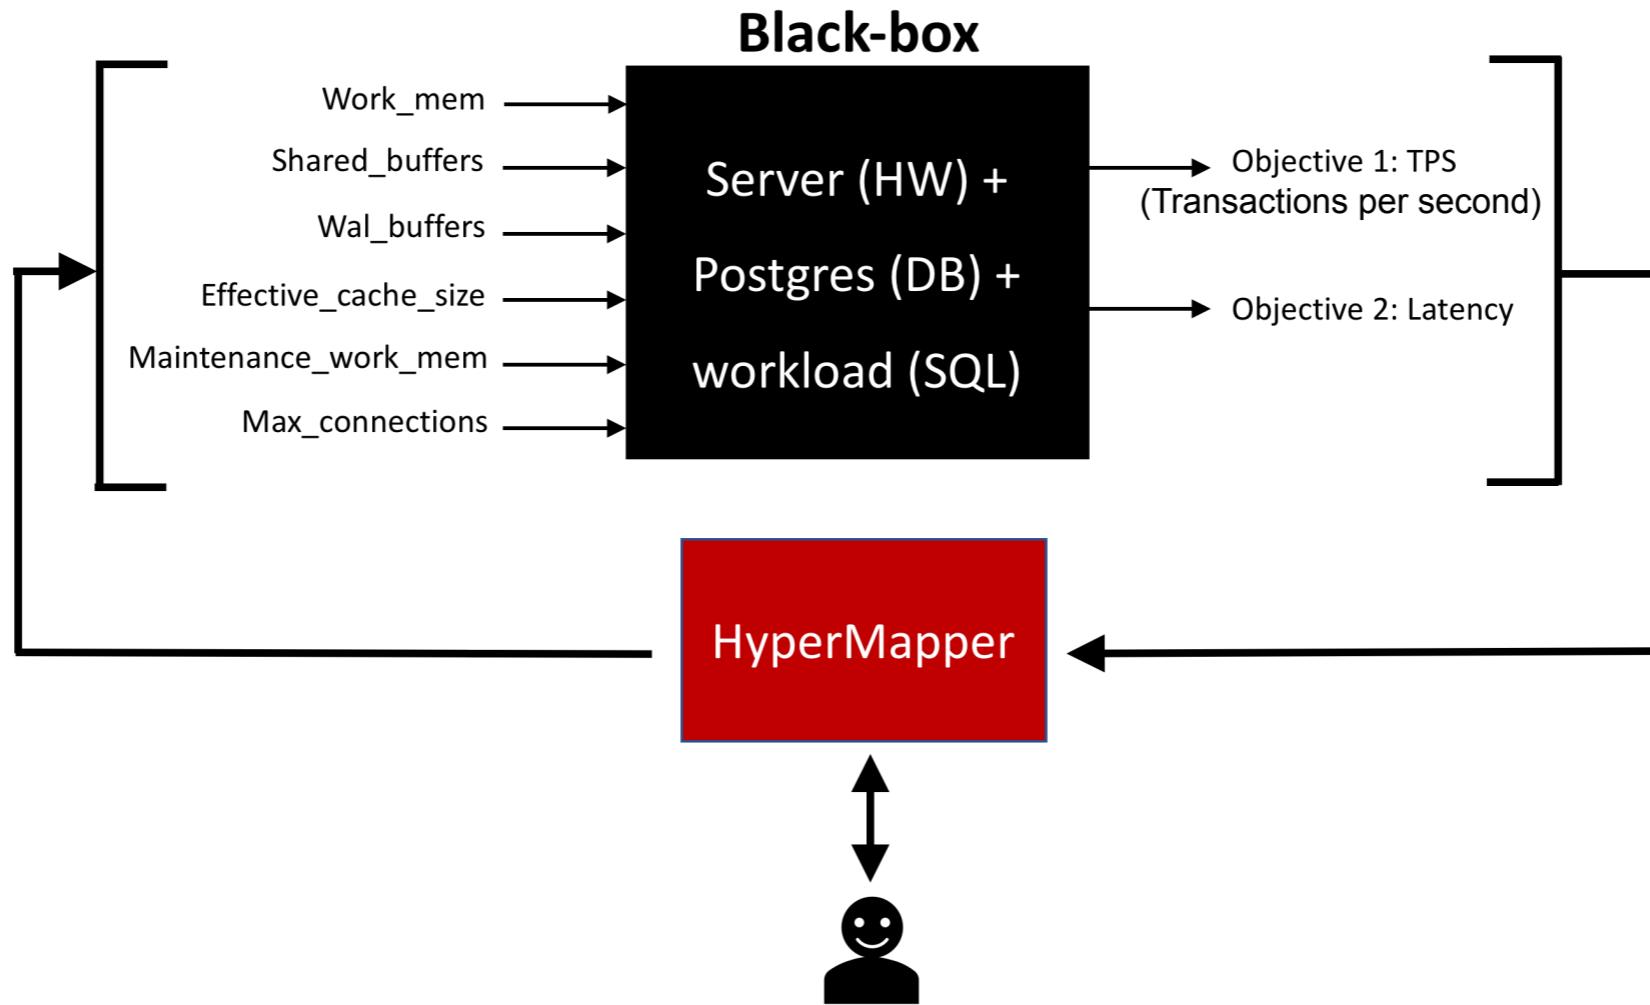
- Define an objective function. Examples:
  - Accuracy, runtime, throughput, latency, area, energy, etc.
- Standard search procedures:
  - Grid search
  - Random search (very simple, works surprisingly well)
  - Heuristics: e.g., genetic algorithms
  - Black magic
- Painful:
  - Search may be very expensive (e.g., time or money)
  - Usually limited sampling budget
  - Possibly noisy, unknown feasibility constraints, multi-objective
- Alternative approach: Bayesian Optimization (BO)

# Example: PostgreSQL/Azure Autotuning



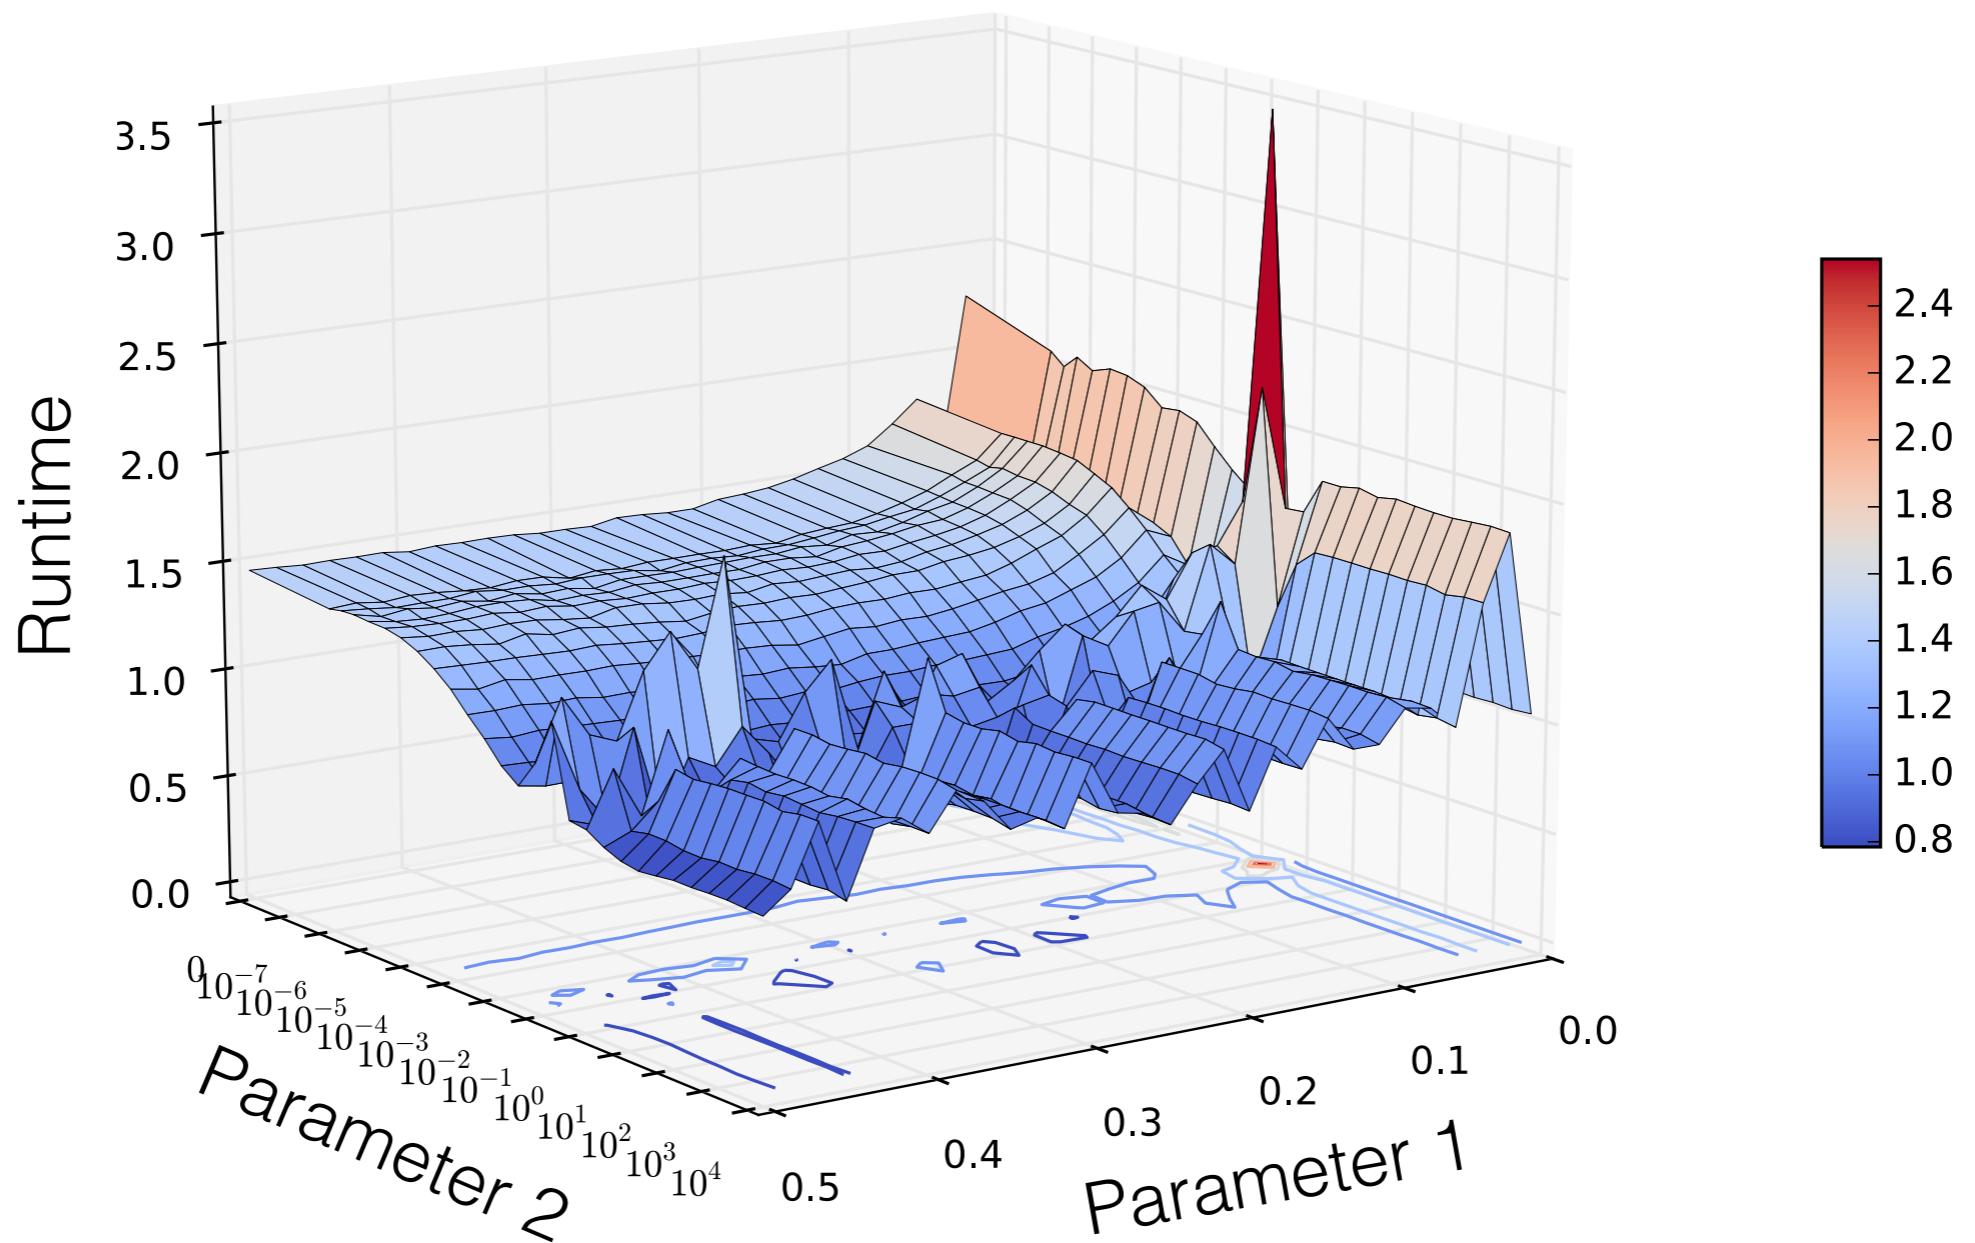
- Estimate of design space size:  $1e+18$
- This is an on going collaboration with **Microsoft** Seattle
- Need to optimize anytime the black-box changes: server, PostgreSQL, DB, queries
- **100,000 PostgreSQL databases** in Microsoft Azure

# Example: PostgreSQL/Azure Autotuning



- Estimate of design space size:  $1e+18$
- This is an on going collaboration with **Microsoft** Seattle
- Need to optimize anytime the black-box changes: server, PostgreSQL, DB, queries
- **100,000 PostgreSQL databases** in Microsoft Azure

# Motivation - Mono-objective



- Benchmark: SLAMBench 1.0 runtime response surface is:  
non-linear, multi-modal and non-smooth

# Practical BBO: Important Features

1. Real, integer, ordinal and categorical variables (RIOC var.)

- Example: parallelism is categorical; # of CPUs is ordinal

# Practical BBO: Important Features

1. Real, integer, ordinal and categorical variables (RIOC var.)

- Example: parallelism is categorical; # of CPUs is ordinal

2. User prior knowledge in the search (Prior)

- Example: # threads on a CPU with 4 cores? Gaussian-shaped around 4

# Practical BBO: Important Features

1. Real, integer, ordinal and categorical variables (RIOC var.)

- Example: parallelism is categorical; # of CPUs is ordinal

2. User prior knowledge in the search (Prior)

- Example: # threads on a CPU with 4 cores? Gaussian-shaped around 4

3. Multi-objective optimization (Multi)

- Example: trade-off accuracy and inference time

# Practical BBO: Important Features

1. Real, integer, ordinal and categorical variables (RIOC var.)

- Example: parallelism is categorical; # of CPUs is ordinal

2. User prior knowledge in the search (Prior)

- Example: # threads on a CPU with 4 cores? Gaussian-shaped around 4

3. Multi-objective optimization (Multi)

- Example: trade-off accuracy and inference time

4. Unknown feasibility constraints (Constr.)

- Example: black-box function crash; a design doesn't fit in the FPGA

# BBO Tools Taxonomy

**None of the tools available support all these BBO features**

We introduce a new framework dubbed **HyperMapper**

Name	Multi	RIOC var.	Constr.	Prior
GpyOpt	✗	✗	✗	✗
OpenTuner	✗	✓	✗	✗
SURF	✗	✓	✗	✗
SMAC	✗	✓	✗	✗
Spearmint	✗	✗	✓	✗
Hyperopt	✗	✓	✗	✓
Hyperband	✗	✓	✗	✗
GPflowOpt	✓	✗	✓	✗
cBO	✗	✗	✓	✗
BOHB	✗	✓	✗	✗
<b>HyperMapper</b>	✓	✓	✓	✓

[Bodin, et al.] "Integrating algorithmic parameters into benchmarking and design space exploration in 3D scene understanding", PACT, 2016.

[Nardi, et al.] "Algorithmic performance-accuracy trade-off in 3D vision applications using HyperMapper", IPDPS-iWAPT, 2017.

[Nardi, et al., 2019] "Practical Design Space Exploration", MASCOTS, 2019.

# Outline

## 1. Black-box Optimization

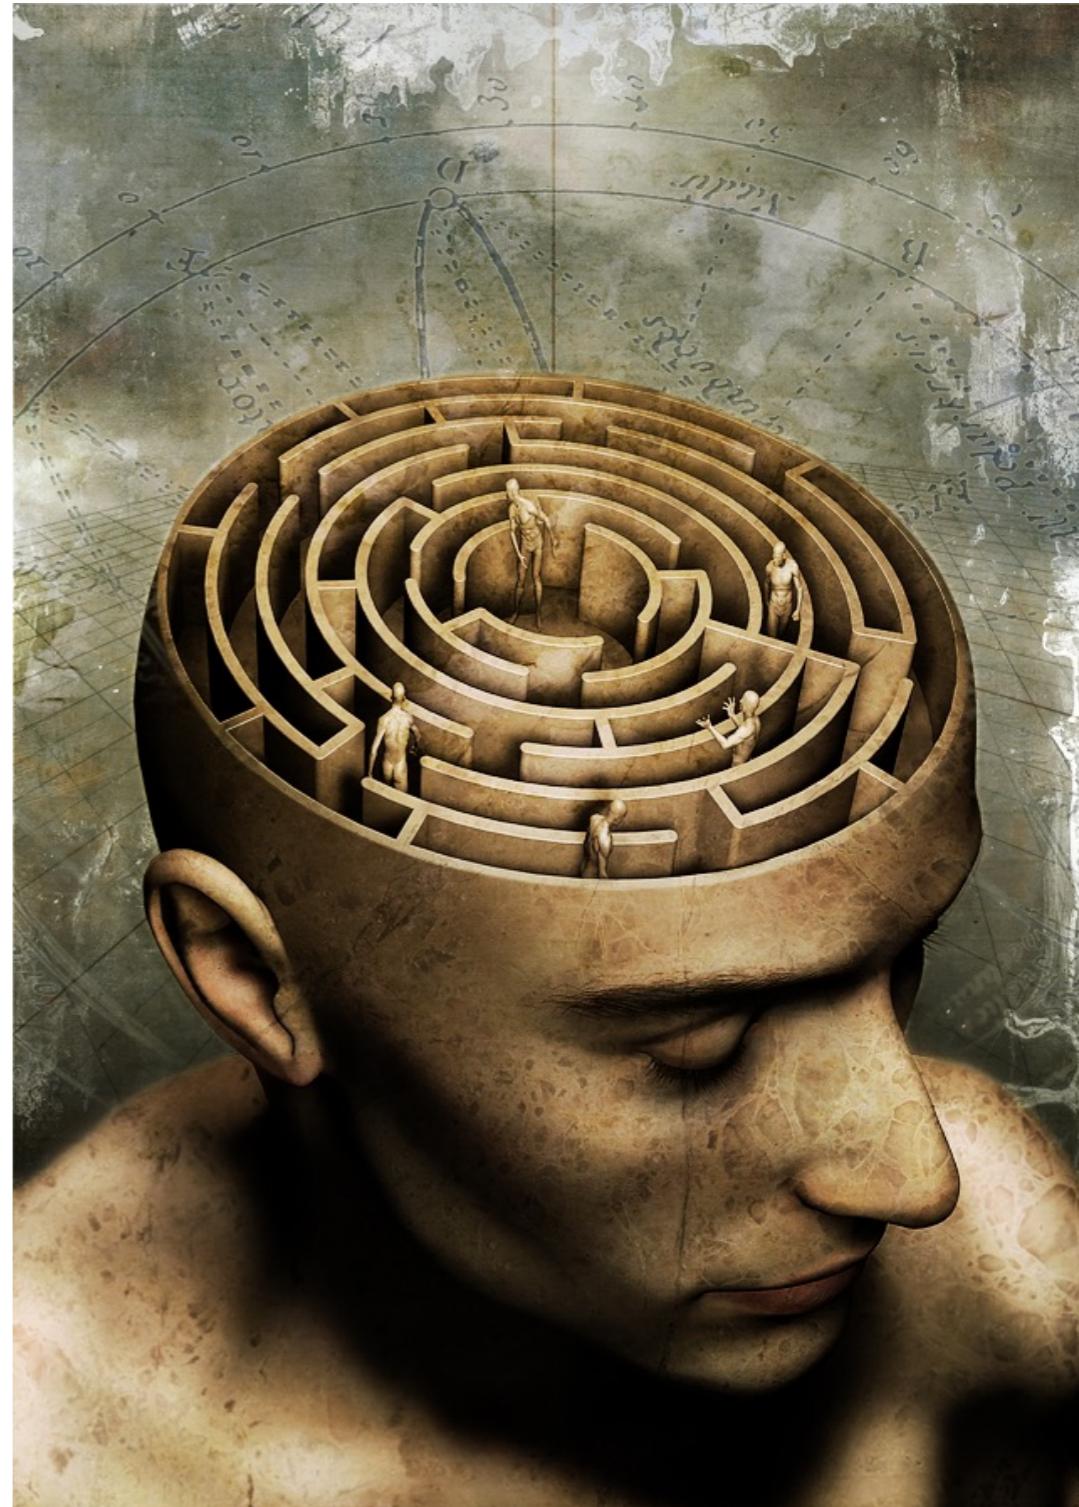
- Problem Setting
- Motivation Examples
- Taxonomy

## 2. The HyperMapper Framework

- Bayesian Optimization
- Pareto and Hypervolume Indicator
- Prior Distribution

## 3. Experimental Results

- Hardware Design
- Computer Vision
- Databases



# Bayesian Optimization

## Setting

Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

# Bayesian Optimization

## Setting

Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

- Find global minimum of objective function  $f$ :  $x^* = \arg \min_{x \in \mathbb{X}} f(x)$

# Bayesian Optimization

## Setting

Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

- Find global minimum of objective function  $f$ :  $x^* = \arg \min_{x \in \mathbb{X}} f(x)$
- We can evaluate the objective  $f$  pointwise
  - But we do not have an easy functional form or gradients

# Bayesian Optimization

## Setting

Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

- Find global minimum of objective function  $f$ :  $x^* = \arg \min_{x \in \mathbb{X}} f(x)$
- We can evaluate the objective  $f$  pointwise
  - But we do not have an easy functional form or gradients
- Build a **probabilistic proxy model** for  $f$ 
  - Using outcomes of past experiments as training data

# Bayesian Optimization

## Setting

Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

- Find global minimum of objective function  $f$ :  $x^* = \arg \min_{x \in \mathbb{X}} f(x)$
- We can evaluate the objective  $f$  pointwise
  - But we do not have an easy functional form or gradients
- Build a **probabilistic proxy model** for  $f$ 
  - Using outcomes of past experiments as training data
- Proxy model is **cheaper to evaluate** than the objective

# Bayesian Optimization

## Setting

Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

- Find global minimum of objective function  $f$ :  $x^* = \arg \min_{x \in \mathbb{X}} f(x)$
- We can evaluate the objective  $f$  pointwise
  - But we do not have an easy functional form or gradients
- Build a **probabilistic proxy model** for  $f$ 
  - Using outcomes of past experiments as training data
- Proxy model is **cheaper to evaluate** than the objective
- **Optimize cheap proxy** function
  - To determine where to evaluate the true objective next

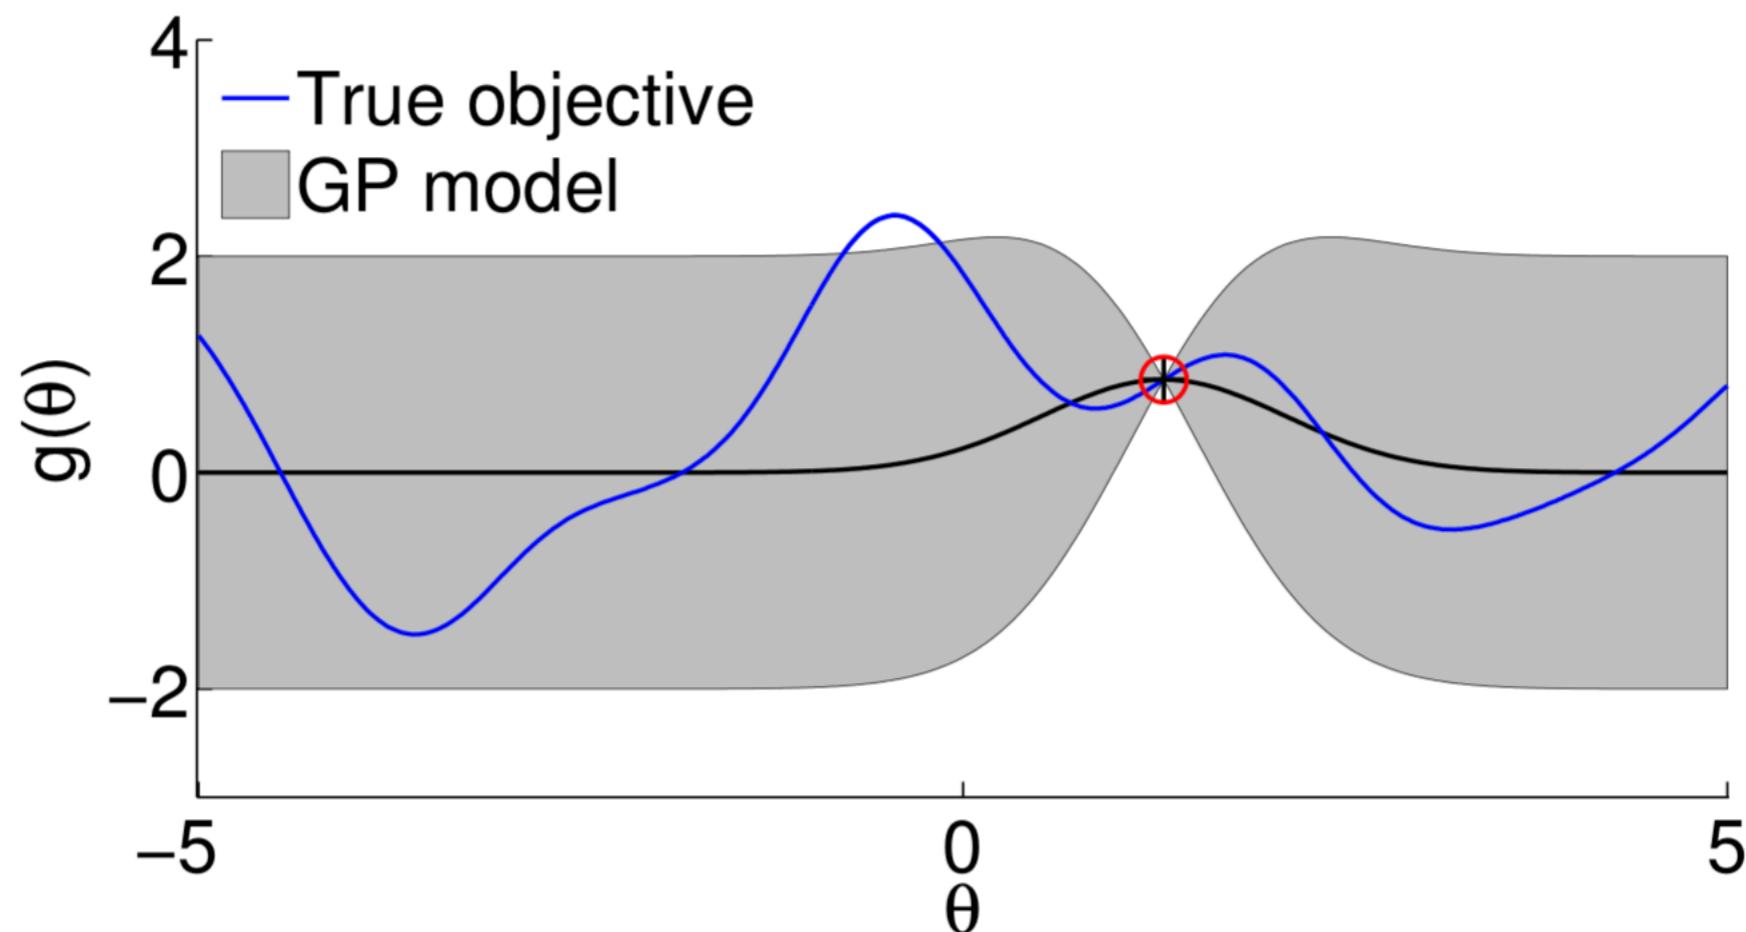
# Bayesian Optimization

## Setting

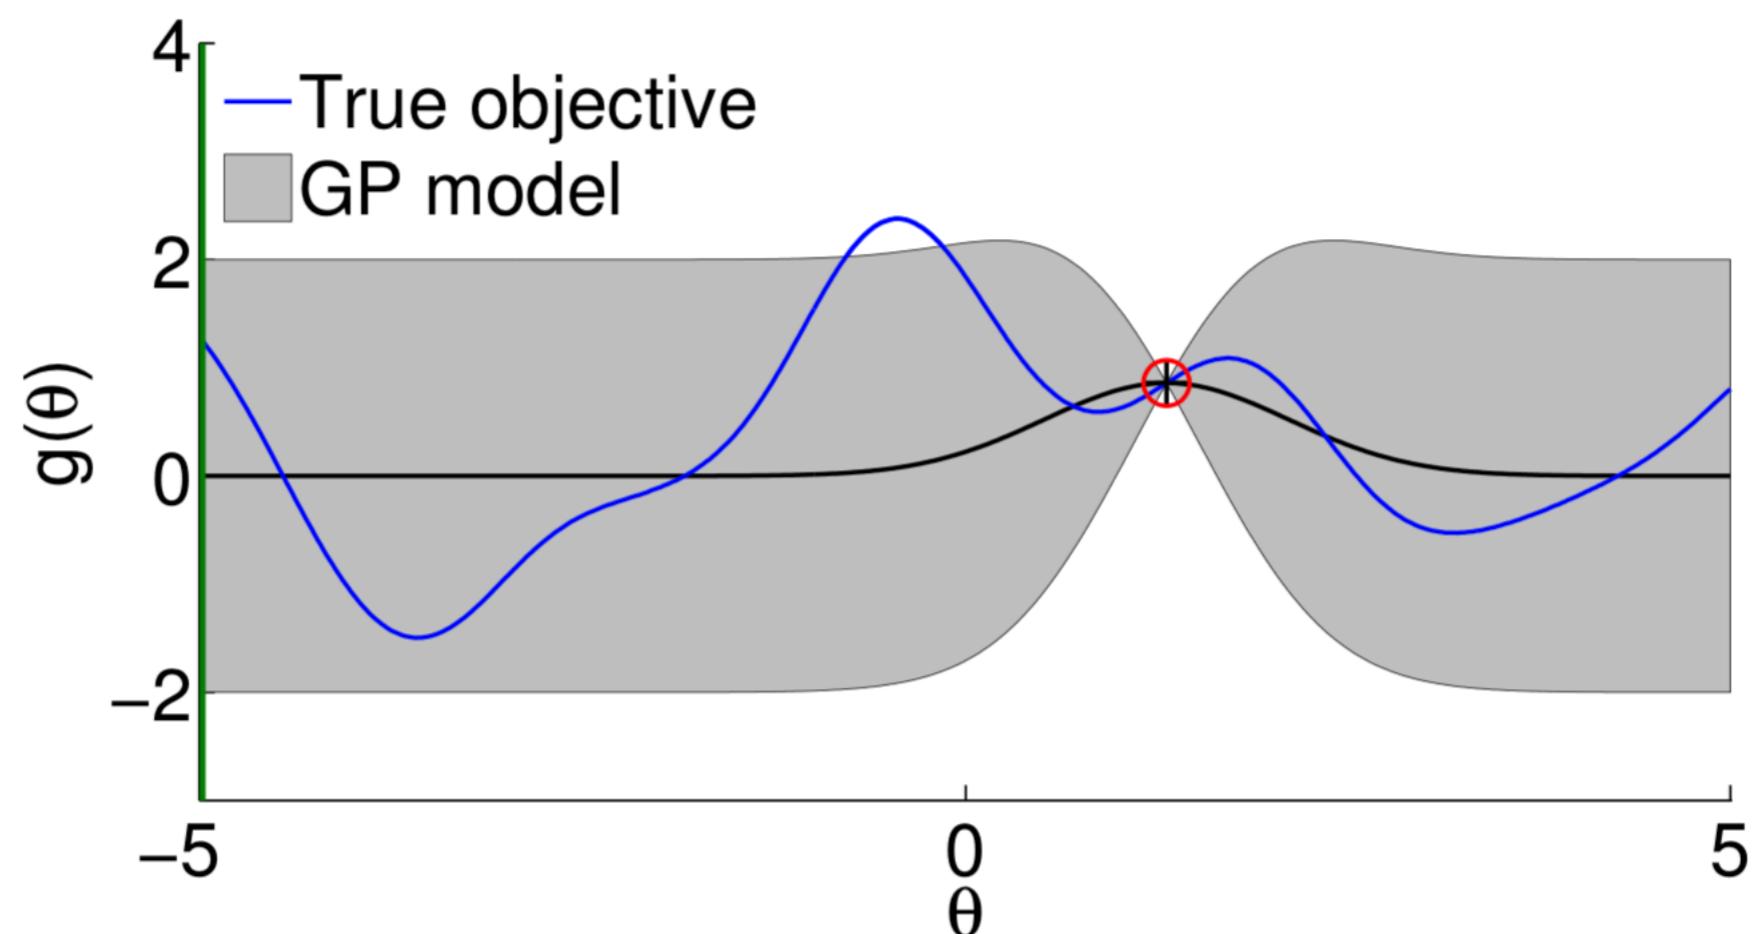
Globally optimize an objective function that is expensive to evaluate  
(e.g., PostgreSQL example above)

- Find global minimum of objective function  $f$ :  $x^* = \arg \min_{x \in \mathbb{X}} f(x)$
- We can evaluate the objective  $f$  pointwise
  - But we do not have an easy functional form or gradients
- Build a **probabilistic proxy model** for  $f$ 
  - Using outcomes of past experiments as training data
- Proxy model is **cheaper to evaluate** than the objective
- **Optimize cheap proxy** function
  - To determine where to evaluate the true objective next
- Standard proxy: Gaussian process
  - But also: Random Forests, Tree Parzen Estimators, etc.

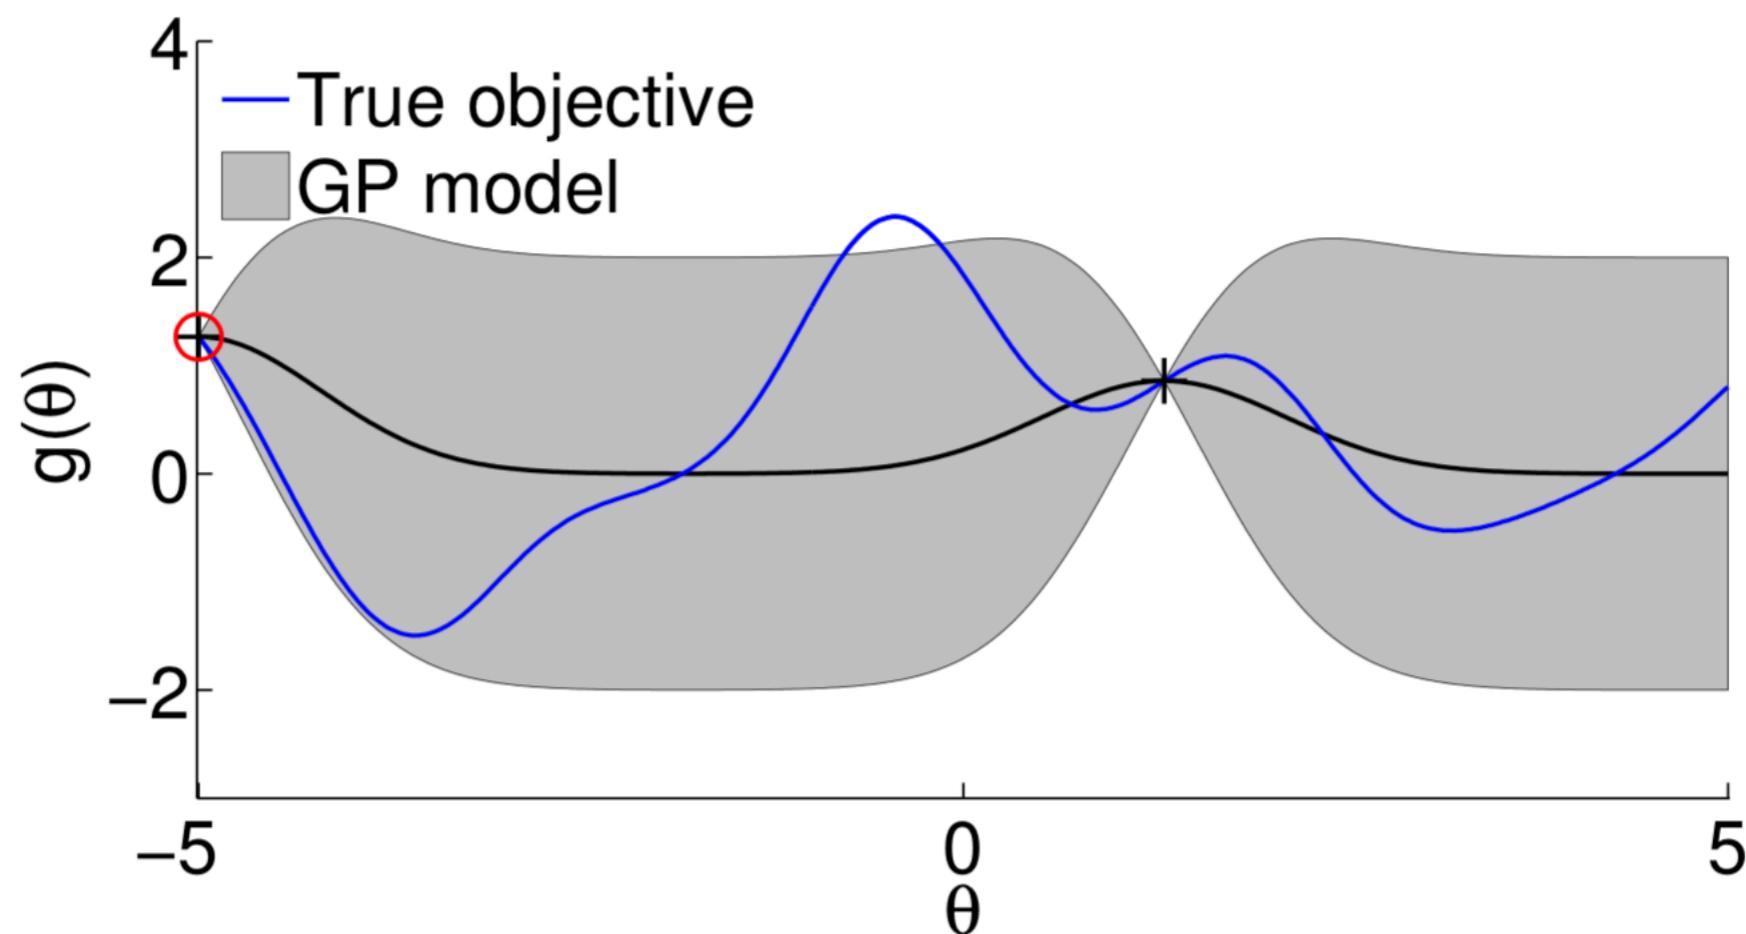
# Bayesian Optimization: Illustration



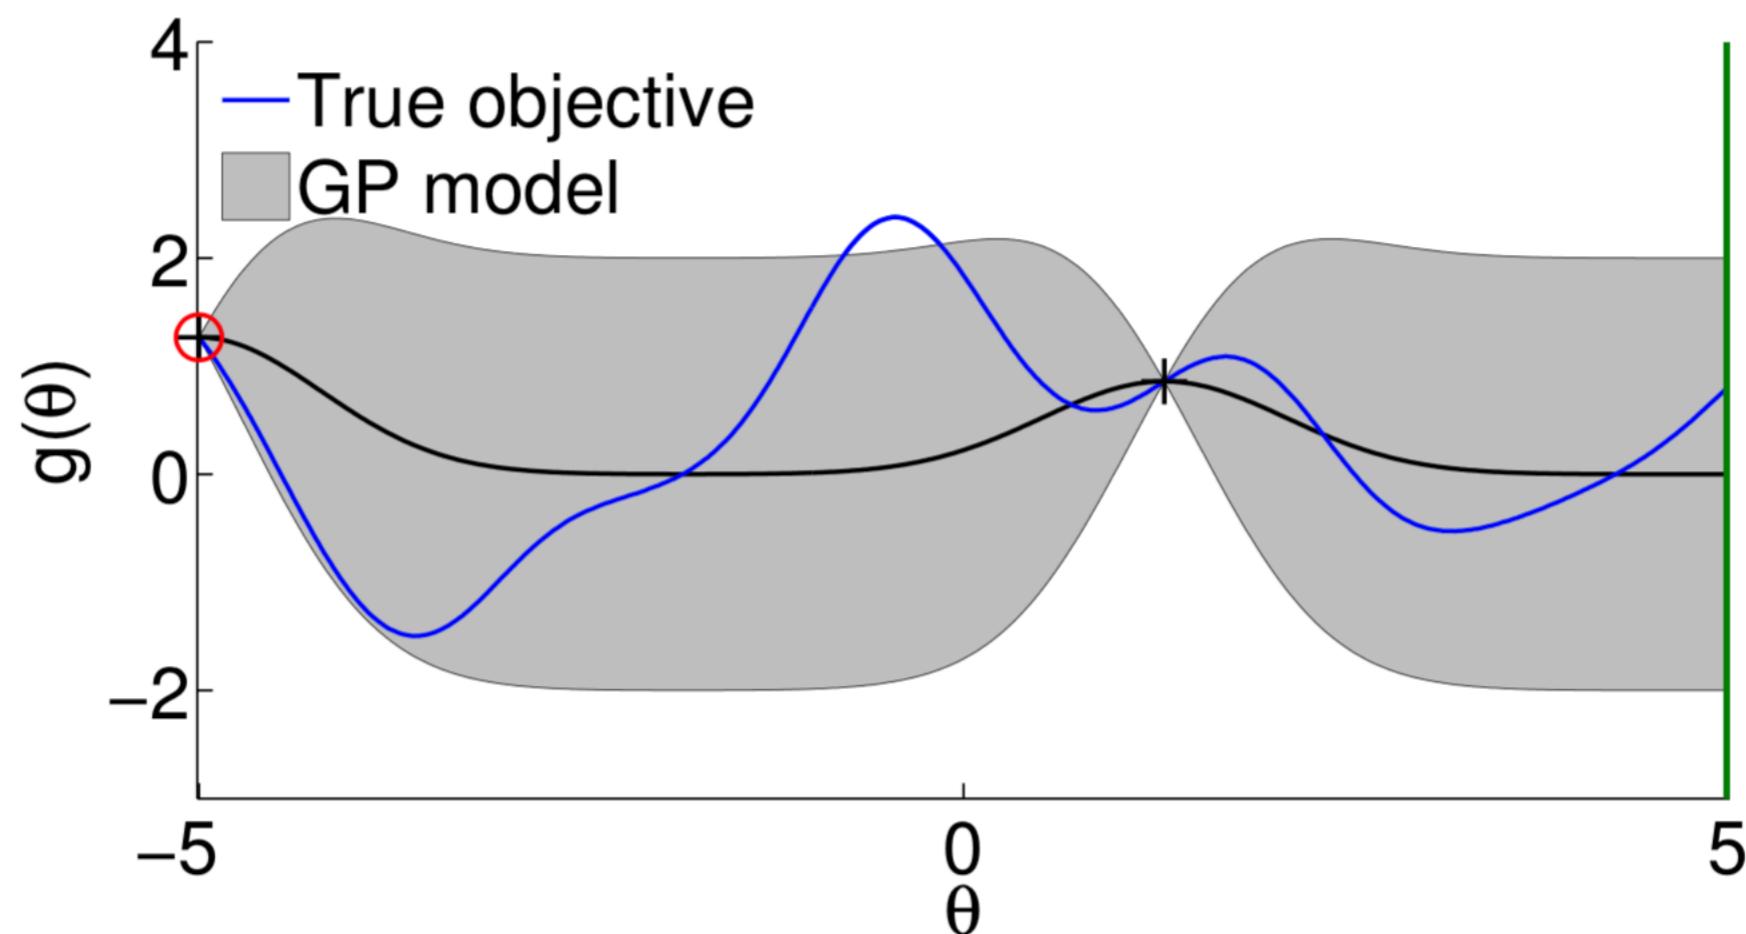
# Bayesian Optimization: Illustration



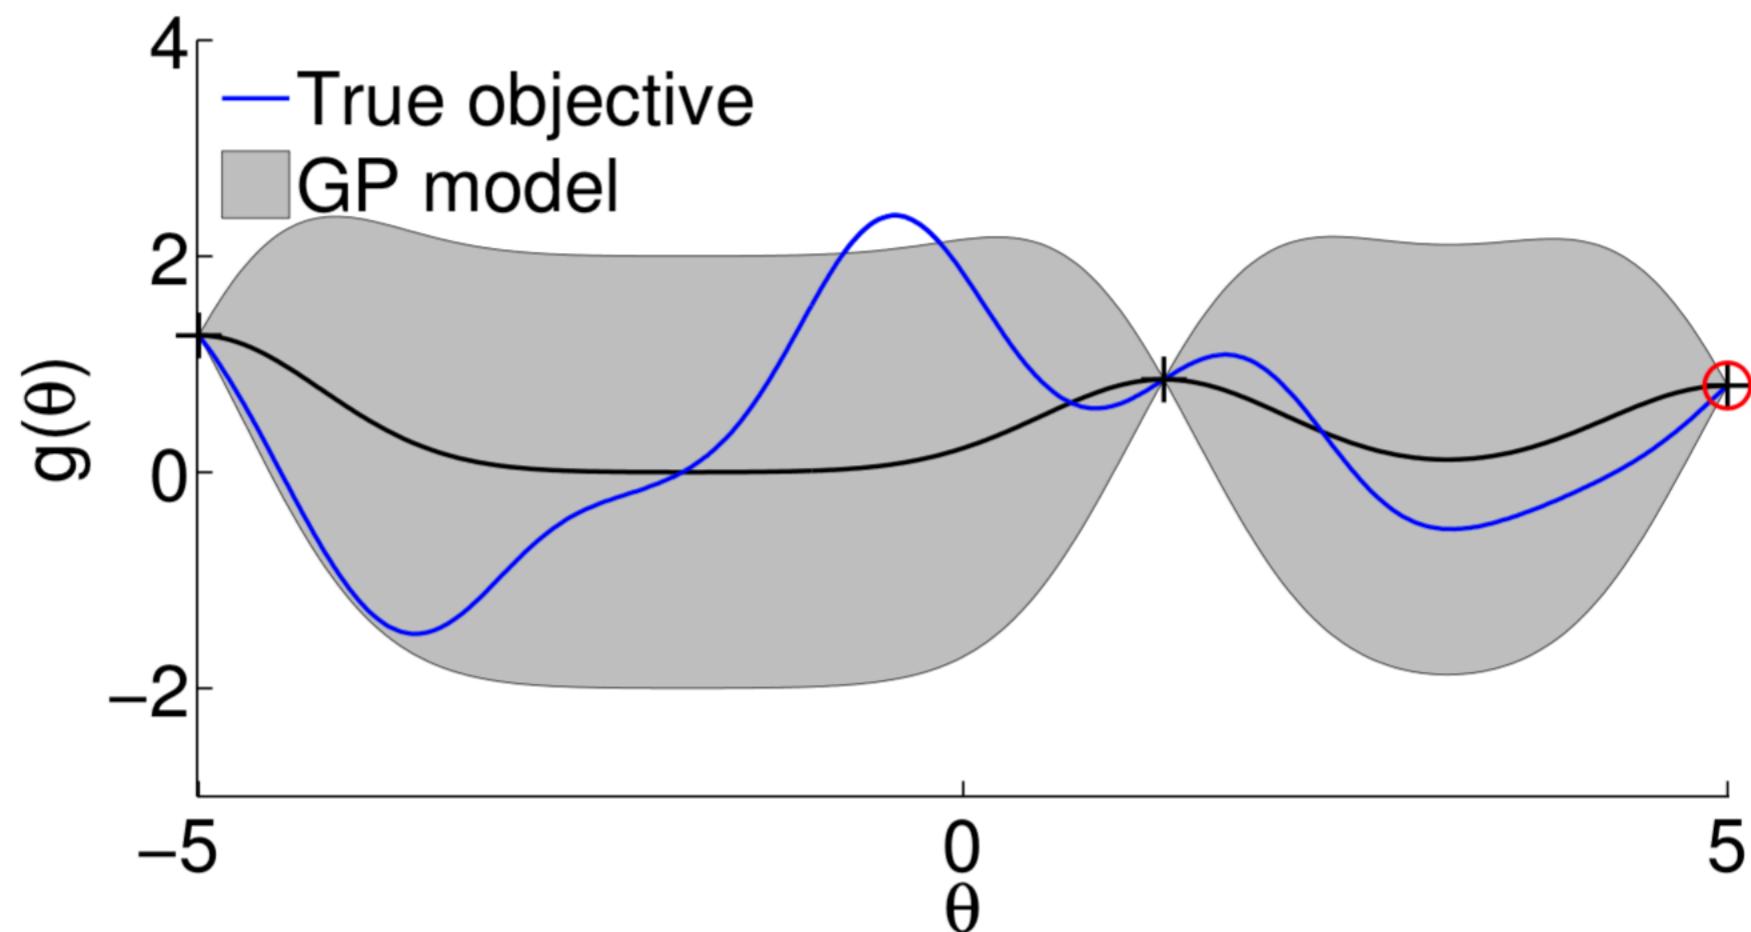
# Bayesian Optimization: Illustration



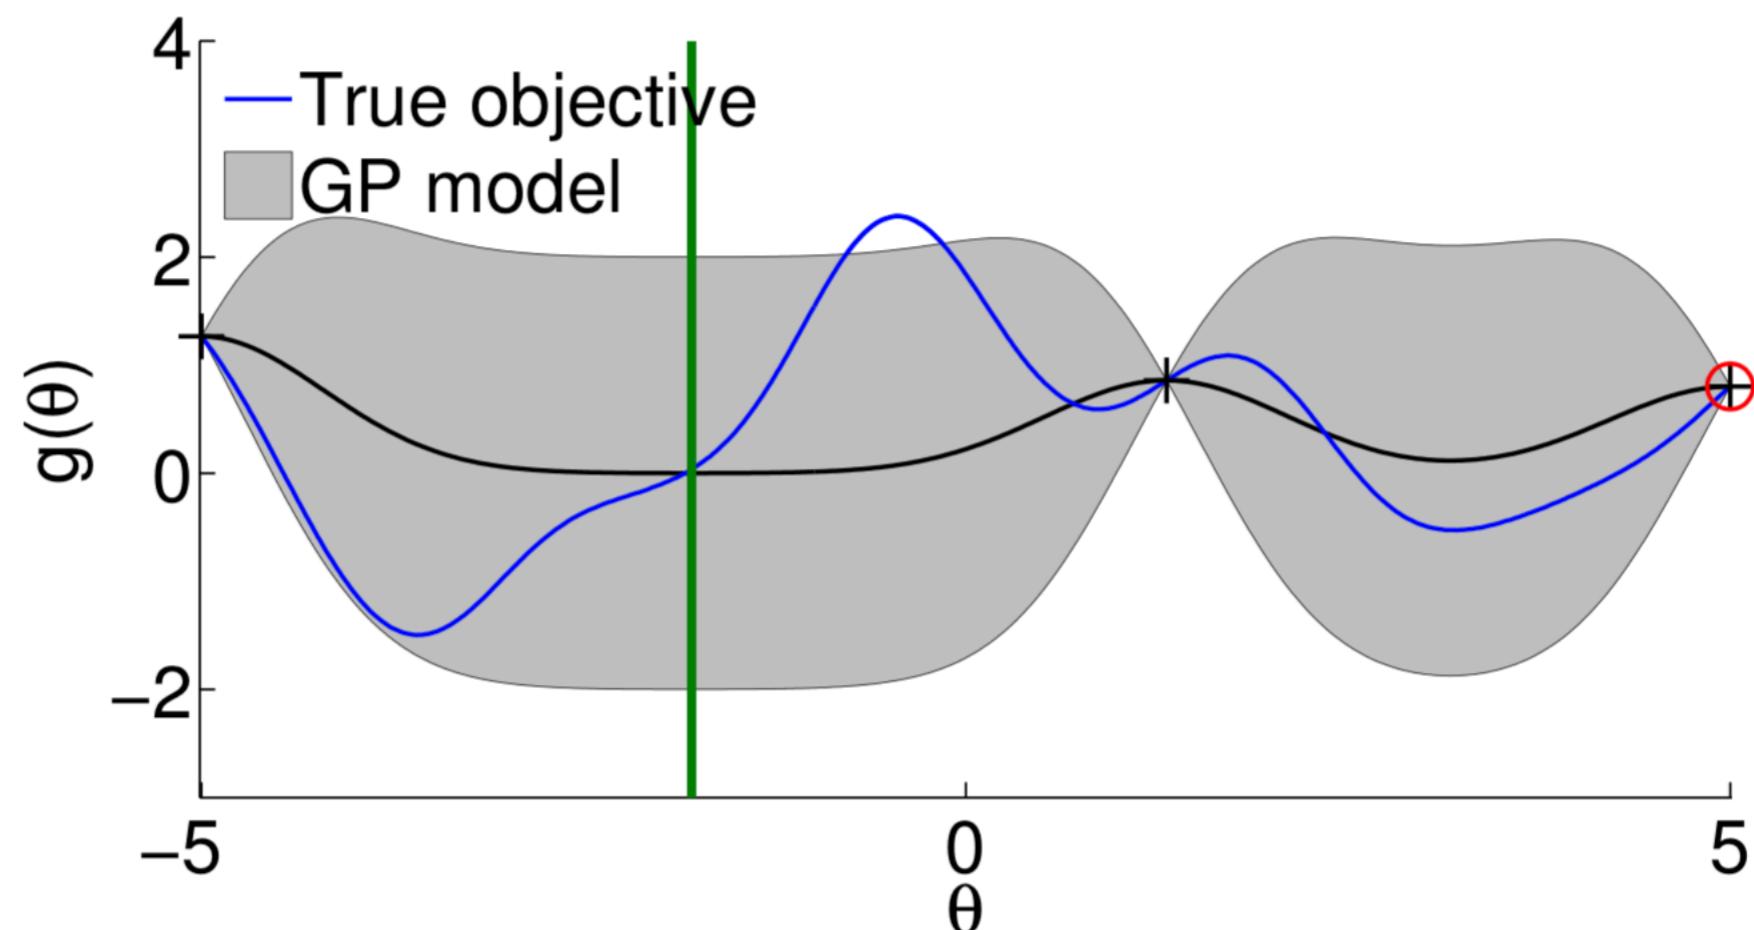
# Bayesian Optimization: Illustration



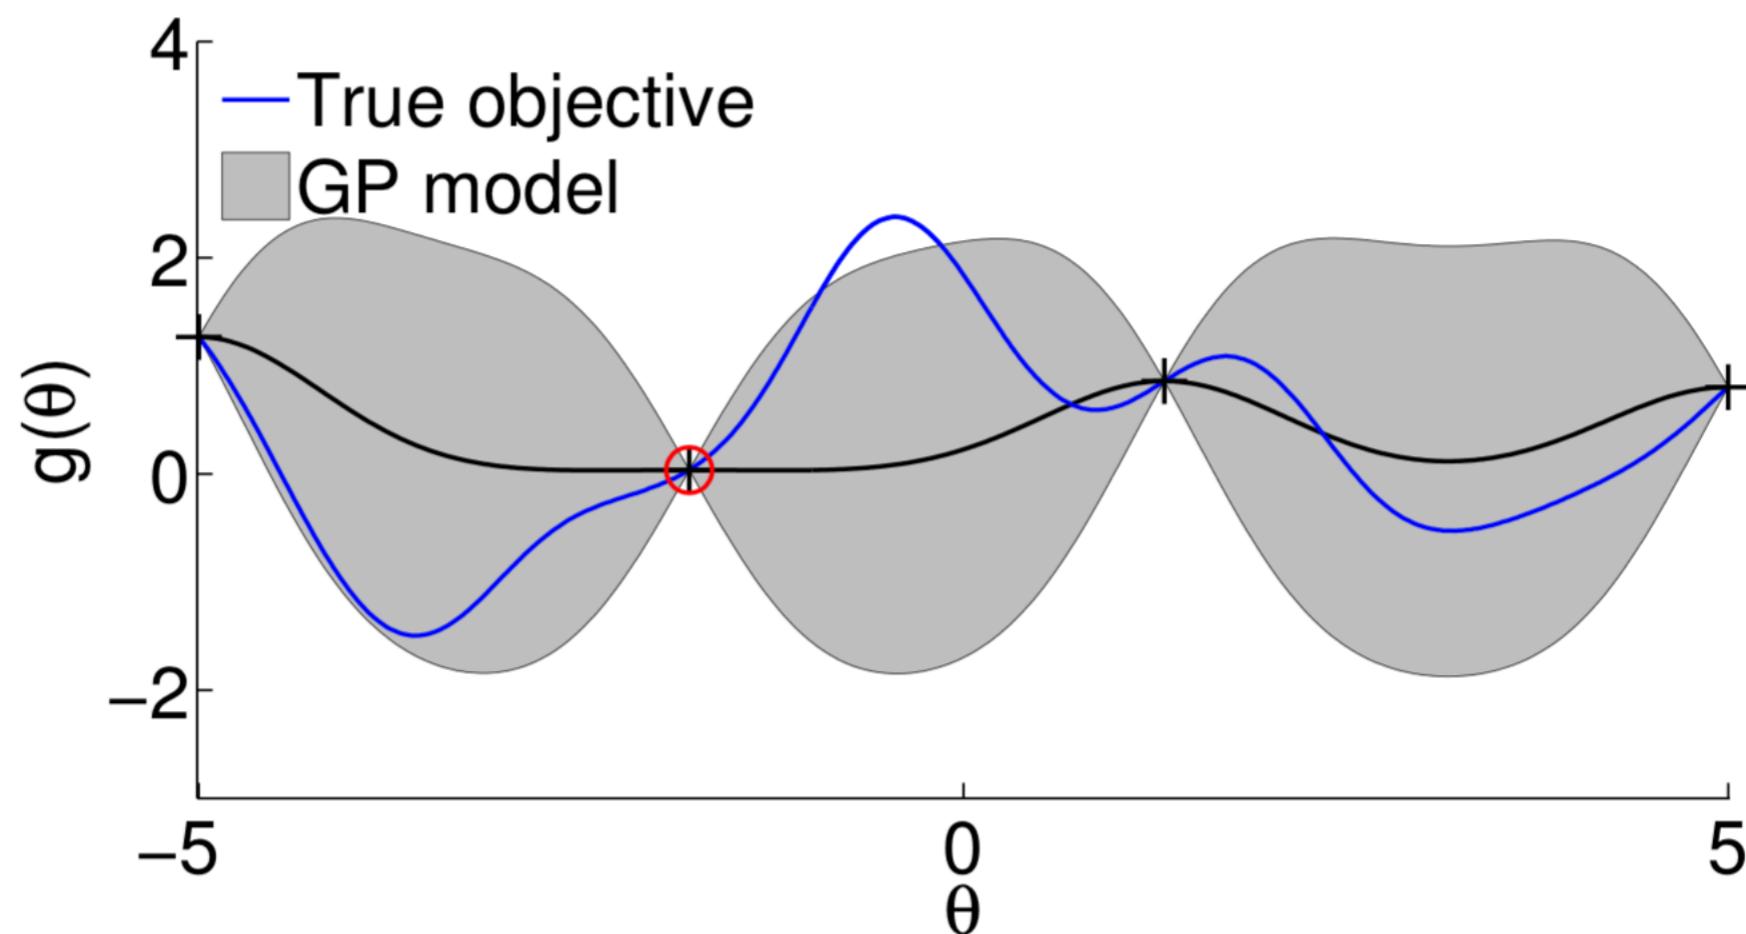
# Bayesian Optimization: Illustration



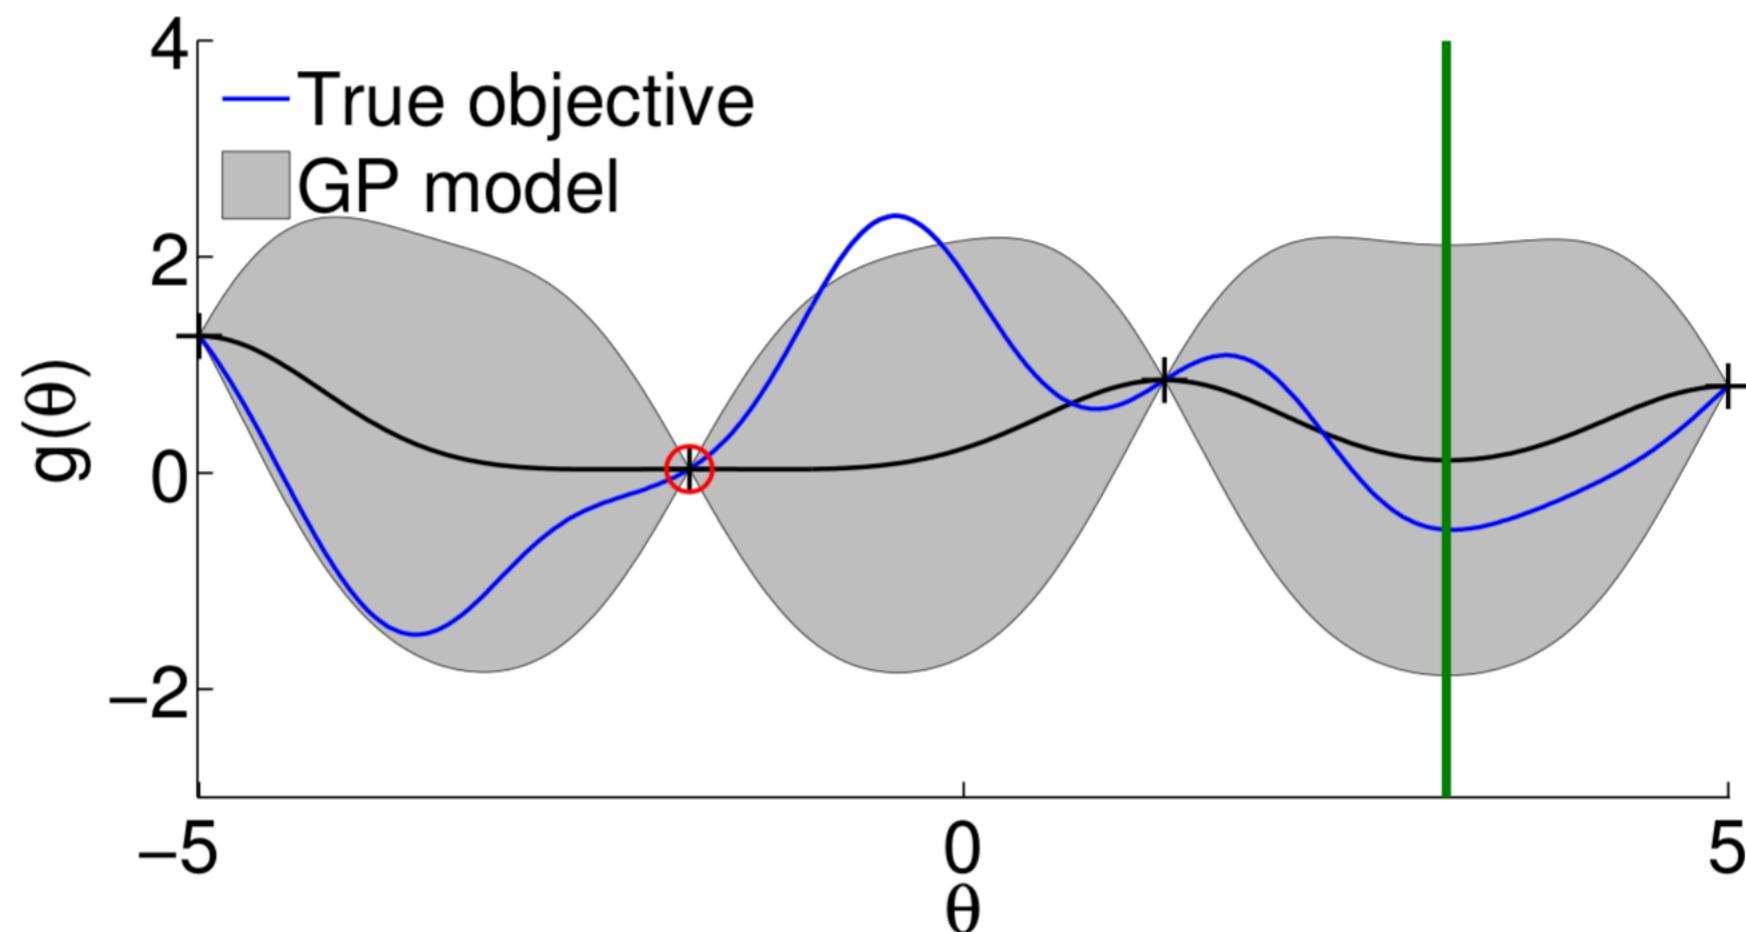
# Bayesian Optimization: Illustration



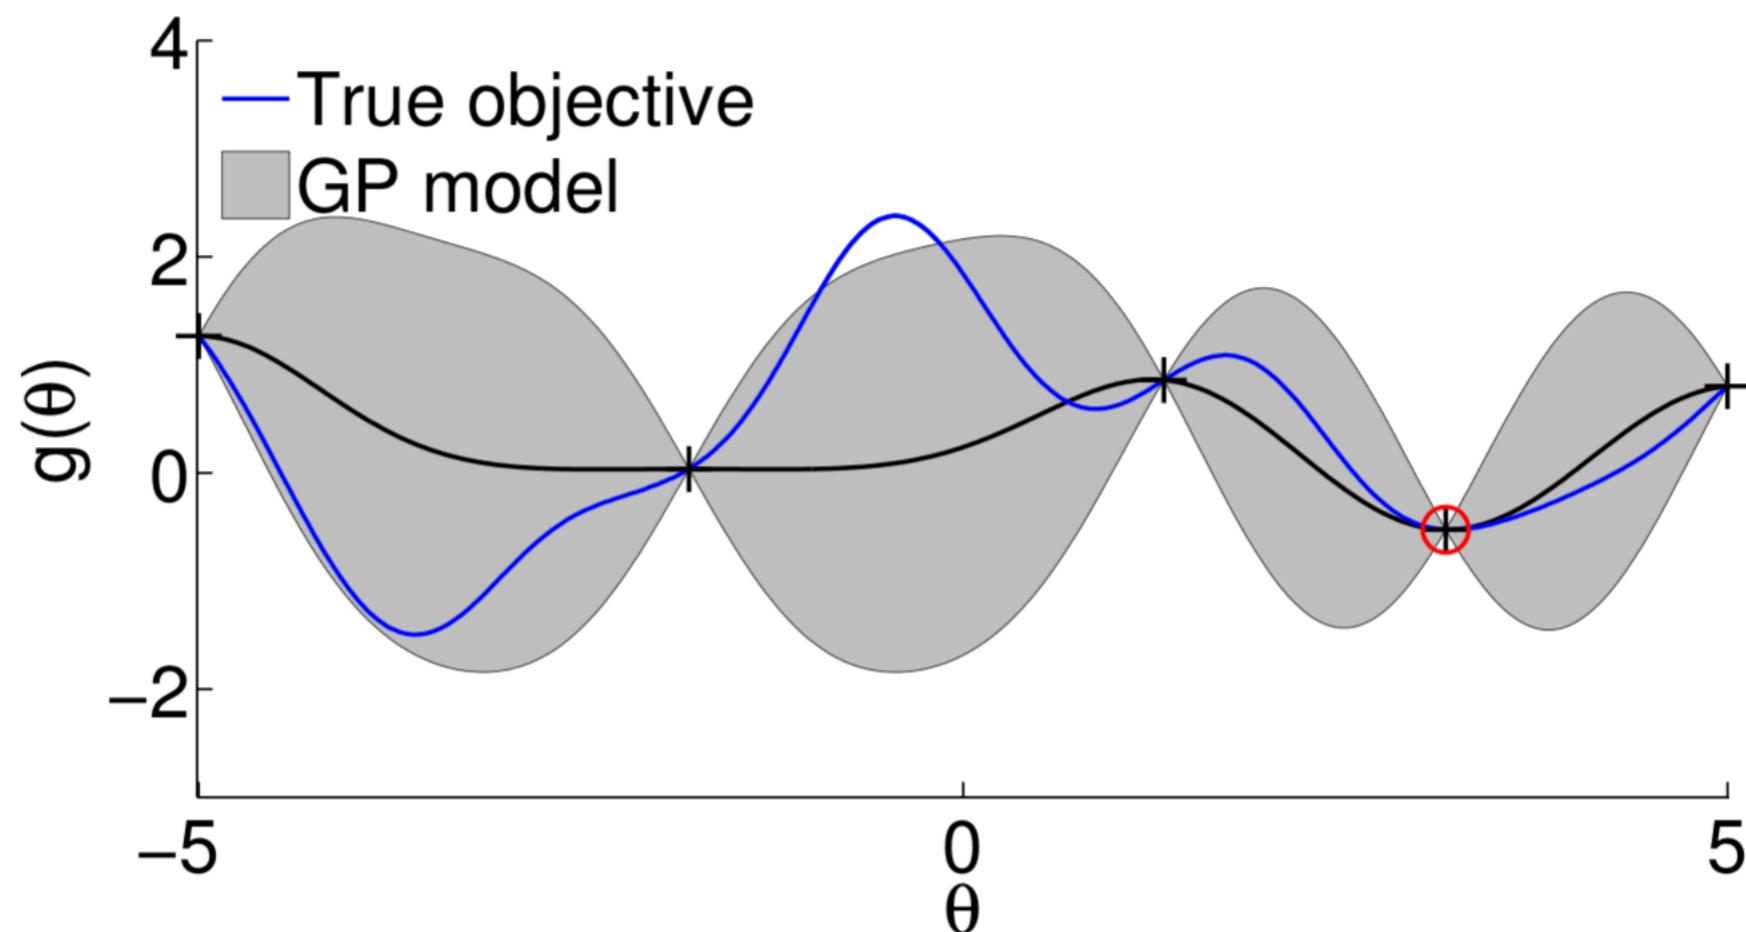
# Bayesian Optimization: Illustration



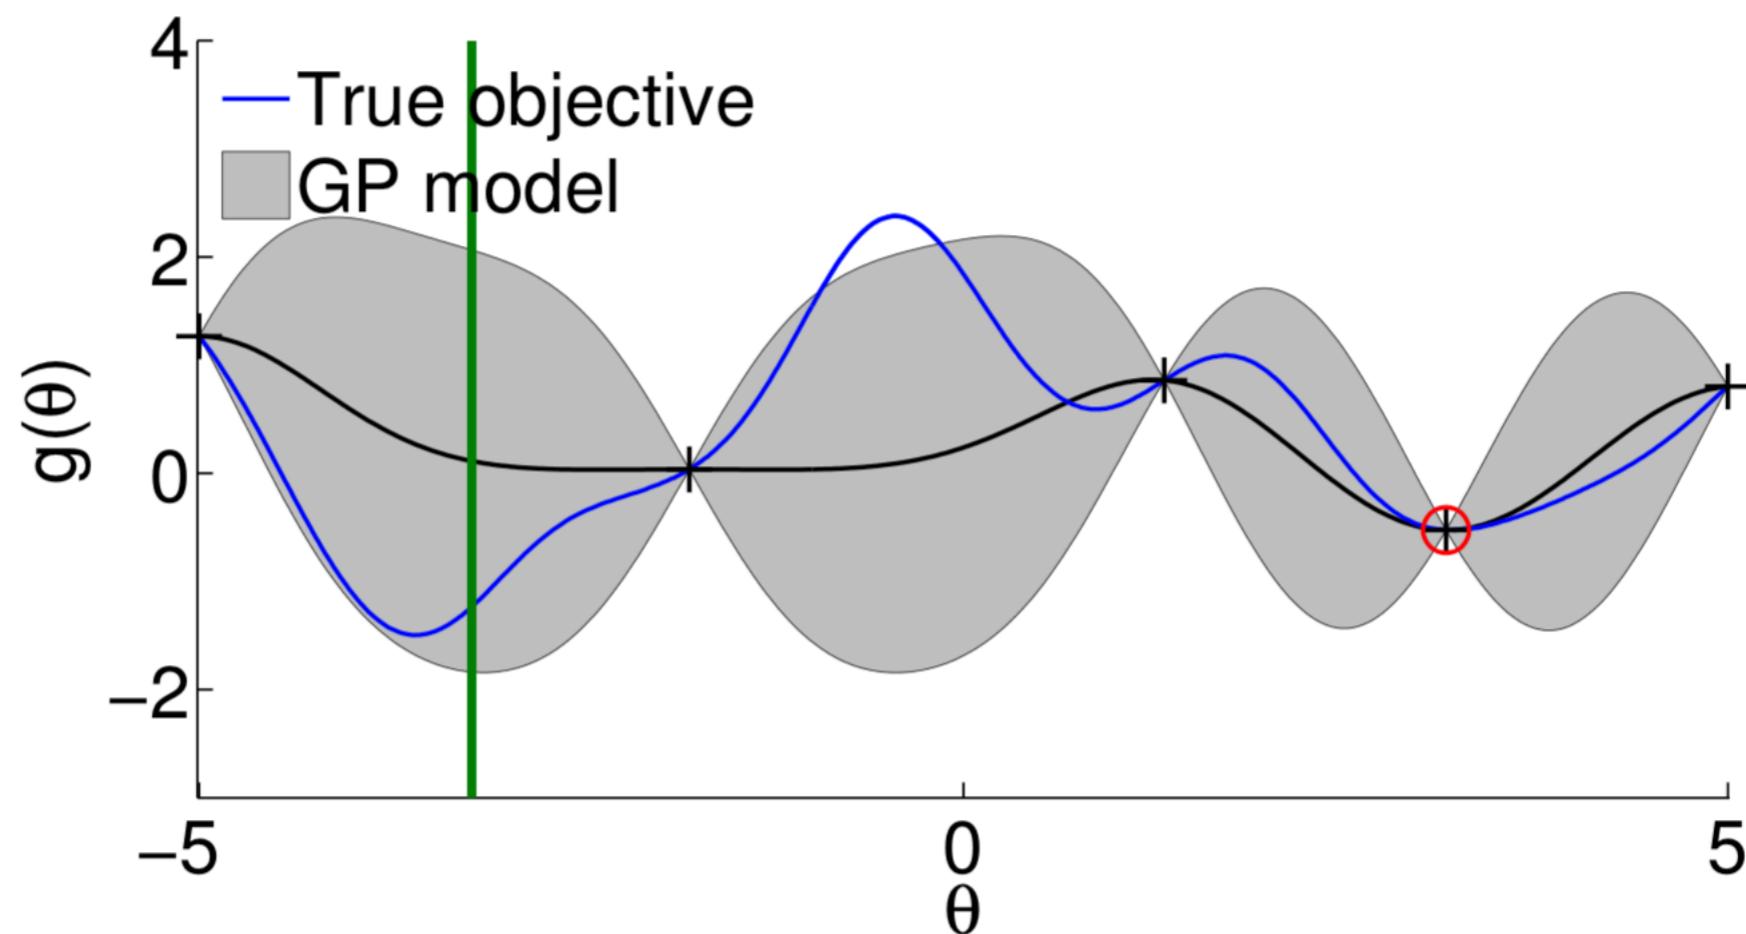
# Bayesian Optimization: Illustration



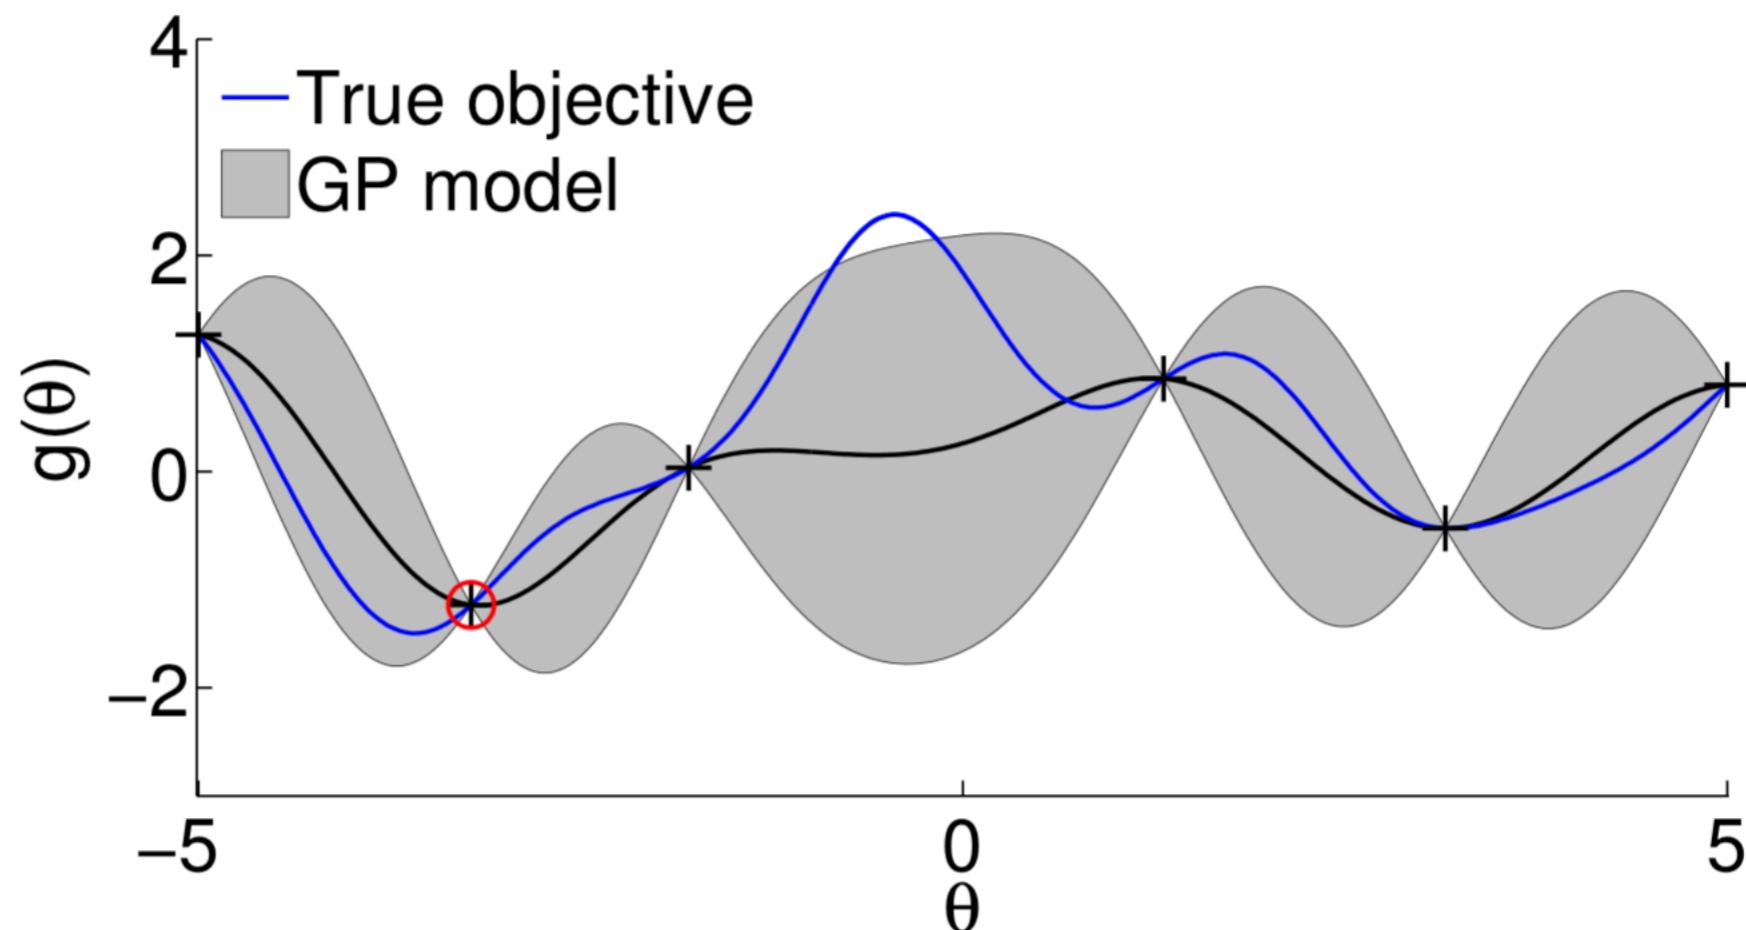
# Bayesian Optimization: Illustration



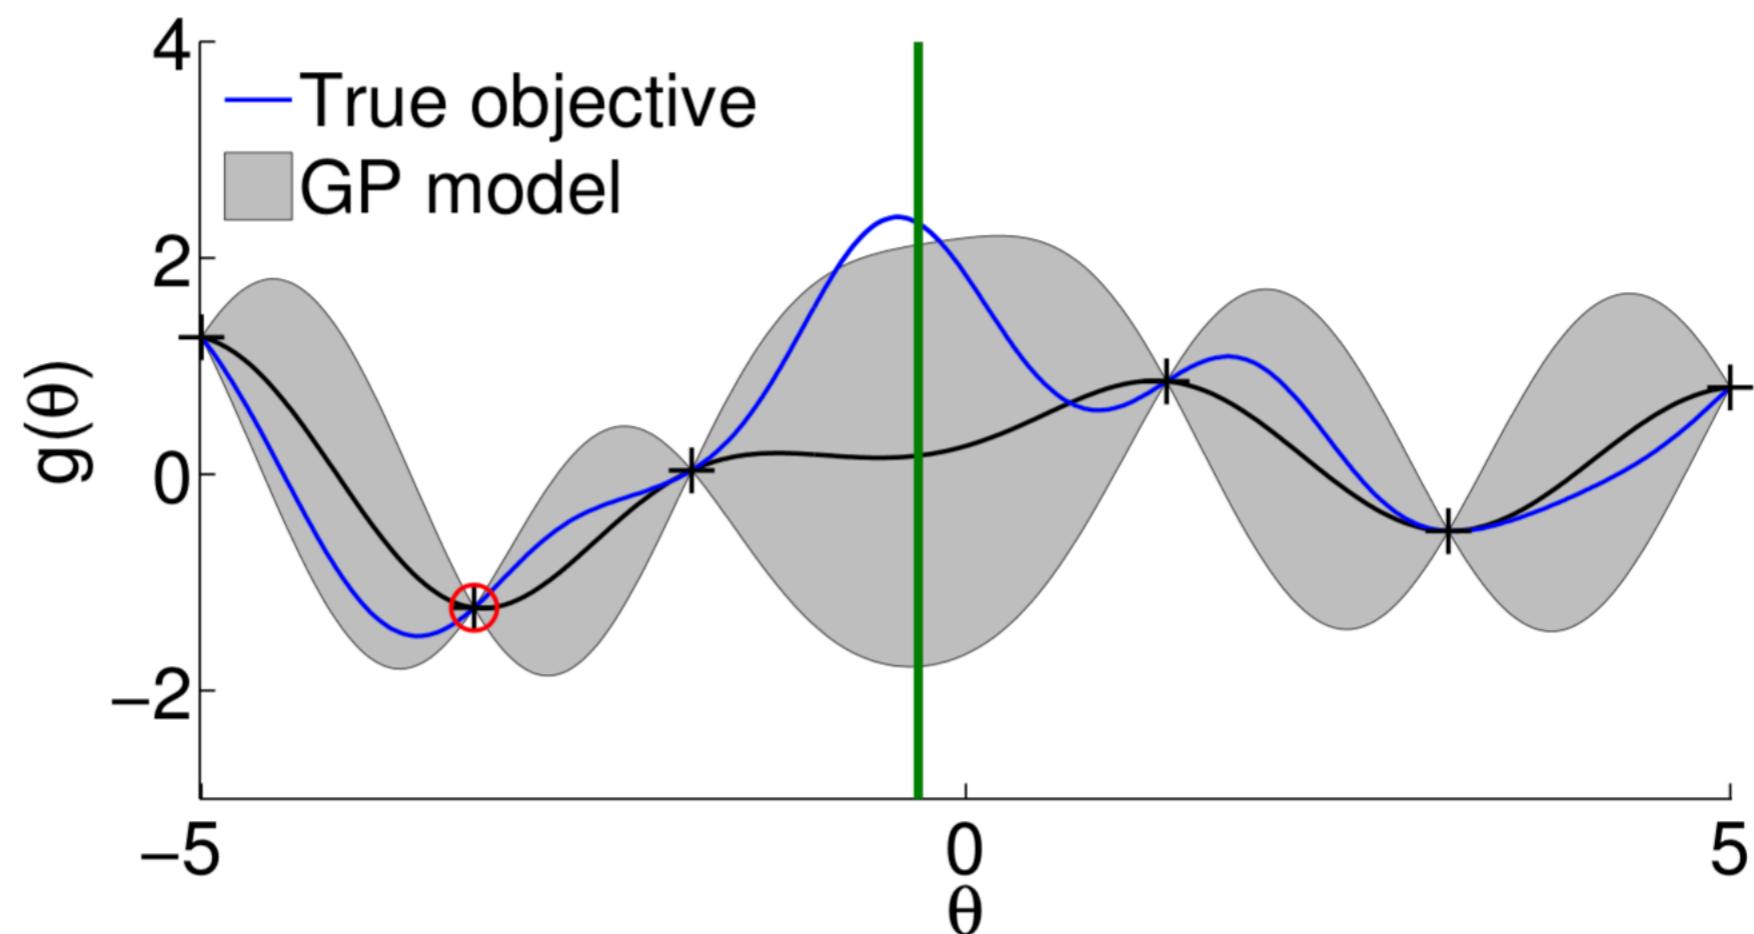
# Bayesian Optimization: Illustration



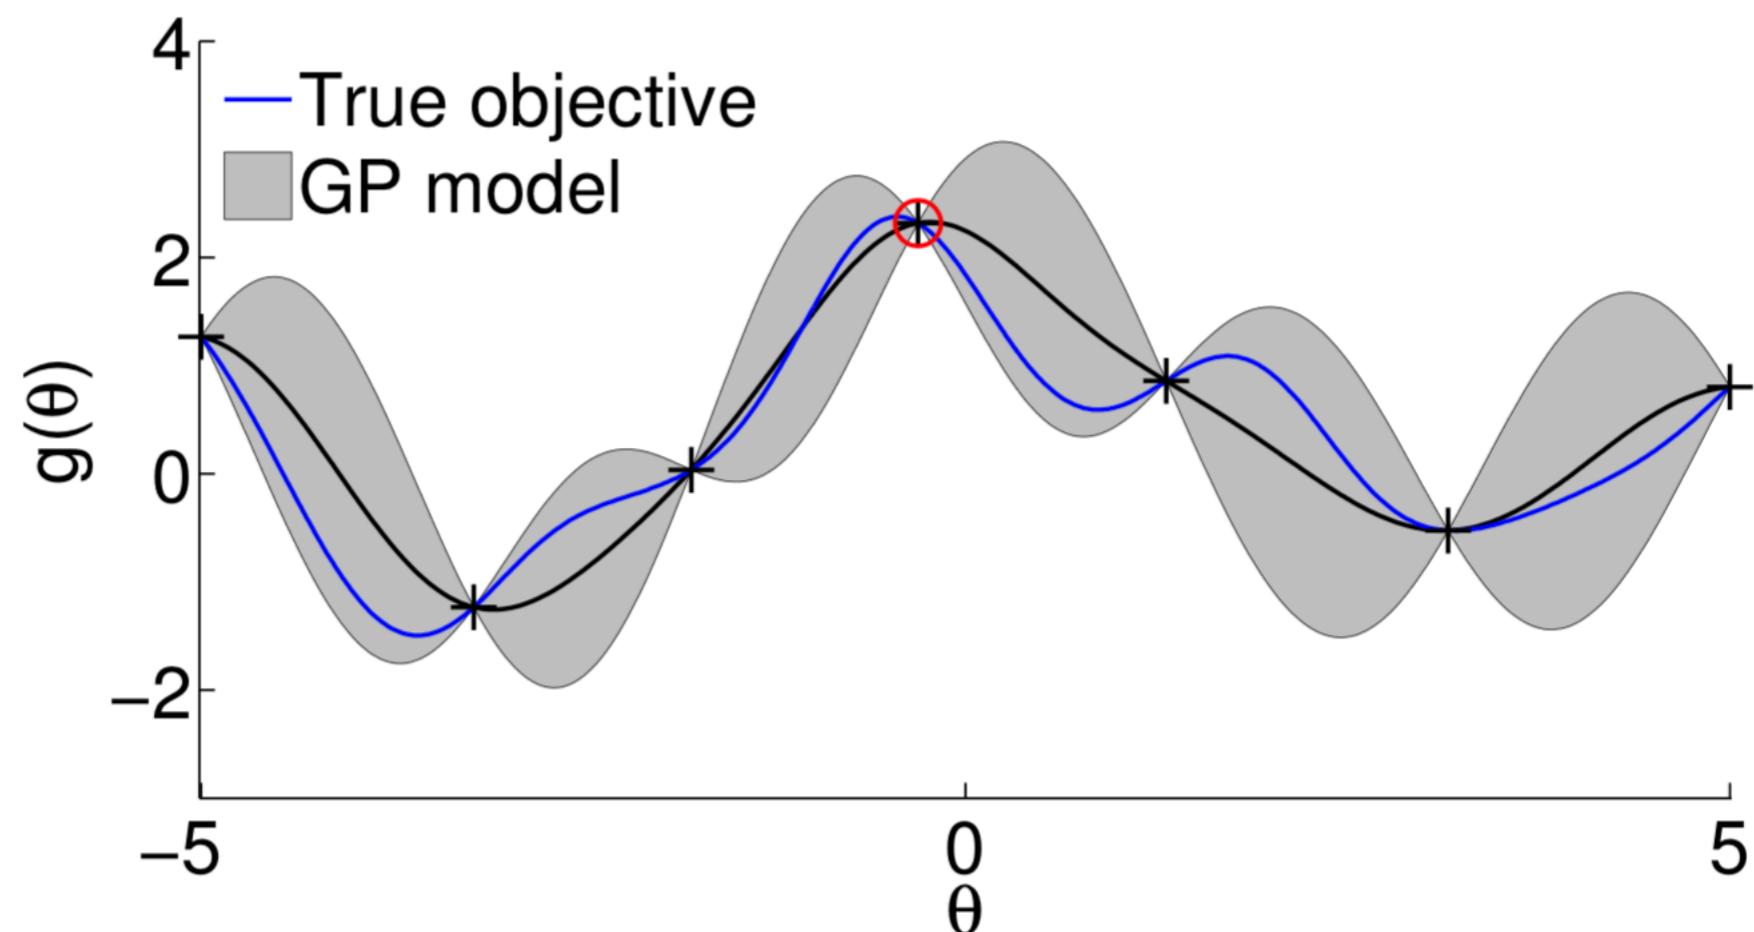
# Bayesian Optimization: Illustration



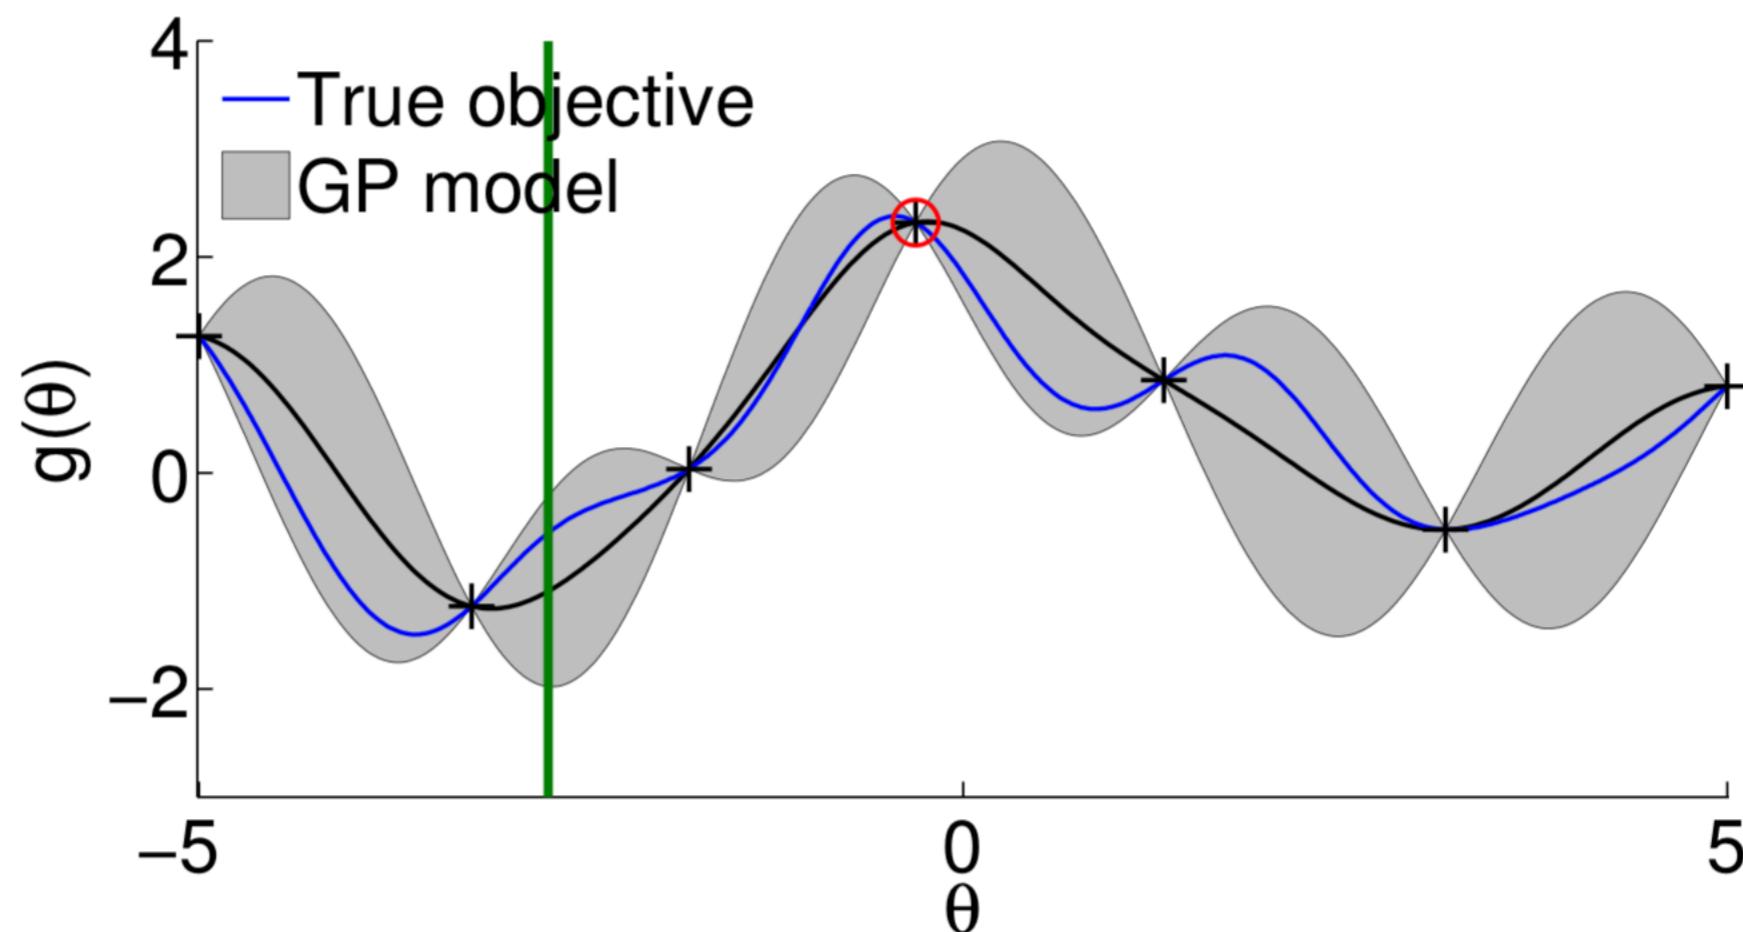
# Bayesian Optimization: Illustration



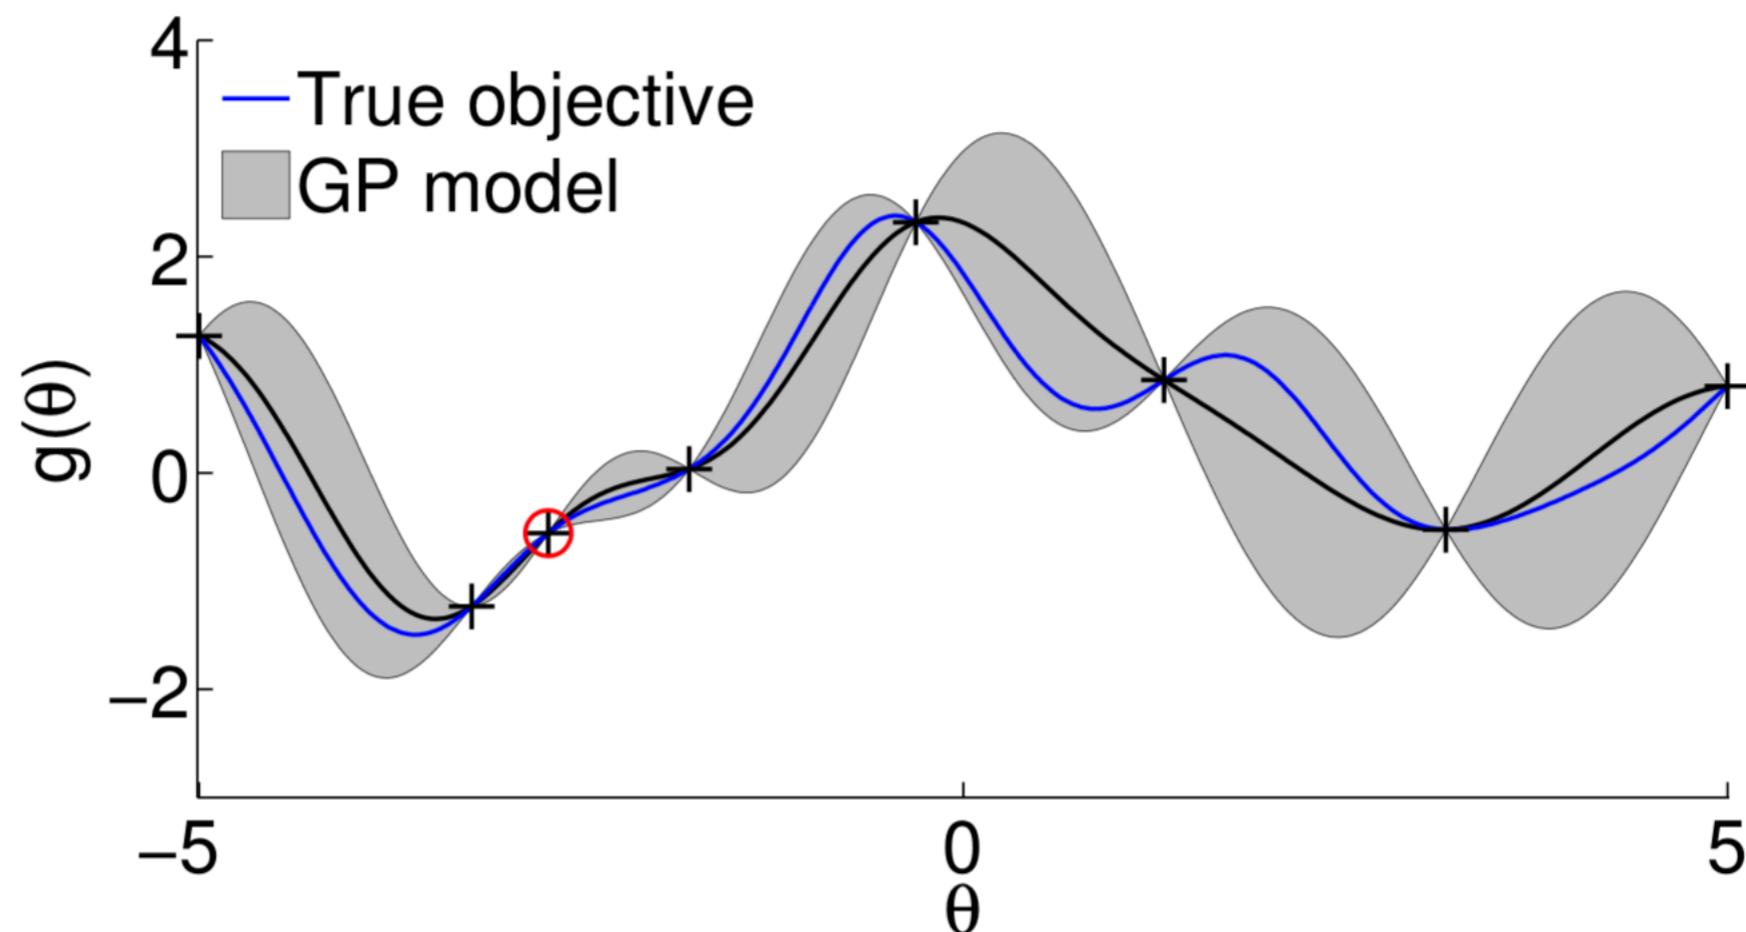
# Bayesian Optimization: Illustration



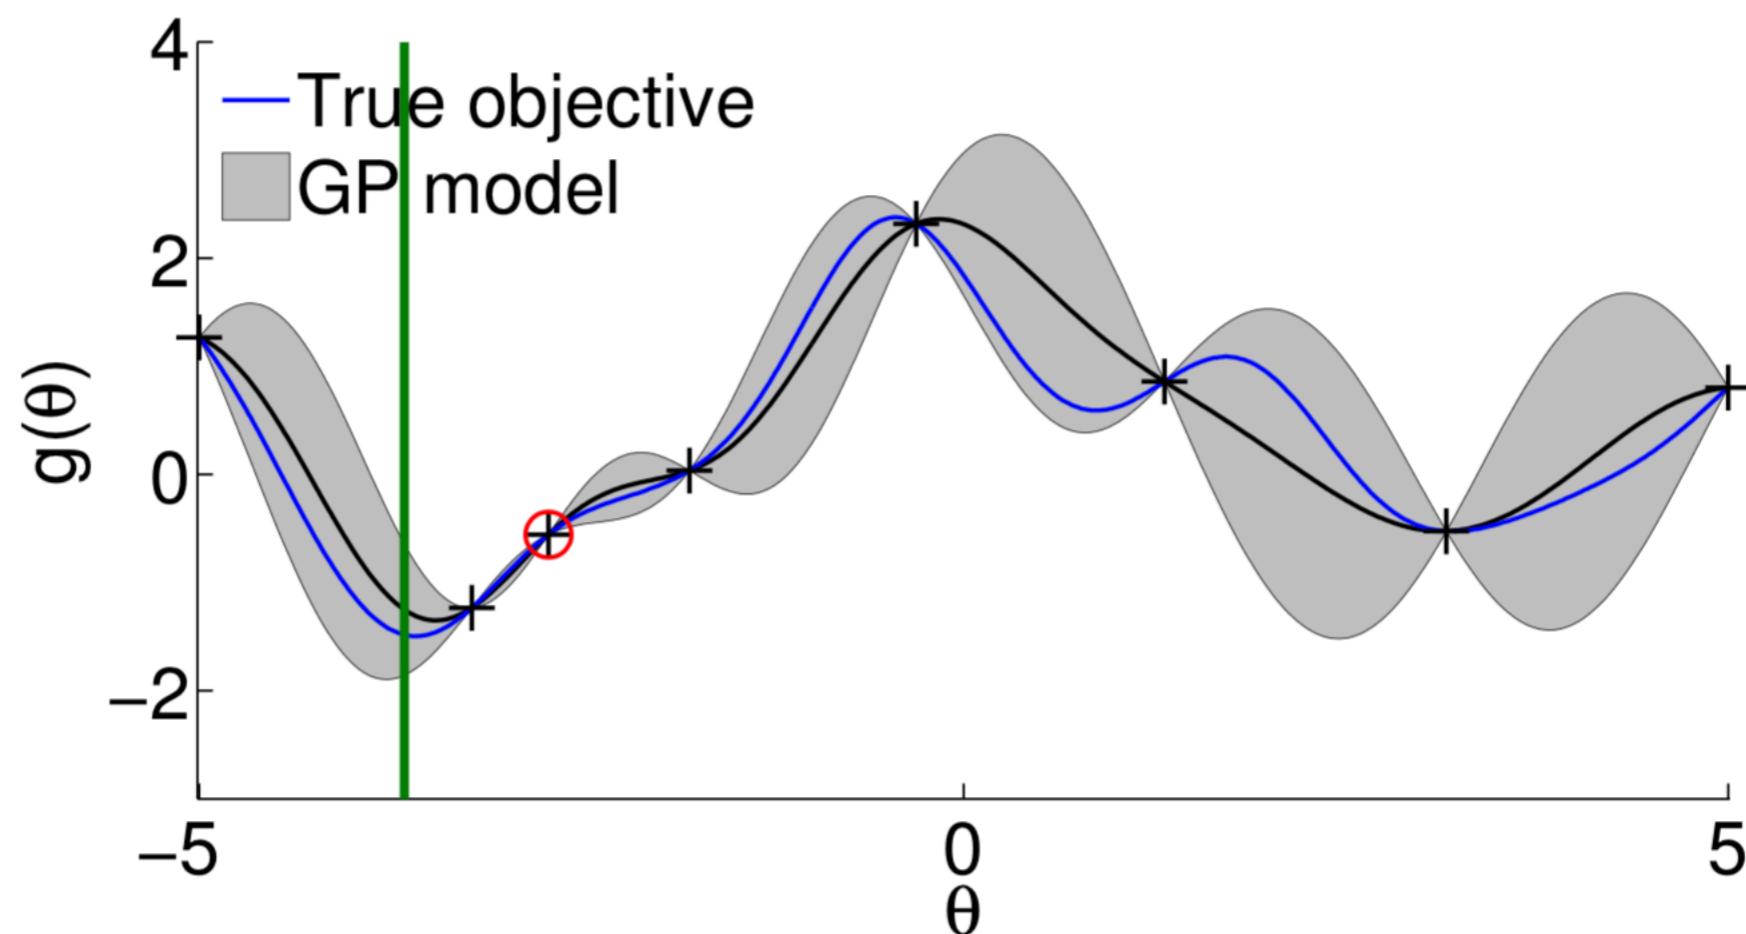
# Bayesian Optimization: Illustration



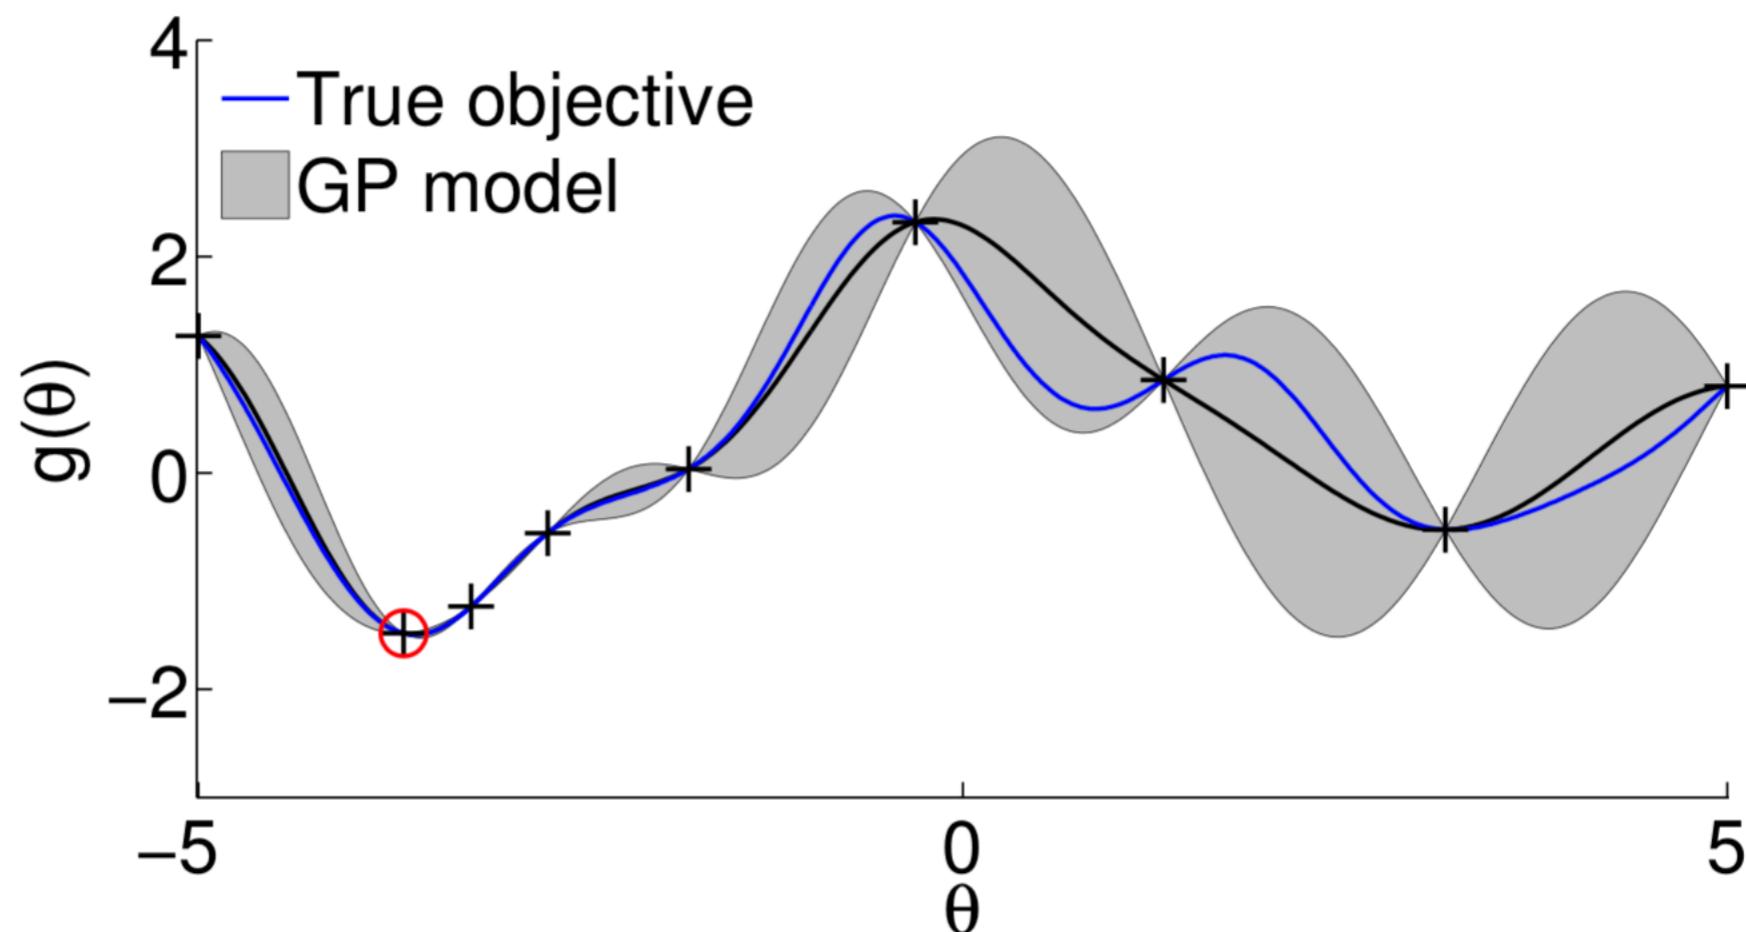
# Bayesian Optimization: Illustration



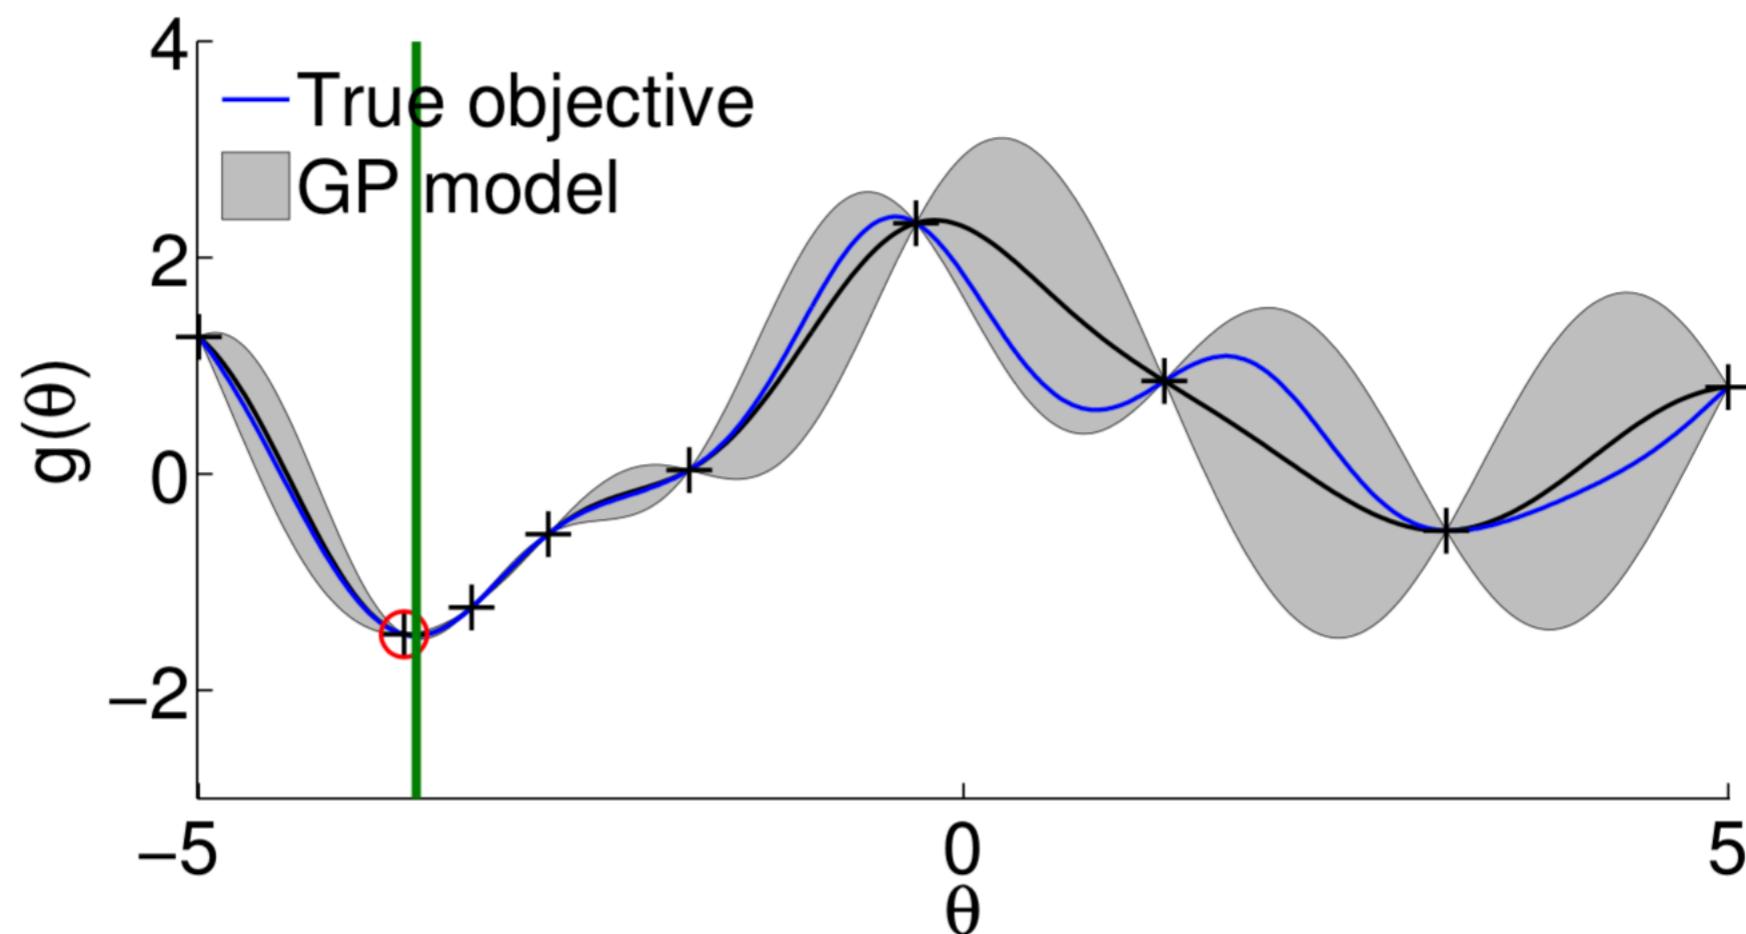
# Bayesian Optimization: Illustration



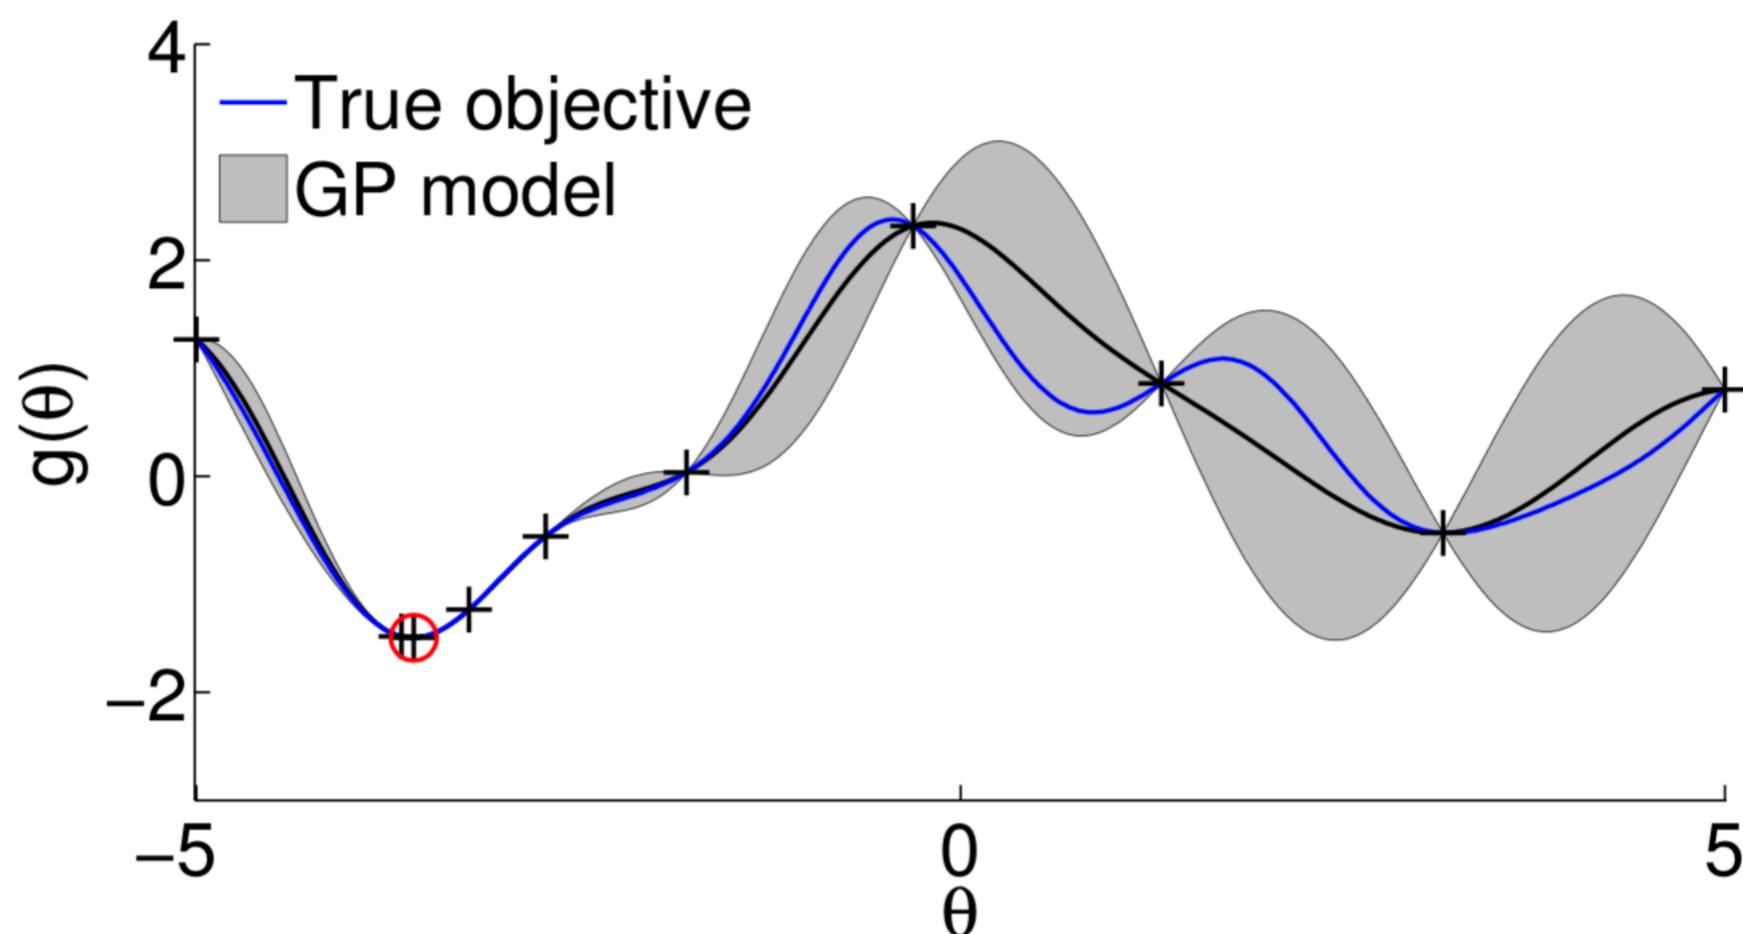
# Bayesian Optimization: Illustration



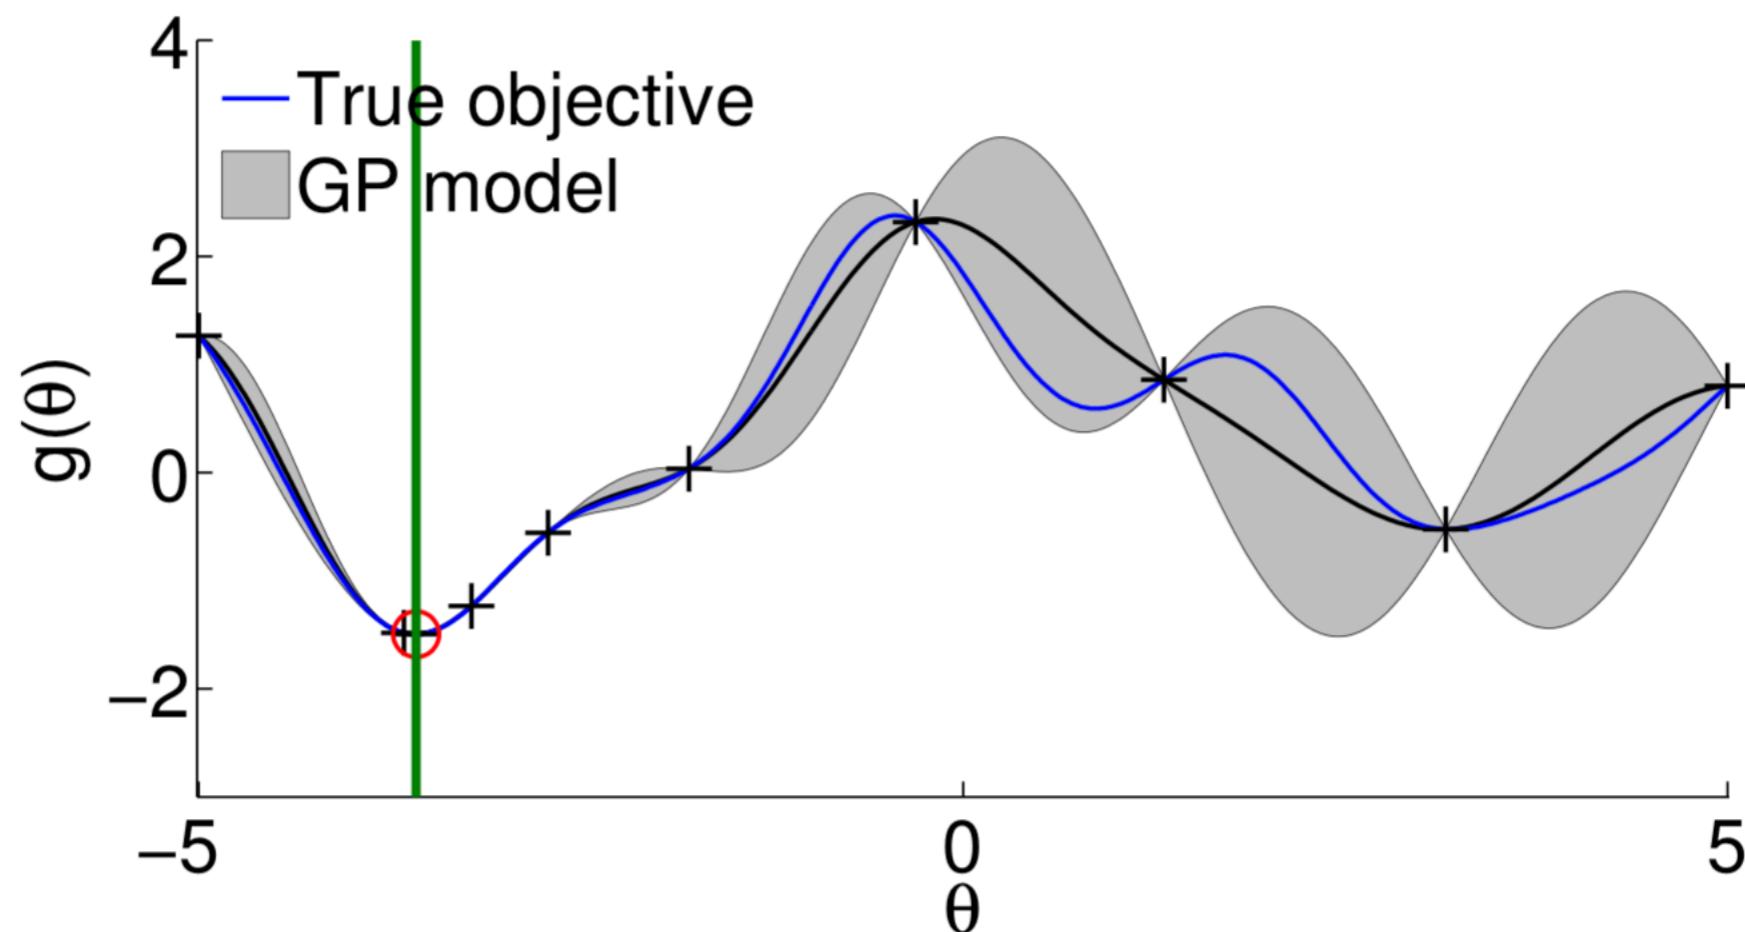
# Bayesian Optimization: Illustration



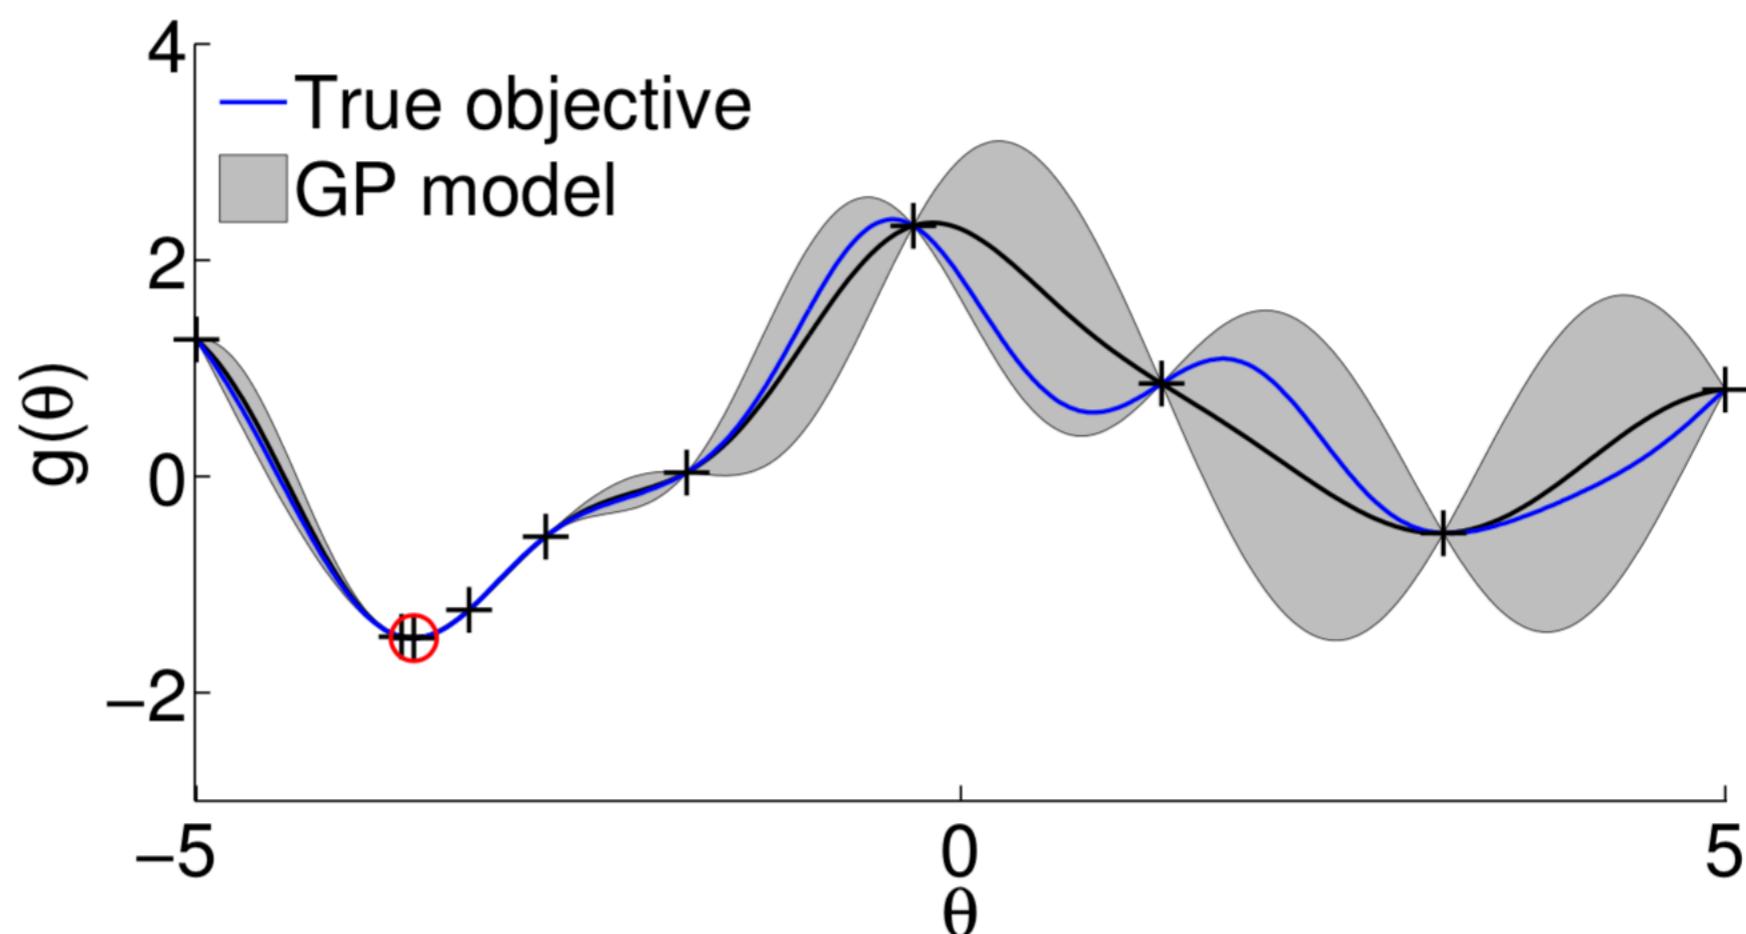
# Bayesian Optimization: Illustration



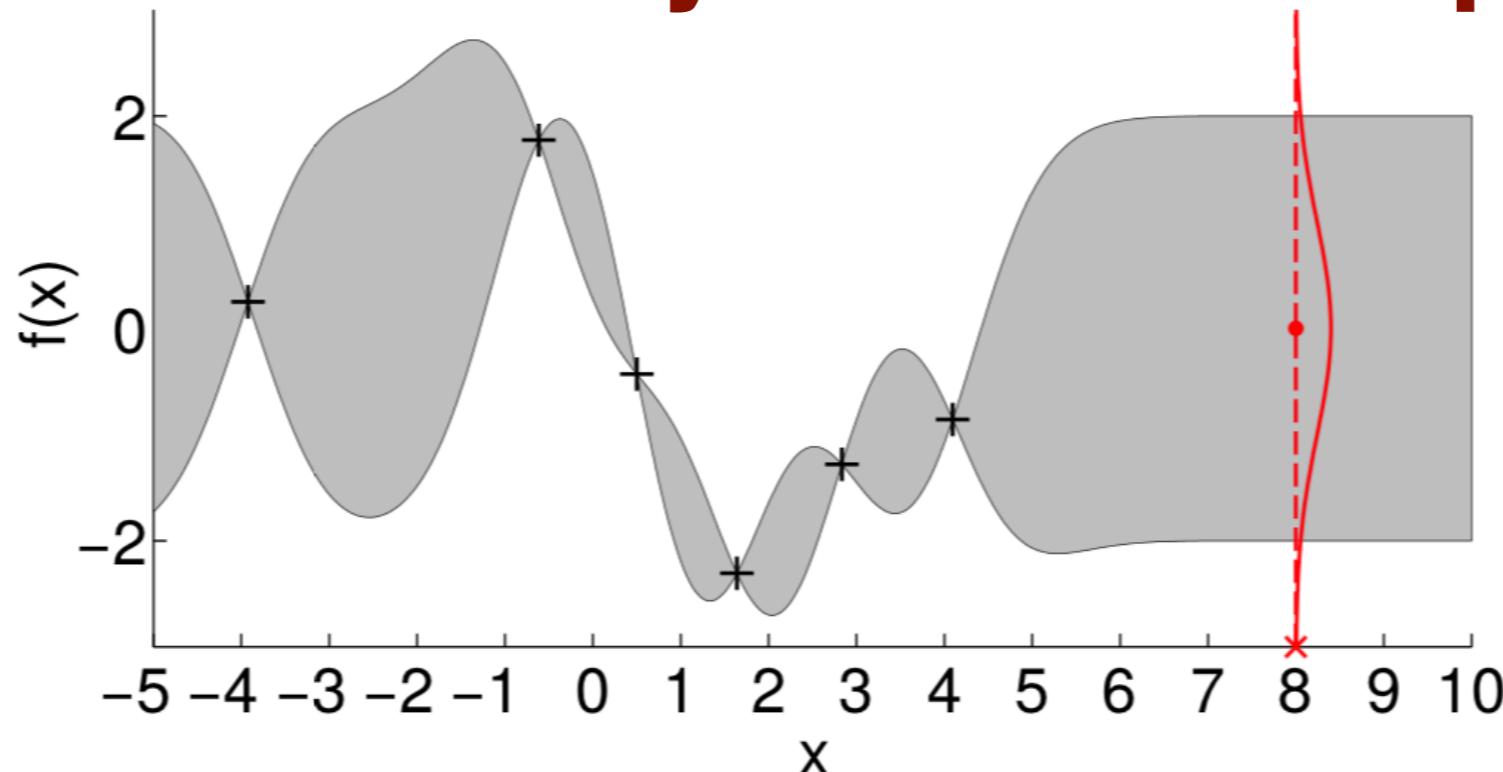
# Bayesian Optimization: Illustration



# Bayesian Optimization: Illustration

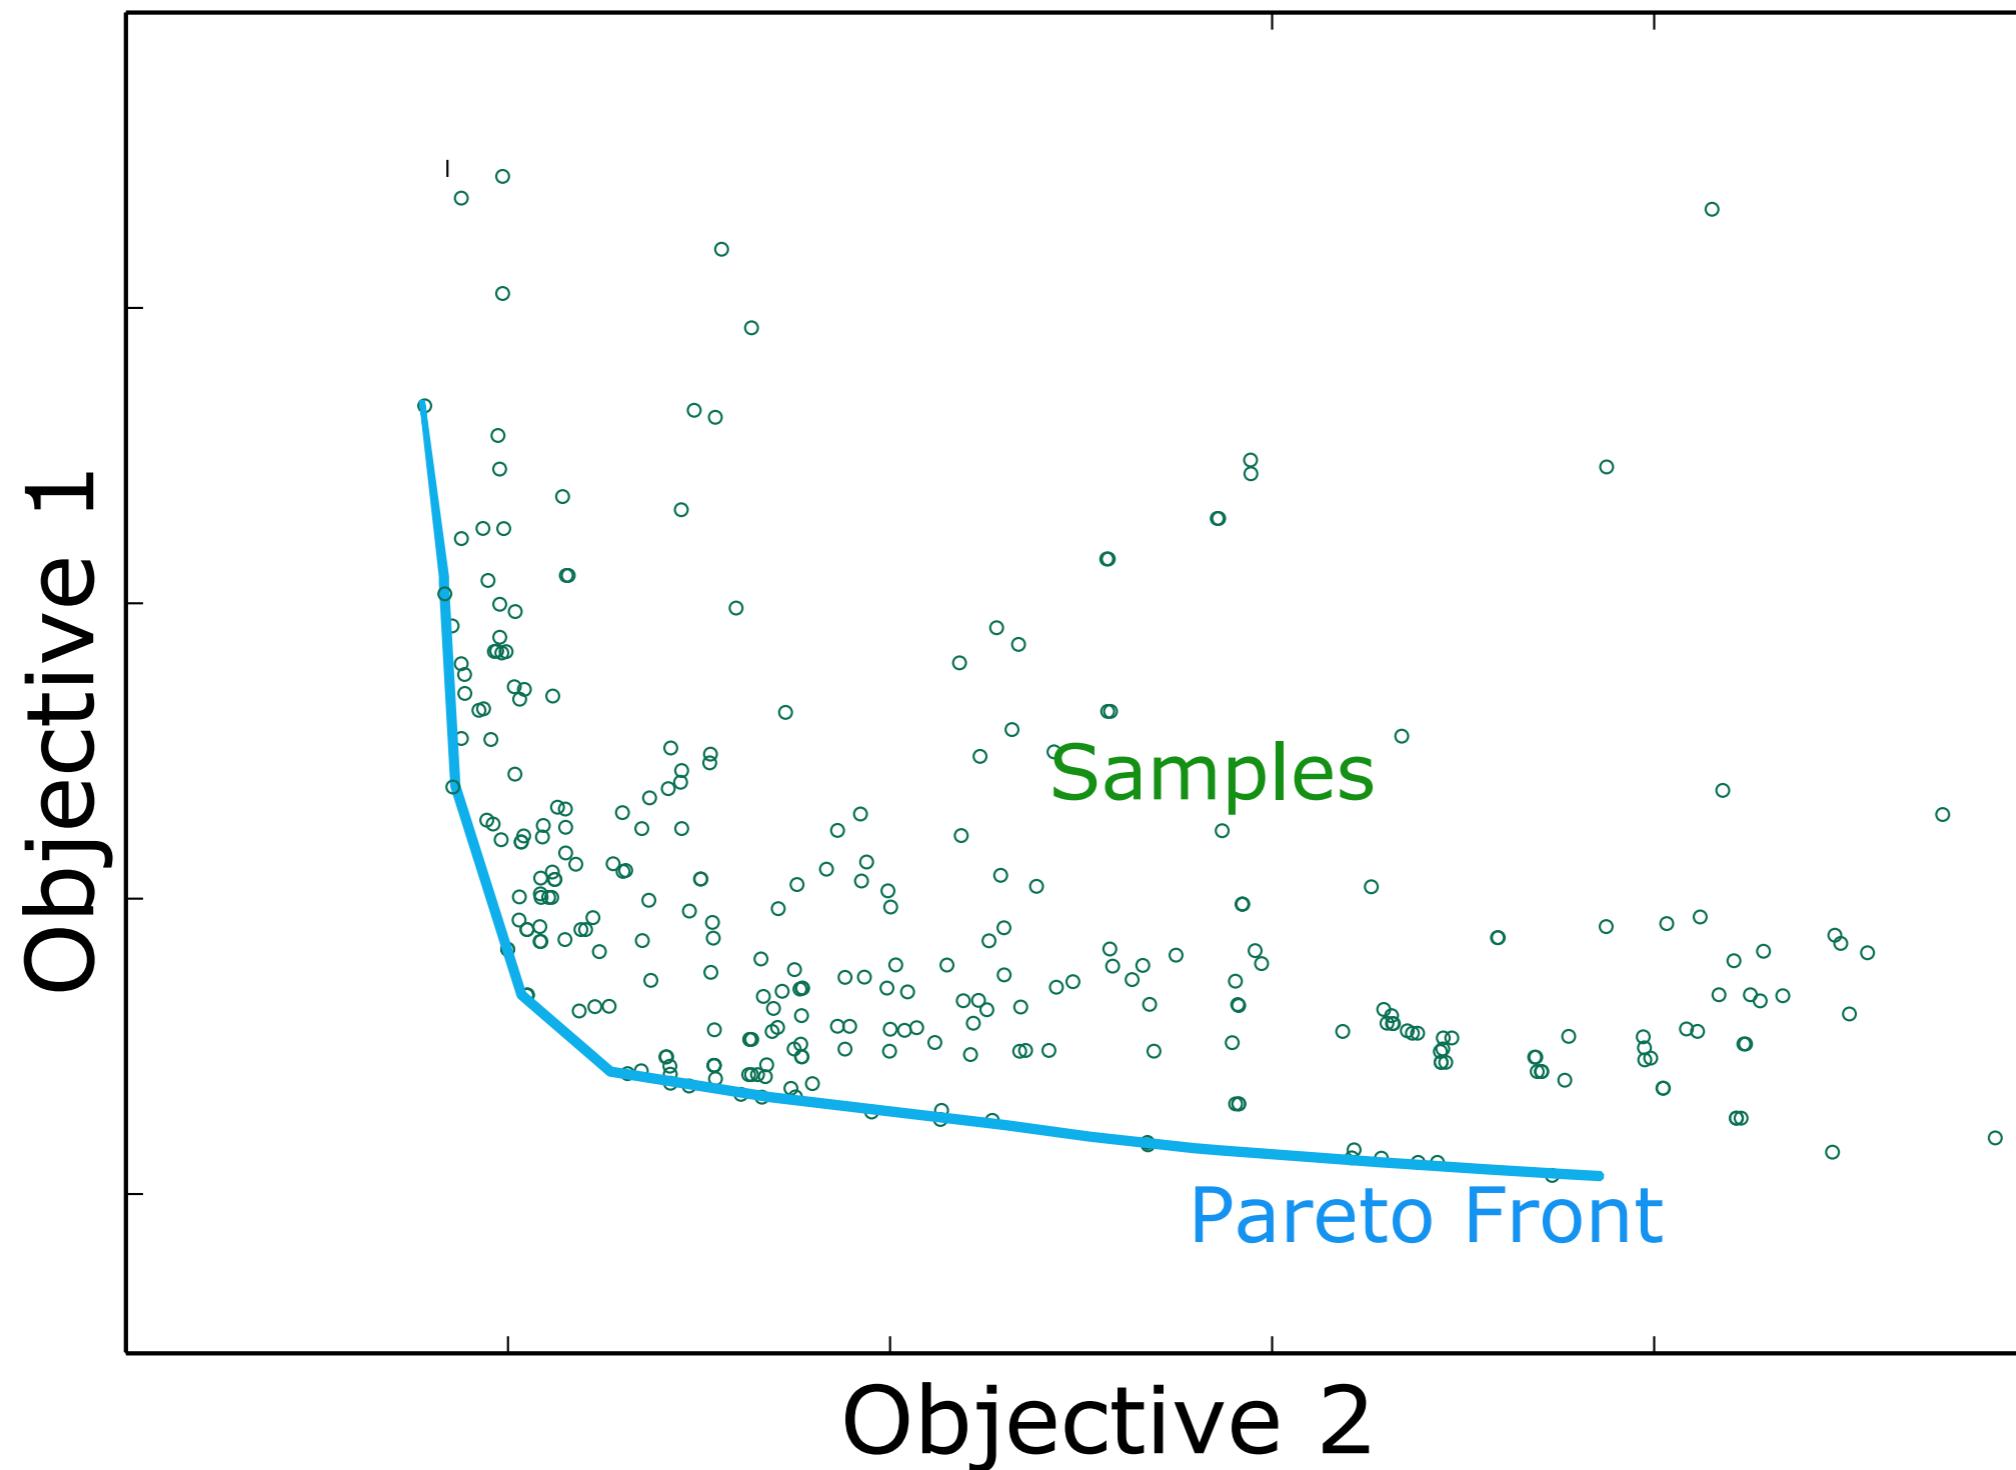


# Using Uncertainty in Global Optimization



- Goal: choosing the next point to evaluate the true objective
- Find a good (global) optimum
  - Need to get out of local optima
  - Probabilistic model gives us means and variances
- Trade off exploration and exploitation
  - **Exploration**: seek places with high variance
  - **Exploitation**: seek places with low mean
- **Acquisition function** combines the two

# Multi-Objective Goal: Pareto Front



# Multi-objective Optimization

**Random scalarizations** [Paria et al.] adapted for:

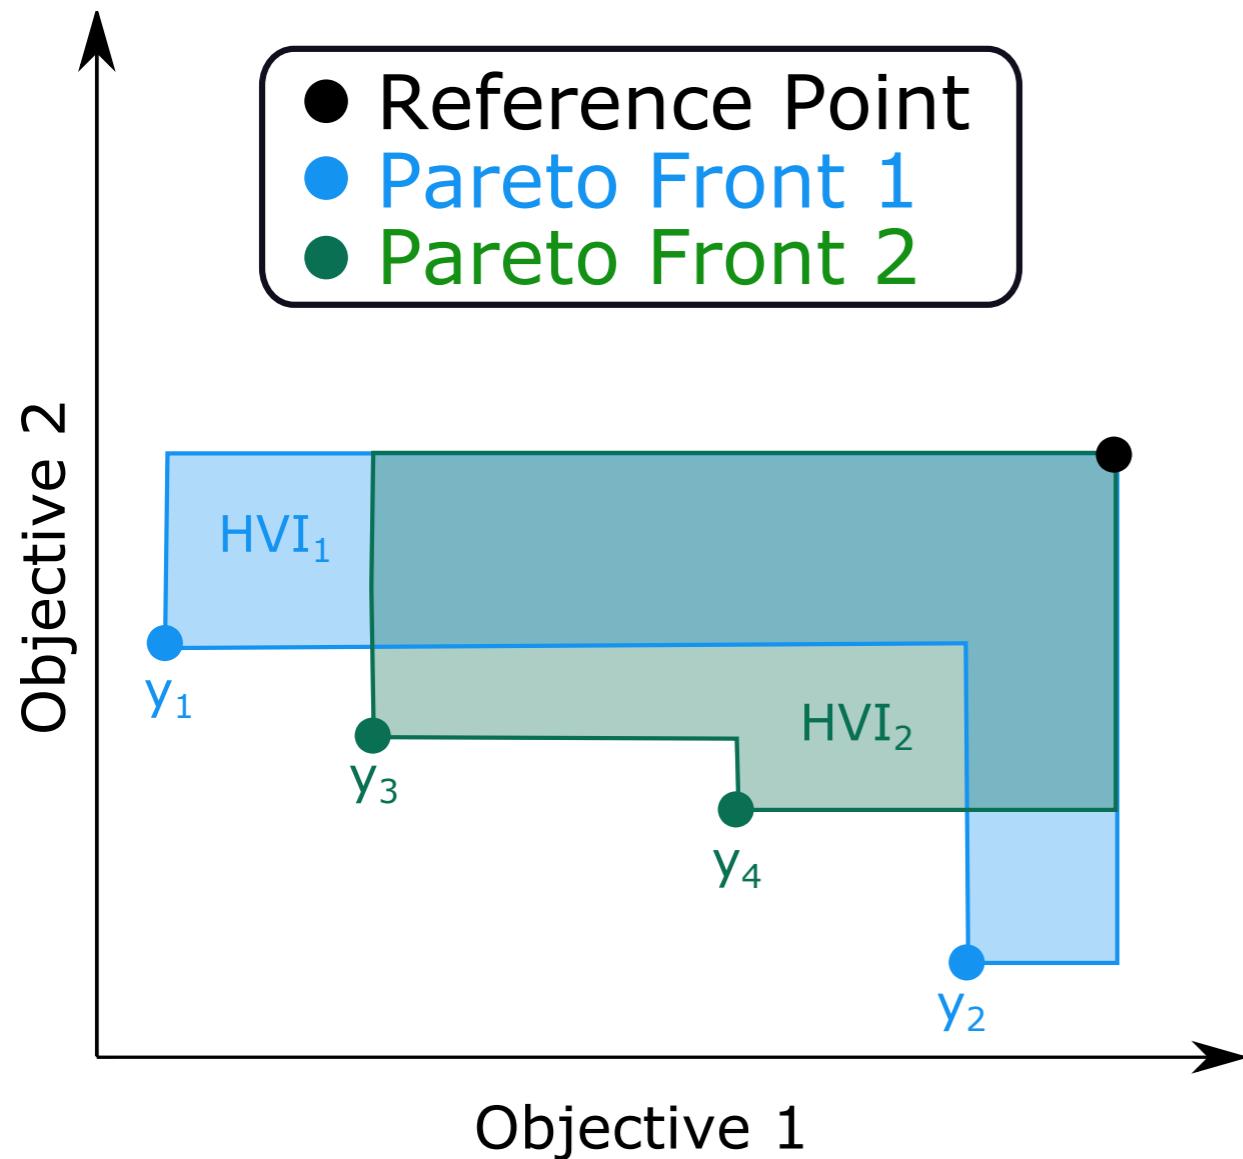
- RIOC variables
- Acquisition functions

Given K objectives  $f_1(\mathbf{x}), \dots, f_K(\mathbf{x})$  and a weight distribution  $\mathcal{L}$  defined as the simplex  $\boldsymbol{\lambda} \in \mathbb{R}_+^K$ ,  $\|\boldsymbol{\lambda}\|_1 = 1$ , a random scalarization is defined as  $g(\boldsymbol{\lambda}, \mathbf{x})$  where  $\boldsymbol{\lambda} \sim \mathcal{L}$

HyperMapper provides two scalarization functions:  
linear and Tchebyshev

**Linear scalarization:**  $g_{lin}(\boldsymbol{\lambda}, \mathbf{x}) = \sum_{k=1}^K \lambda_k f_k(\mathbf{x})$

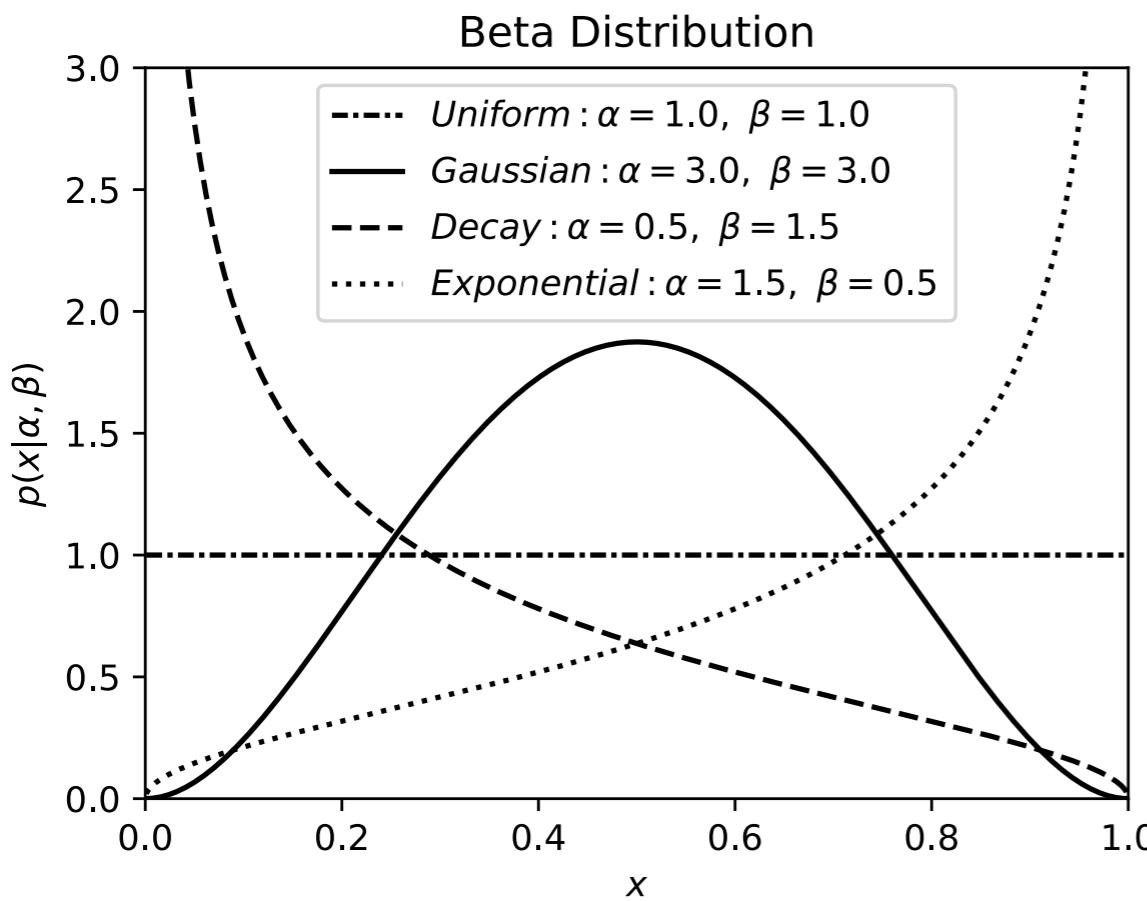
# HyperVolume Indicator (HVI)



- Used to compare 2 Paretos
- HVI is a hypervolume
  - with respect to a reference
- It is always a scalar
  - even with multiple objectives
- The lower the better:  
**HVI<sub>total</sub> - HVI<sub>i</sub>**

# Injecting Prior Knowledge

- Need a probability distribution that:
  1. Has a finite domain
  2. Can flexibly model shapes including:  
Bell-shape and J-shape
- Beta distribution - parameters alpha and beta
- In addition, need of a discrete distribution for categorical variables



PDF given by:

$$f(x|\alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1}$$

for  $x \in [0, 1]$  and  $\alpha, \beta > 0$ ,

where  $\Gamma$  is the Gamma function

# Probabilistic Model in HyperMapper

- Predicts an outcome given an input  $x$

# Probabilistic Model in HyperMapper

- Predicts an outcome given an input  $x$
- One surrogate model per objective
  - Random forests (RF) for regression

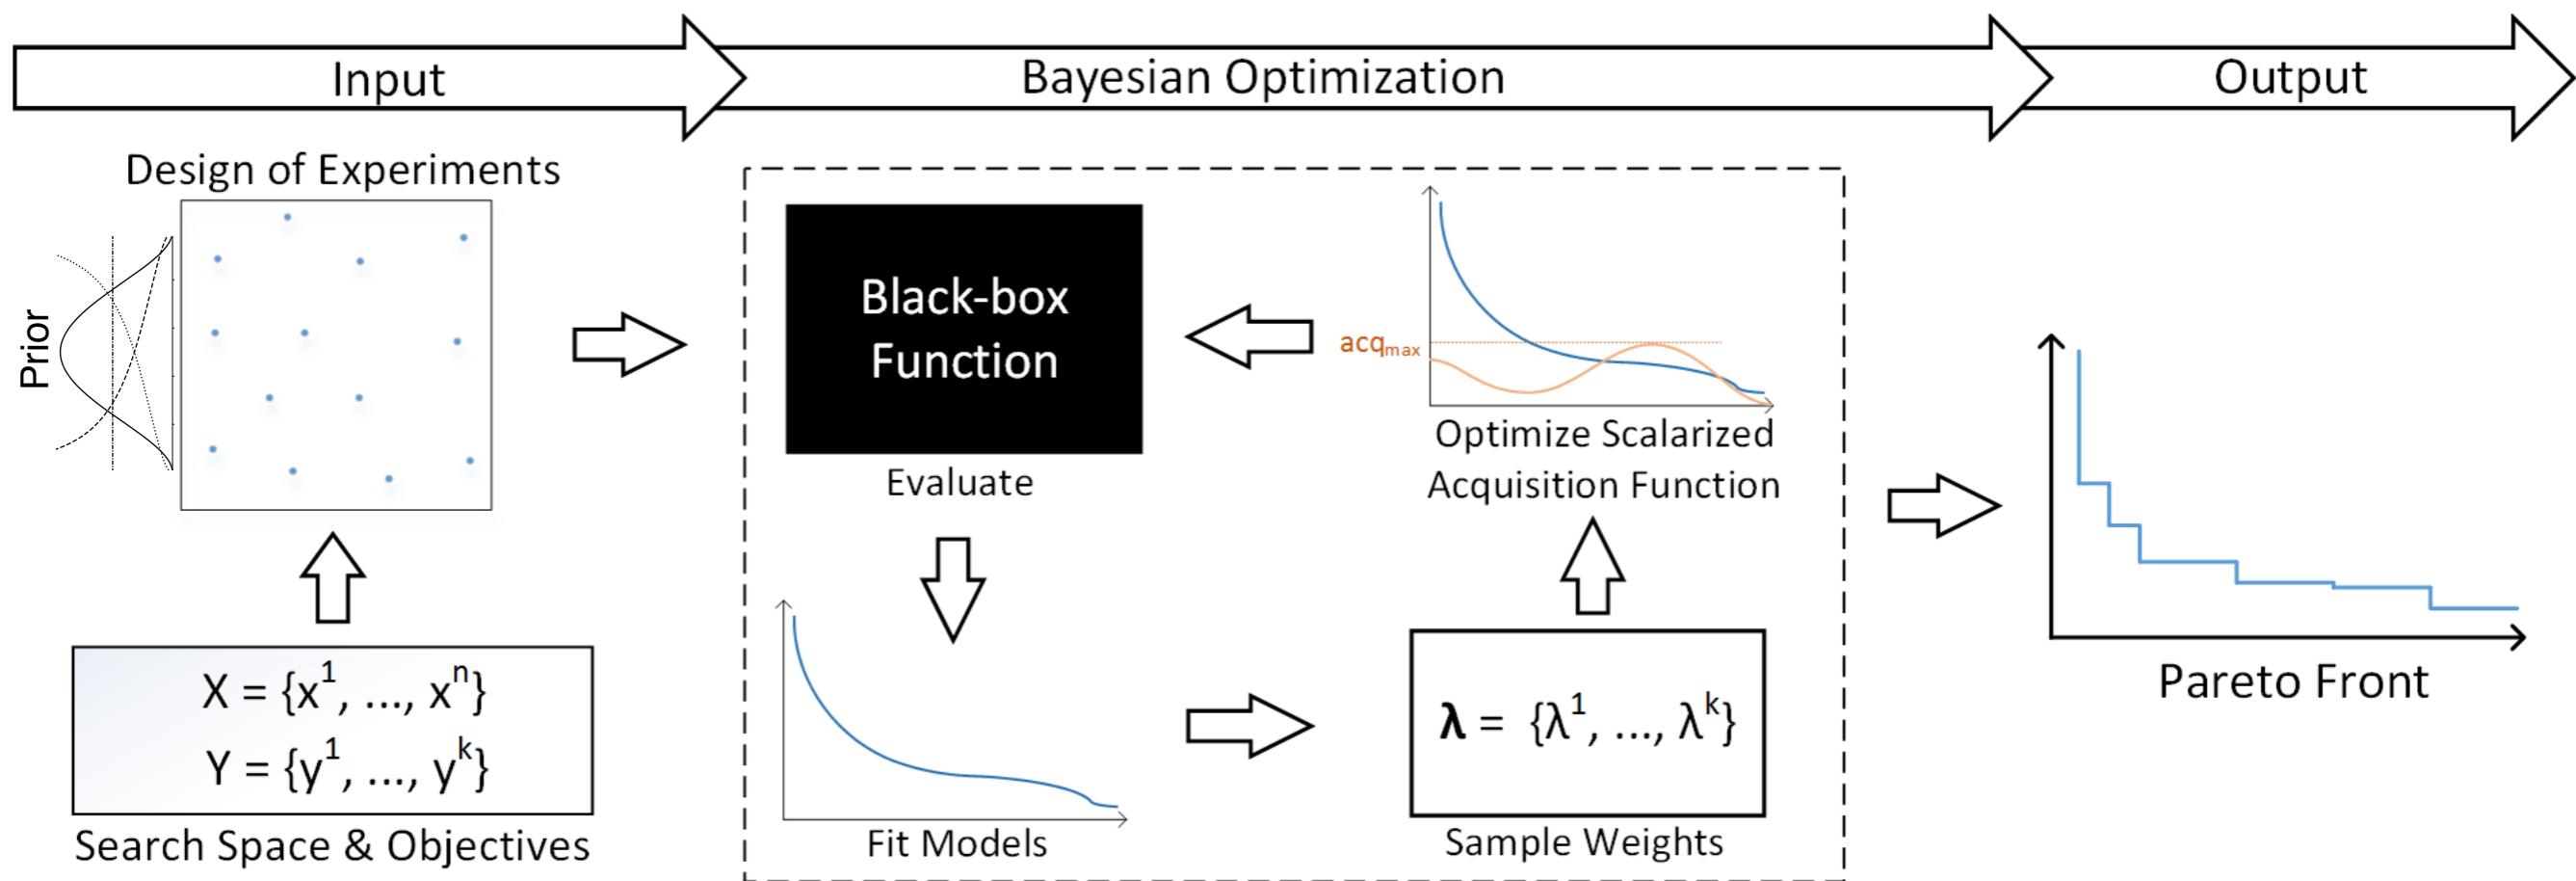
# Probabilistic Model in HyperMapper

- Predicts an outcome given an input  $x$
- One surrogate model per objective
  - Random forests (RF) for regression
- Intuition of why RF is a good model:
  - Good at non-linearity, multi-modality and non-smoothness

# Probabilistic Model in HyperMapper

- Predicts an outcome given an input  $x$
- One surrogate model per objective
  - Random forests (RF) for regression
- Intuition of why RF is a good model:
  - Good at non-linearity, multi-modality and non-smoothness
- The model has not to be perfect
  - Interested in optima not the best model
  - Coefficient of determination ( $R^2$ ) is the wrong indicator of success

# The HyperMapper Framework



# Outline

## 1. Black-box Optimization

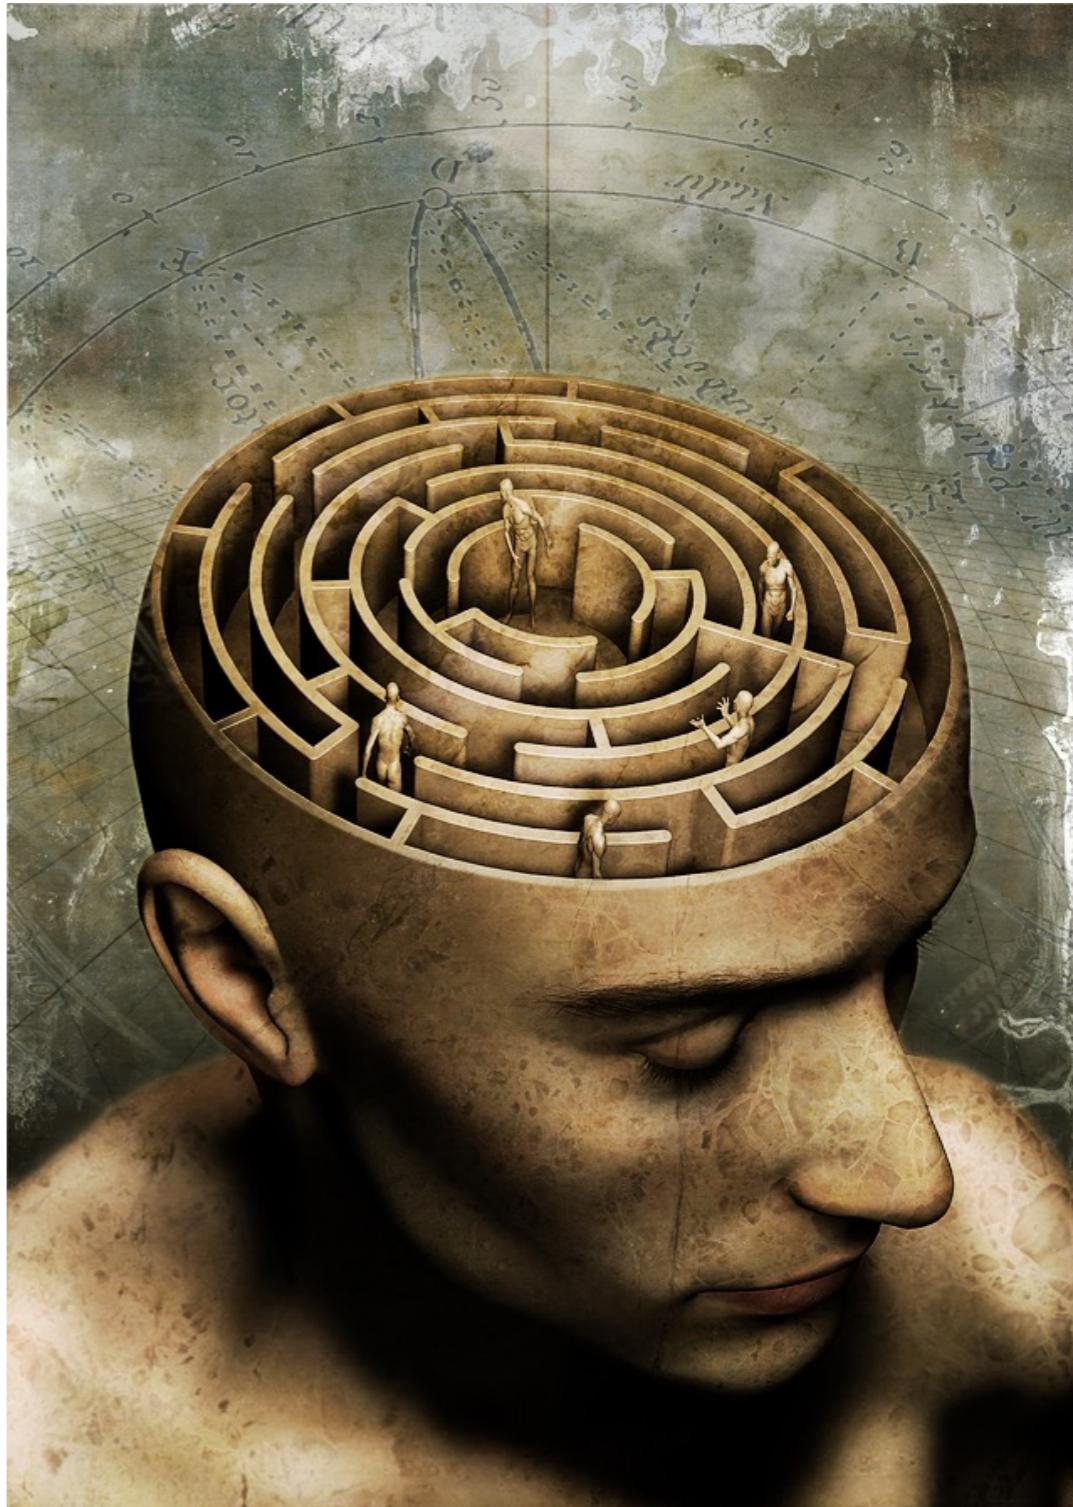
- Problem Setting
- Motivation Examples
- Taxonomy

## 2. The HyperMapper Framework

- Bayesian Optimization
- Pareto and Hypervolume Indicator
- Prior Distribution

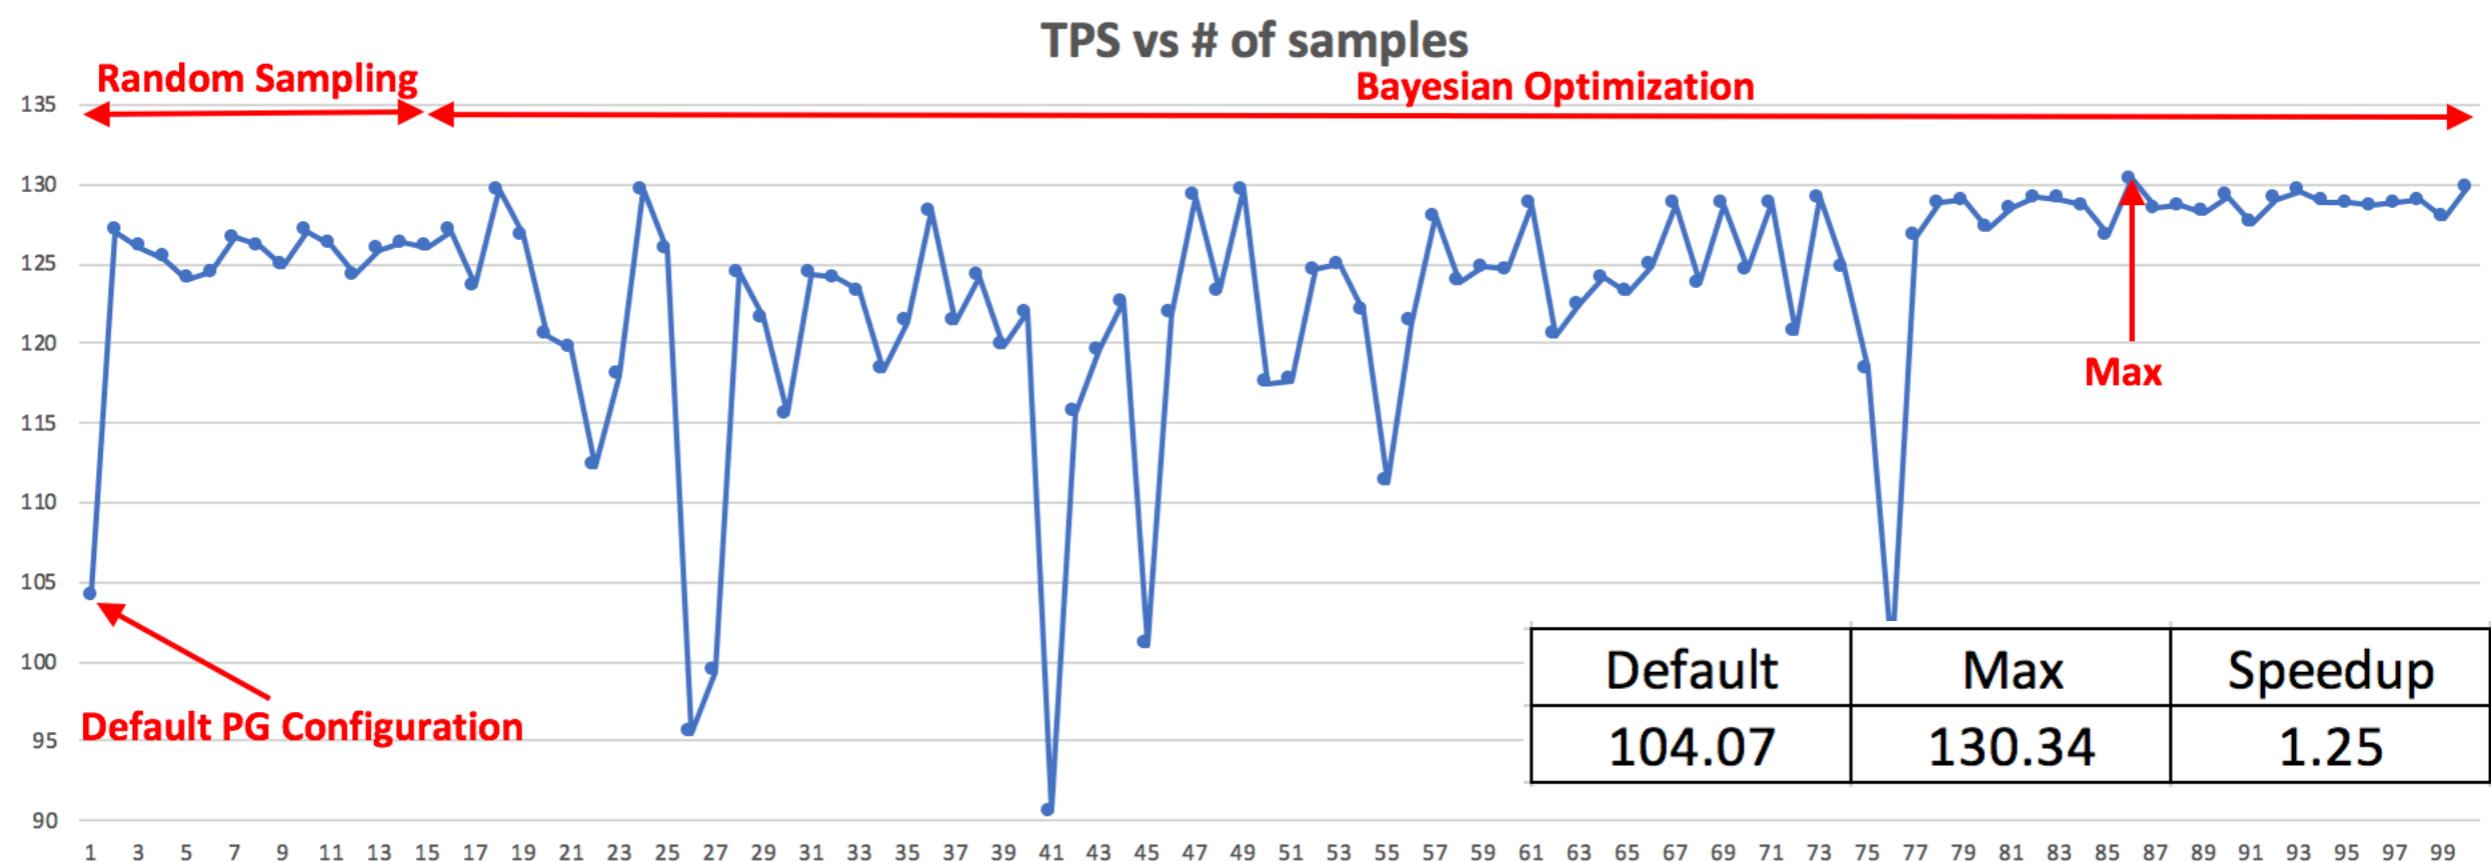
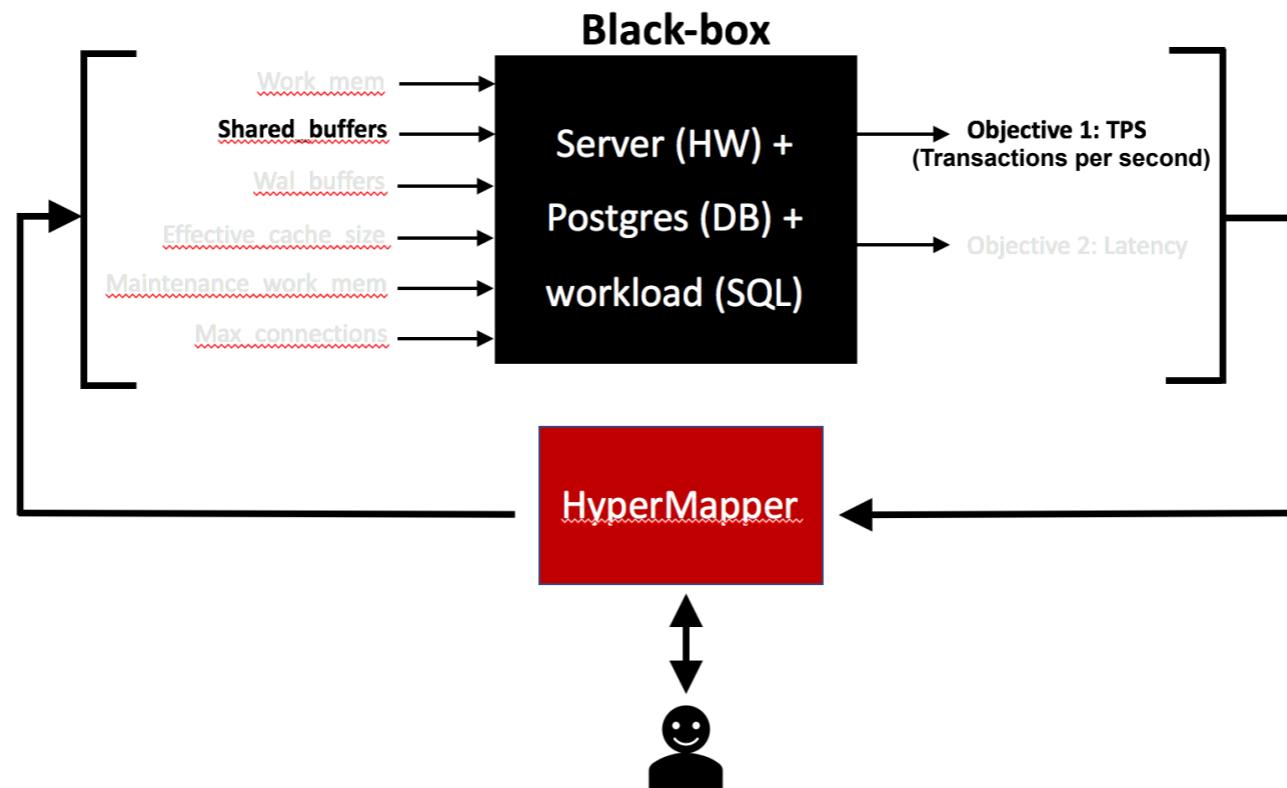
## 3. Experimental Results

- **Hardware Design**
- **Computer Vision**
- **Databases**



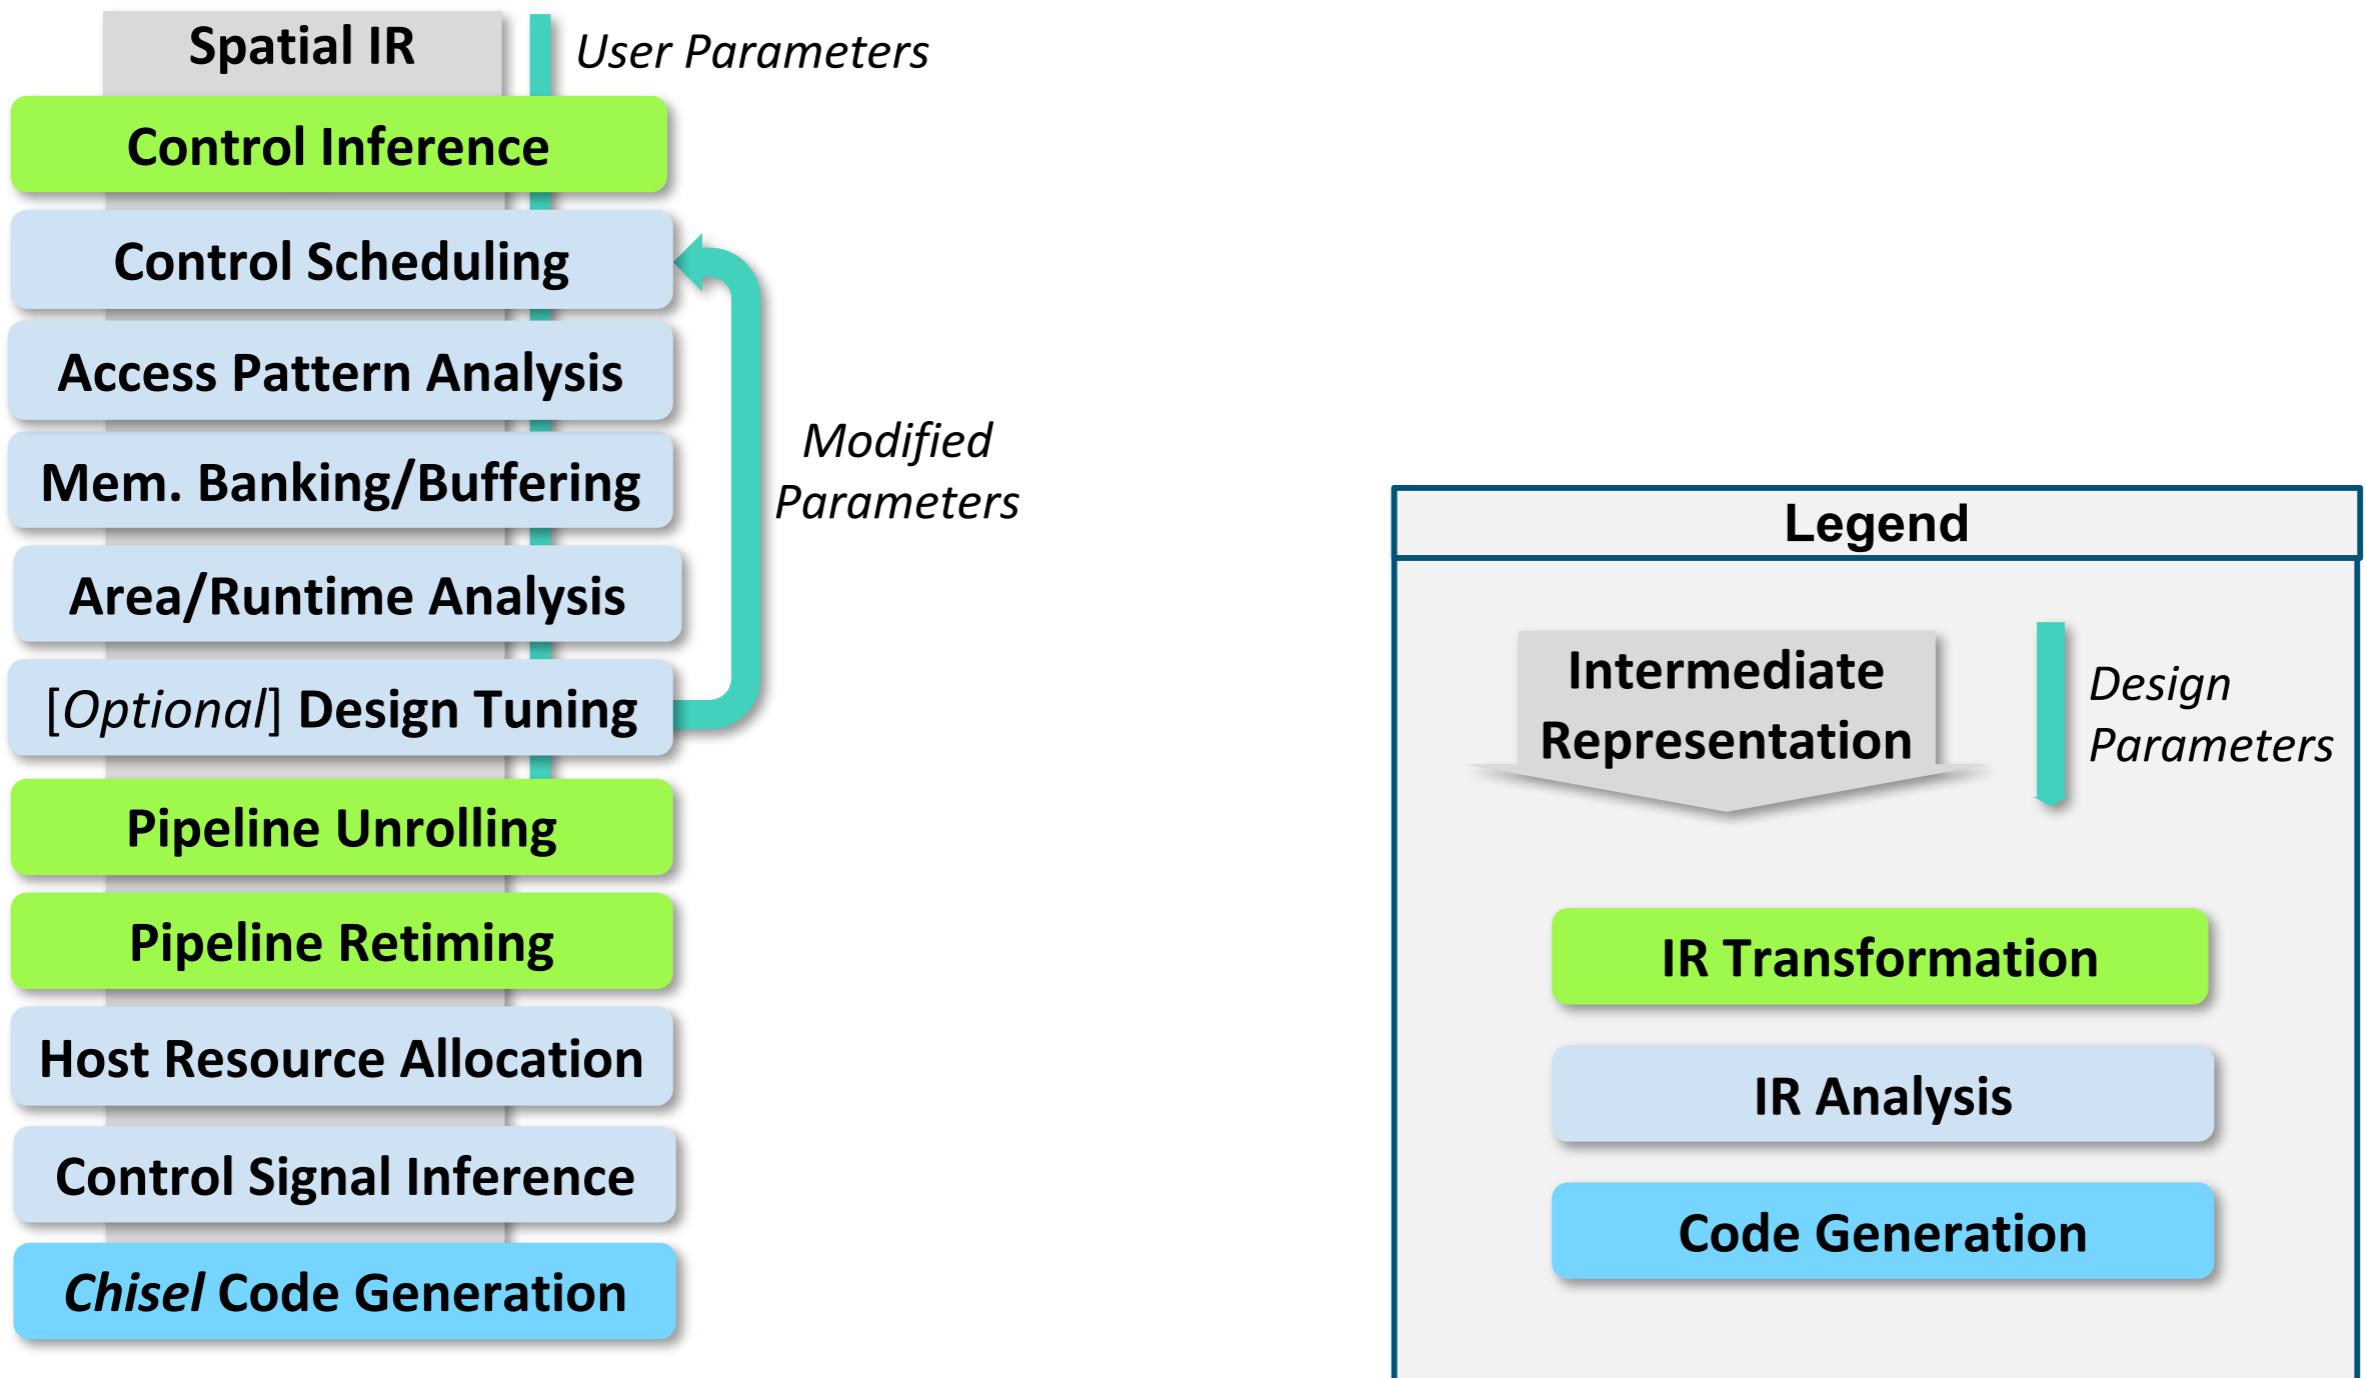
# Application 1 - Microsoft Collaboration

## Preliminary Results on a Simple Scenario



# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



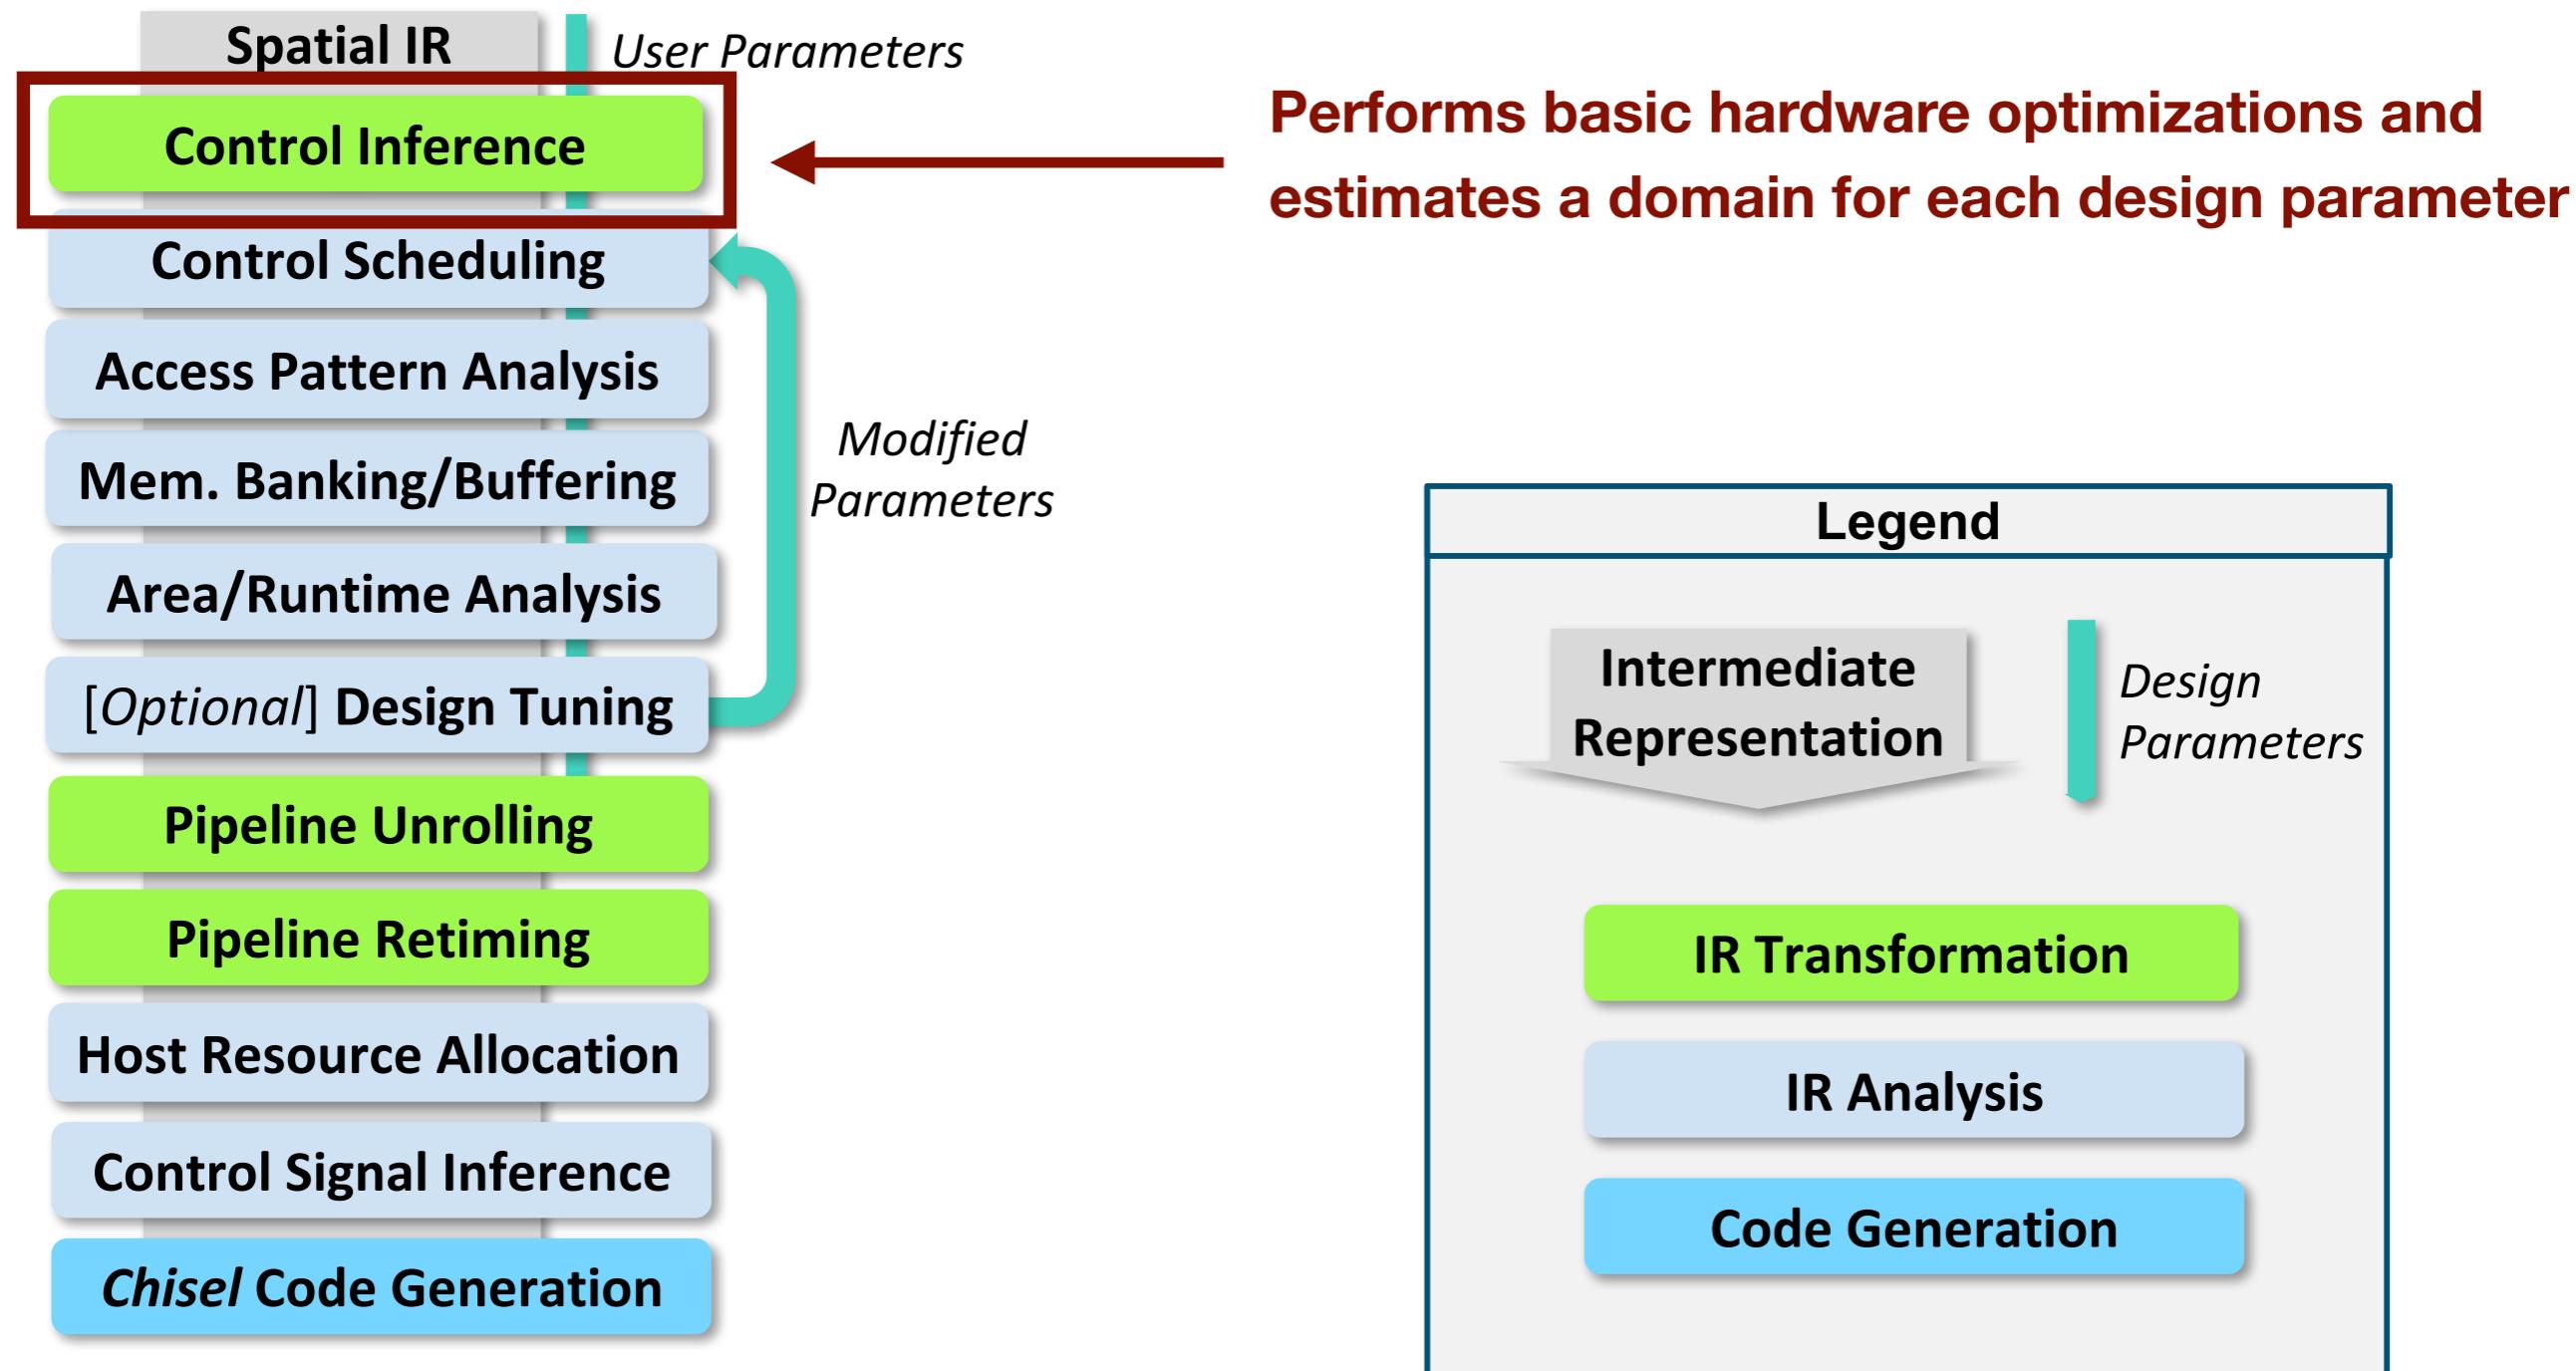
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



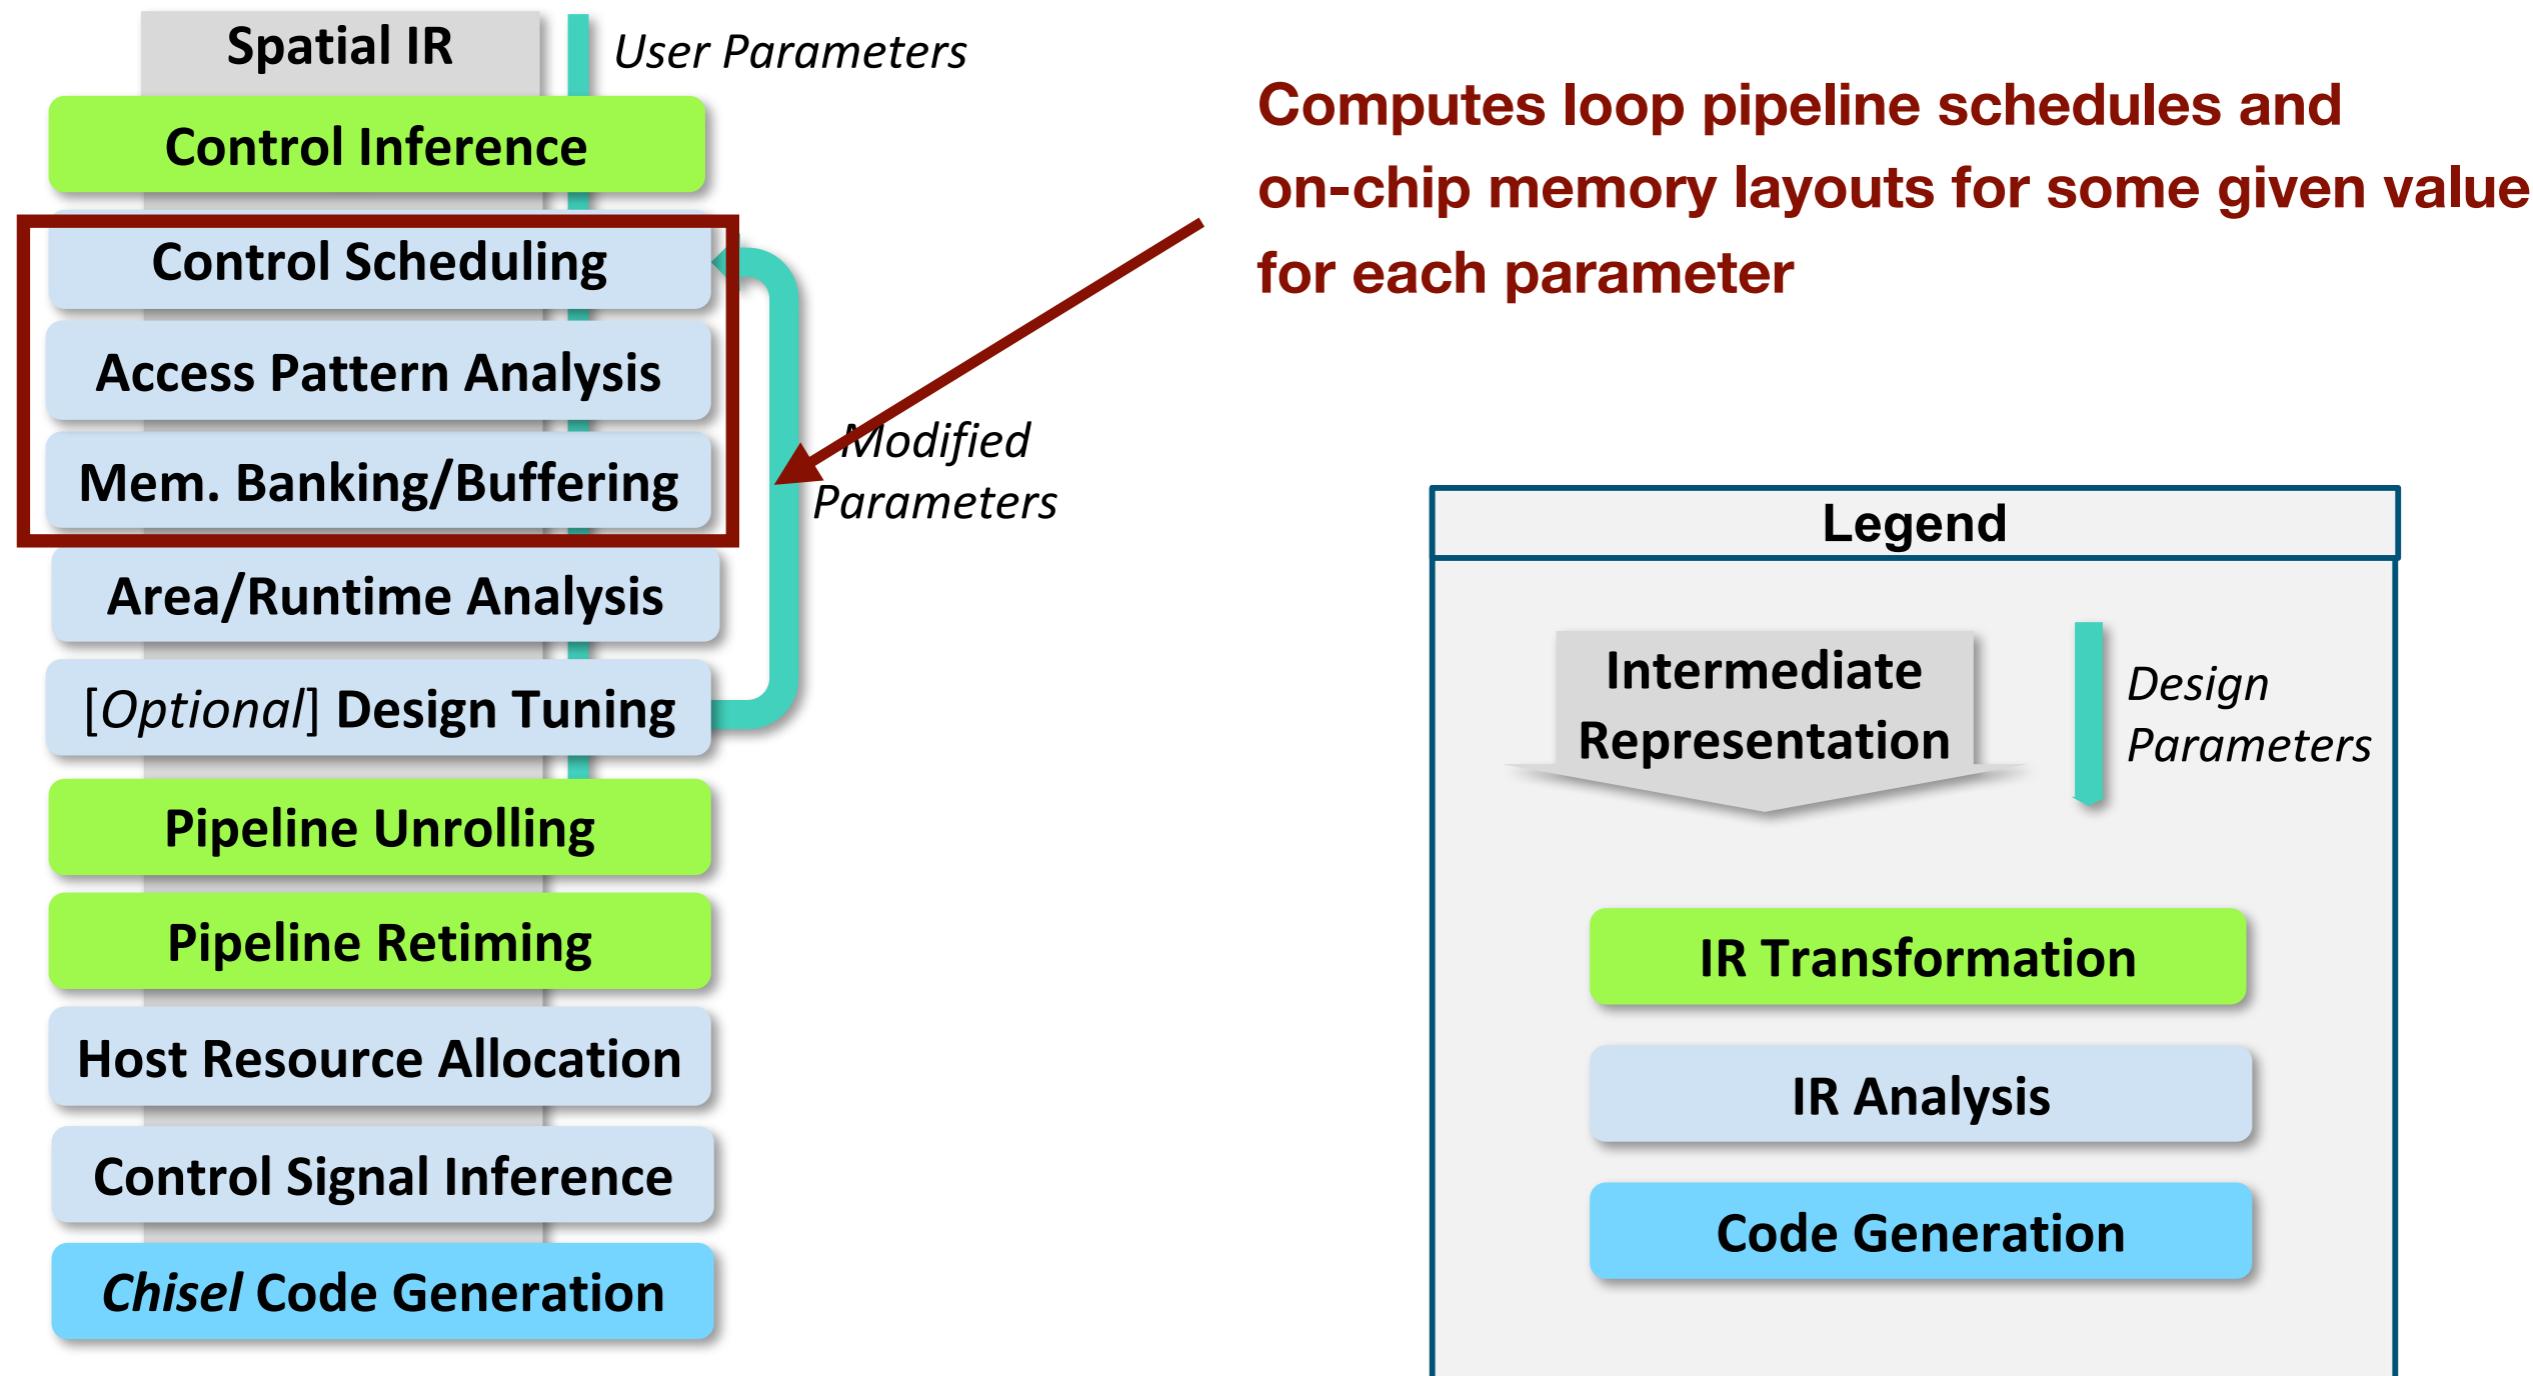
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



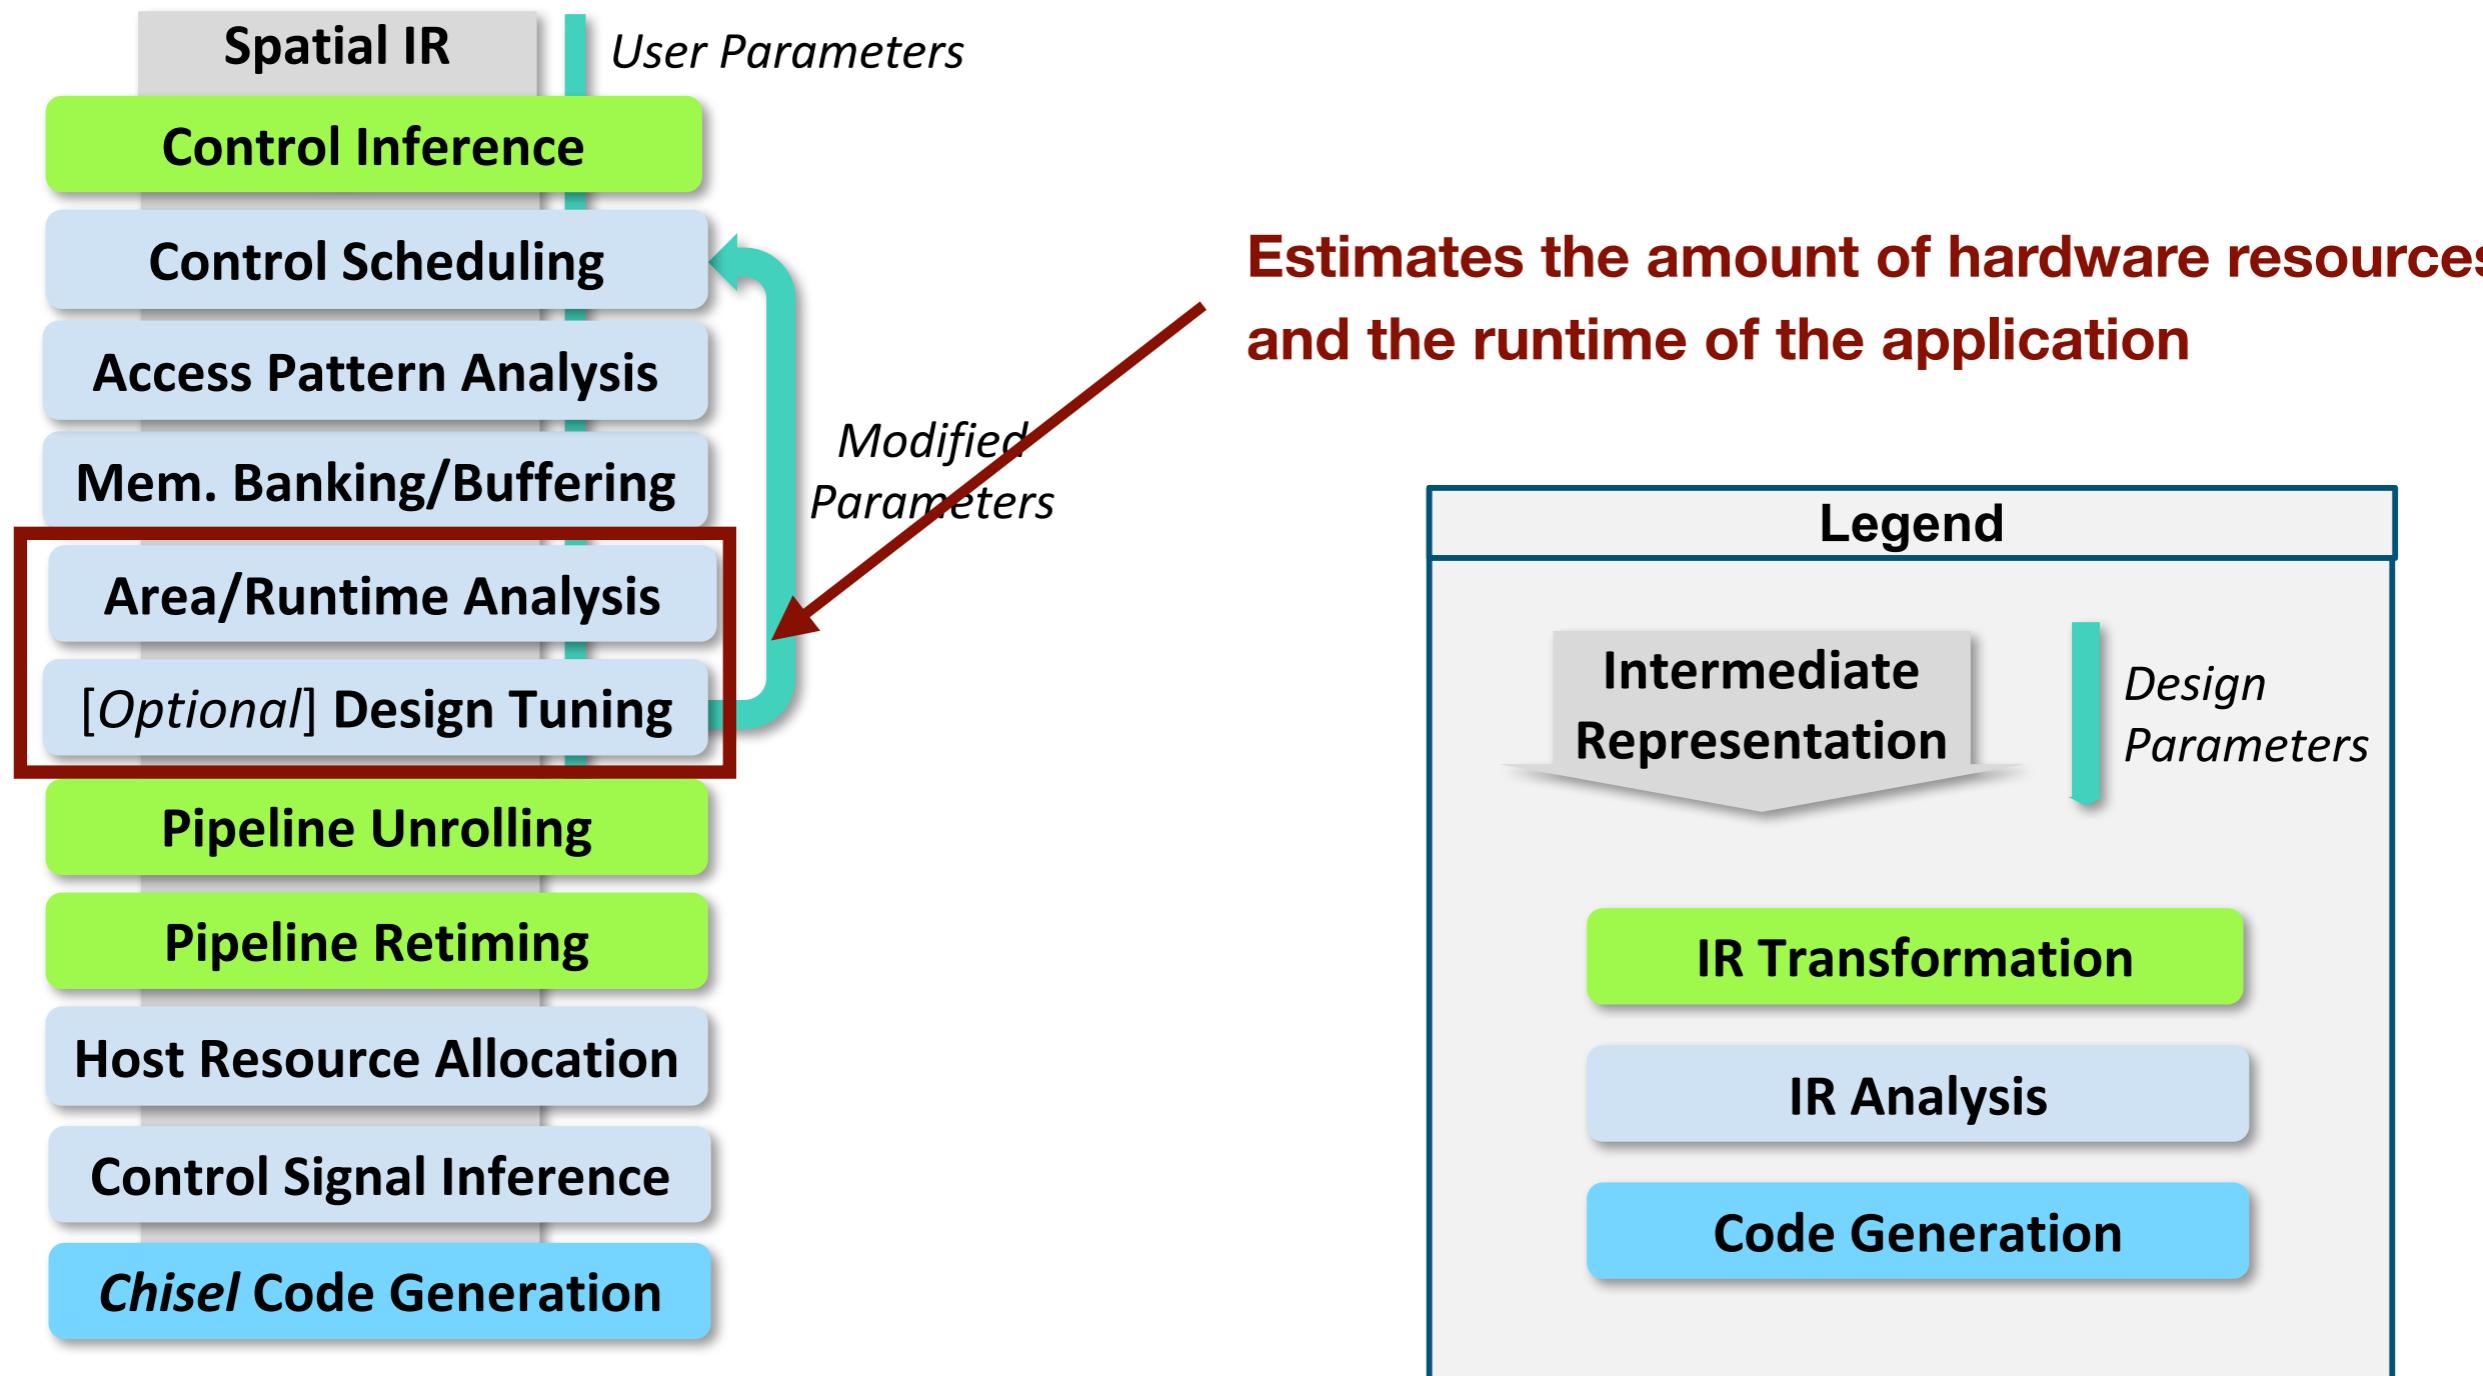
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



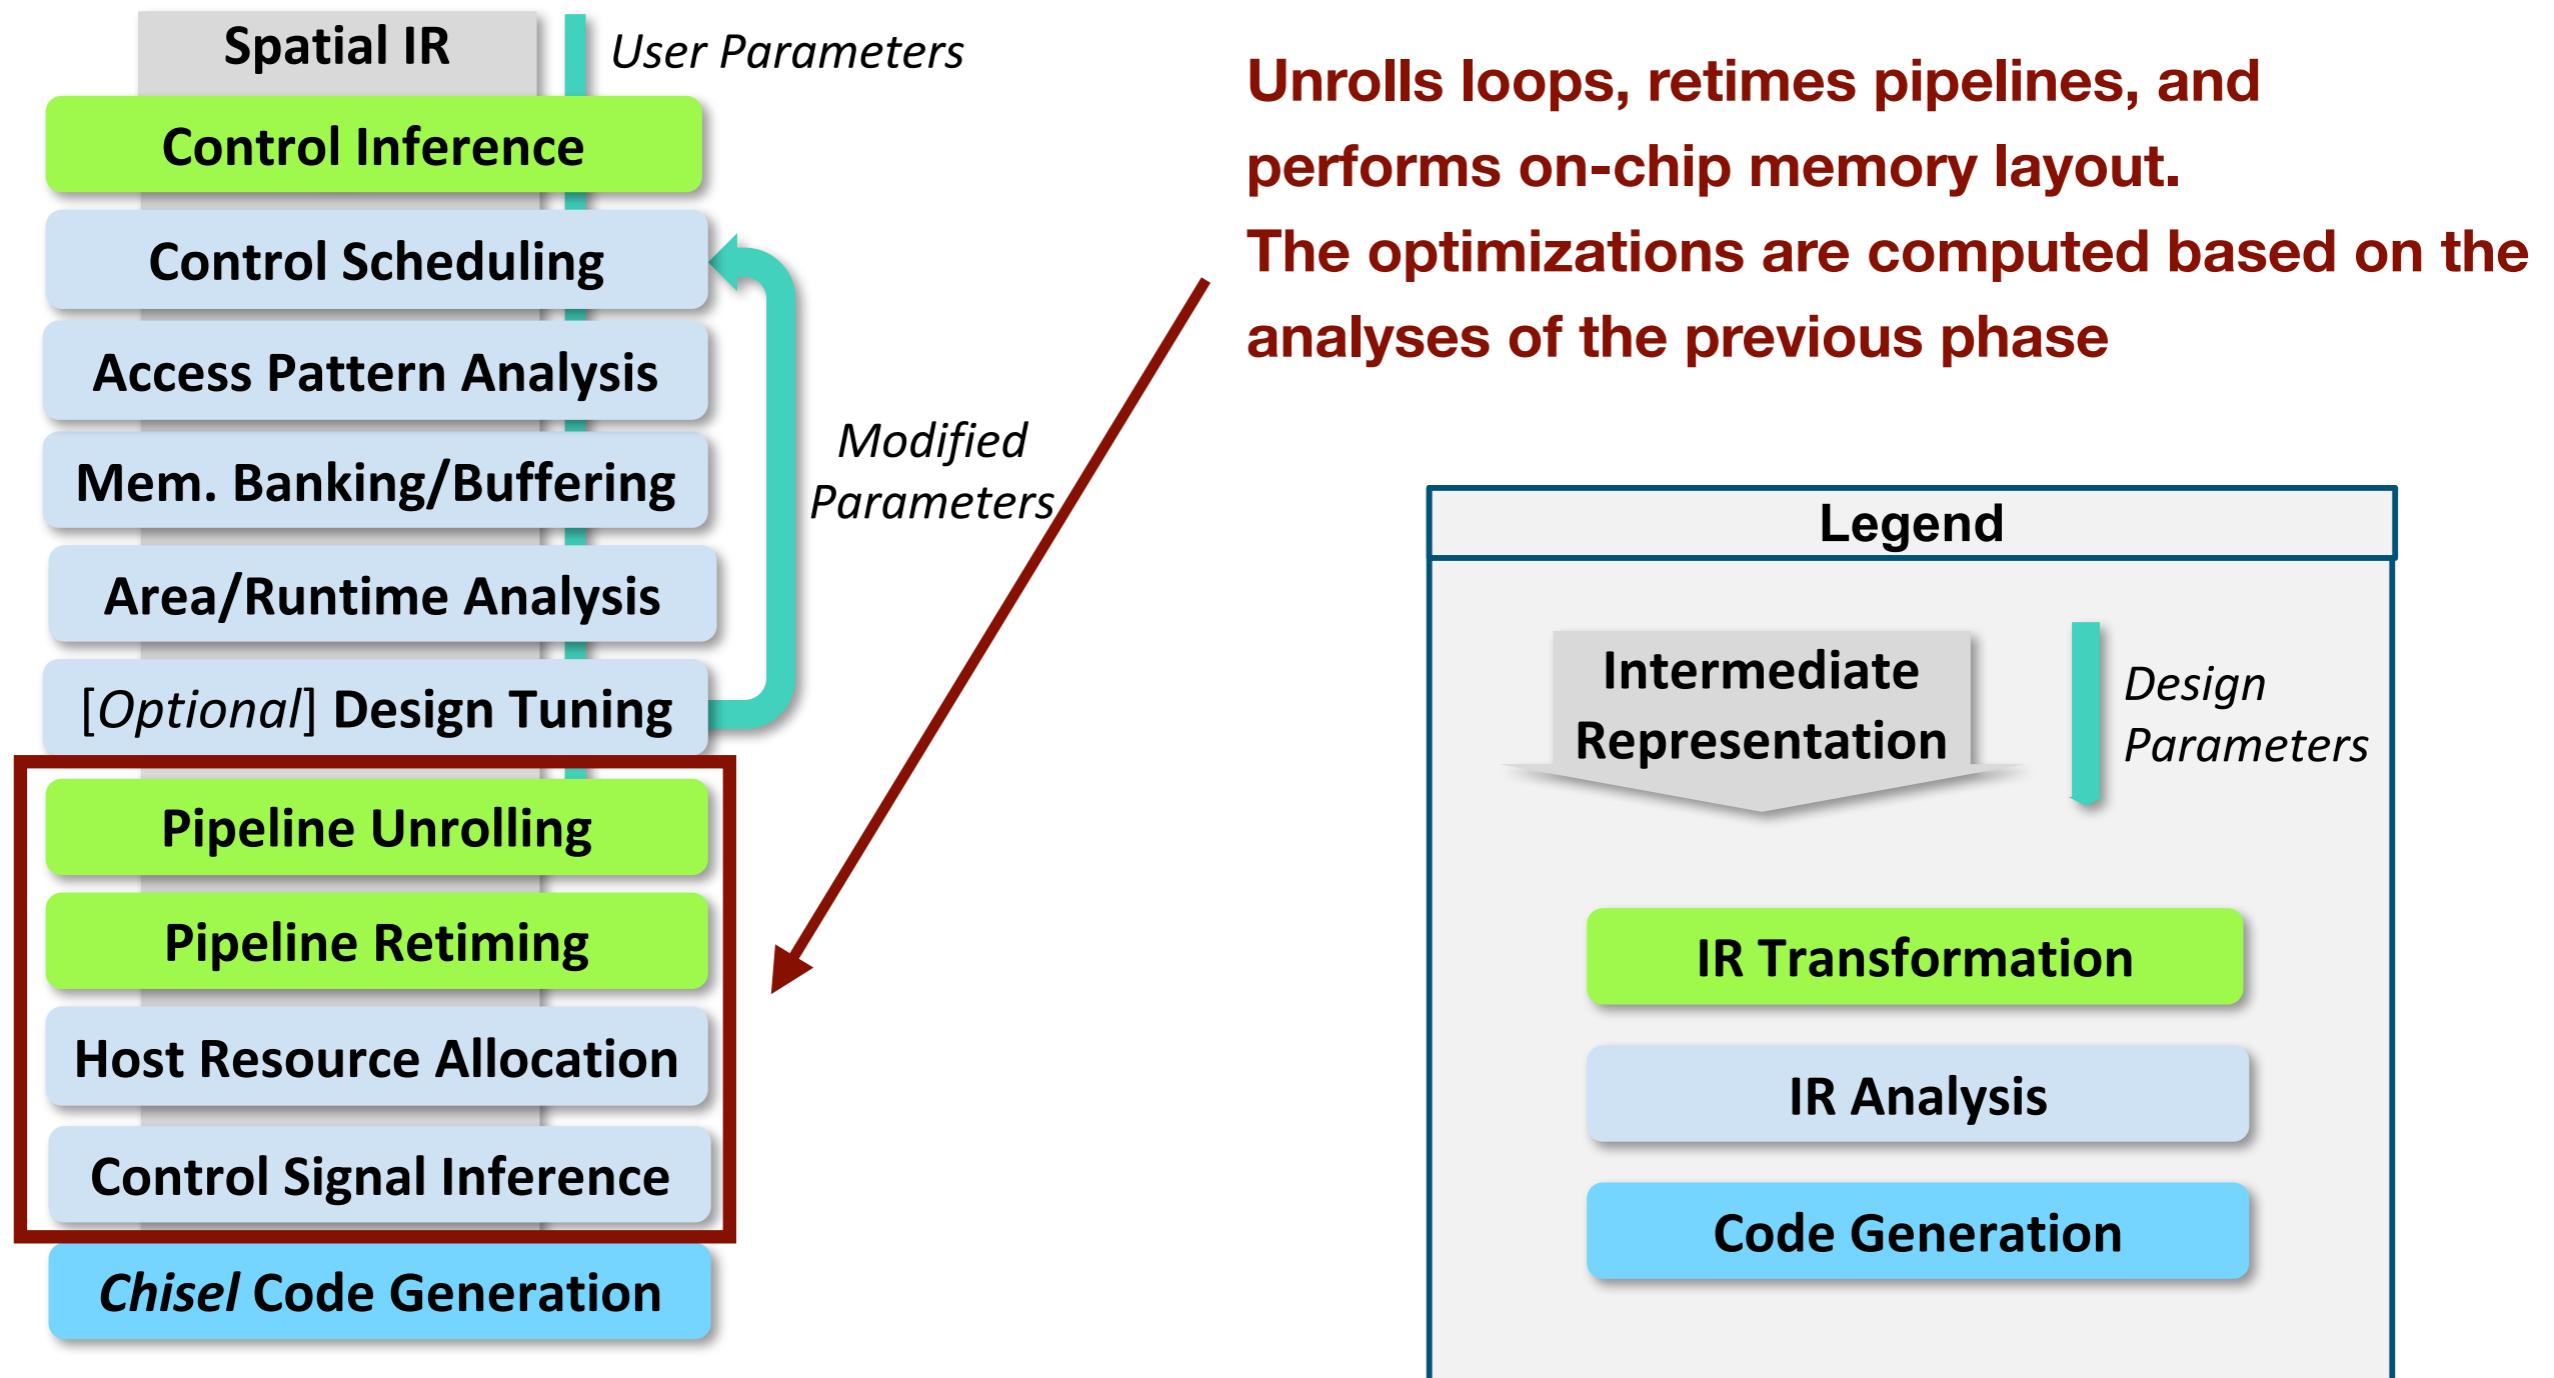
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



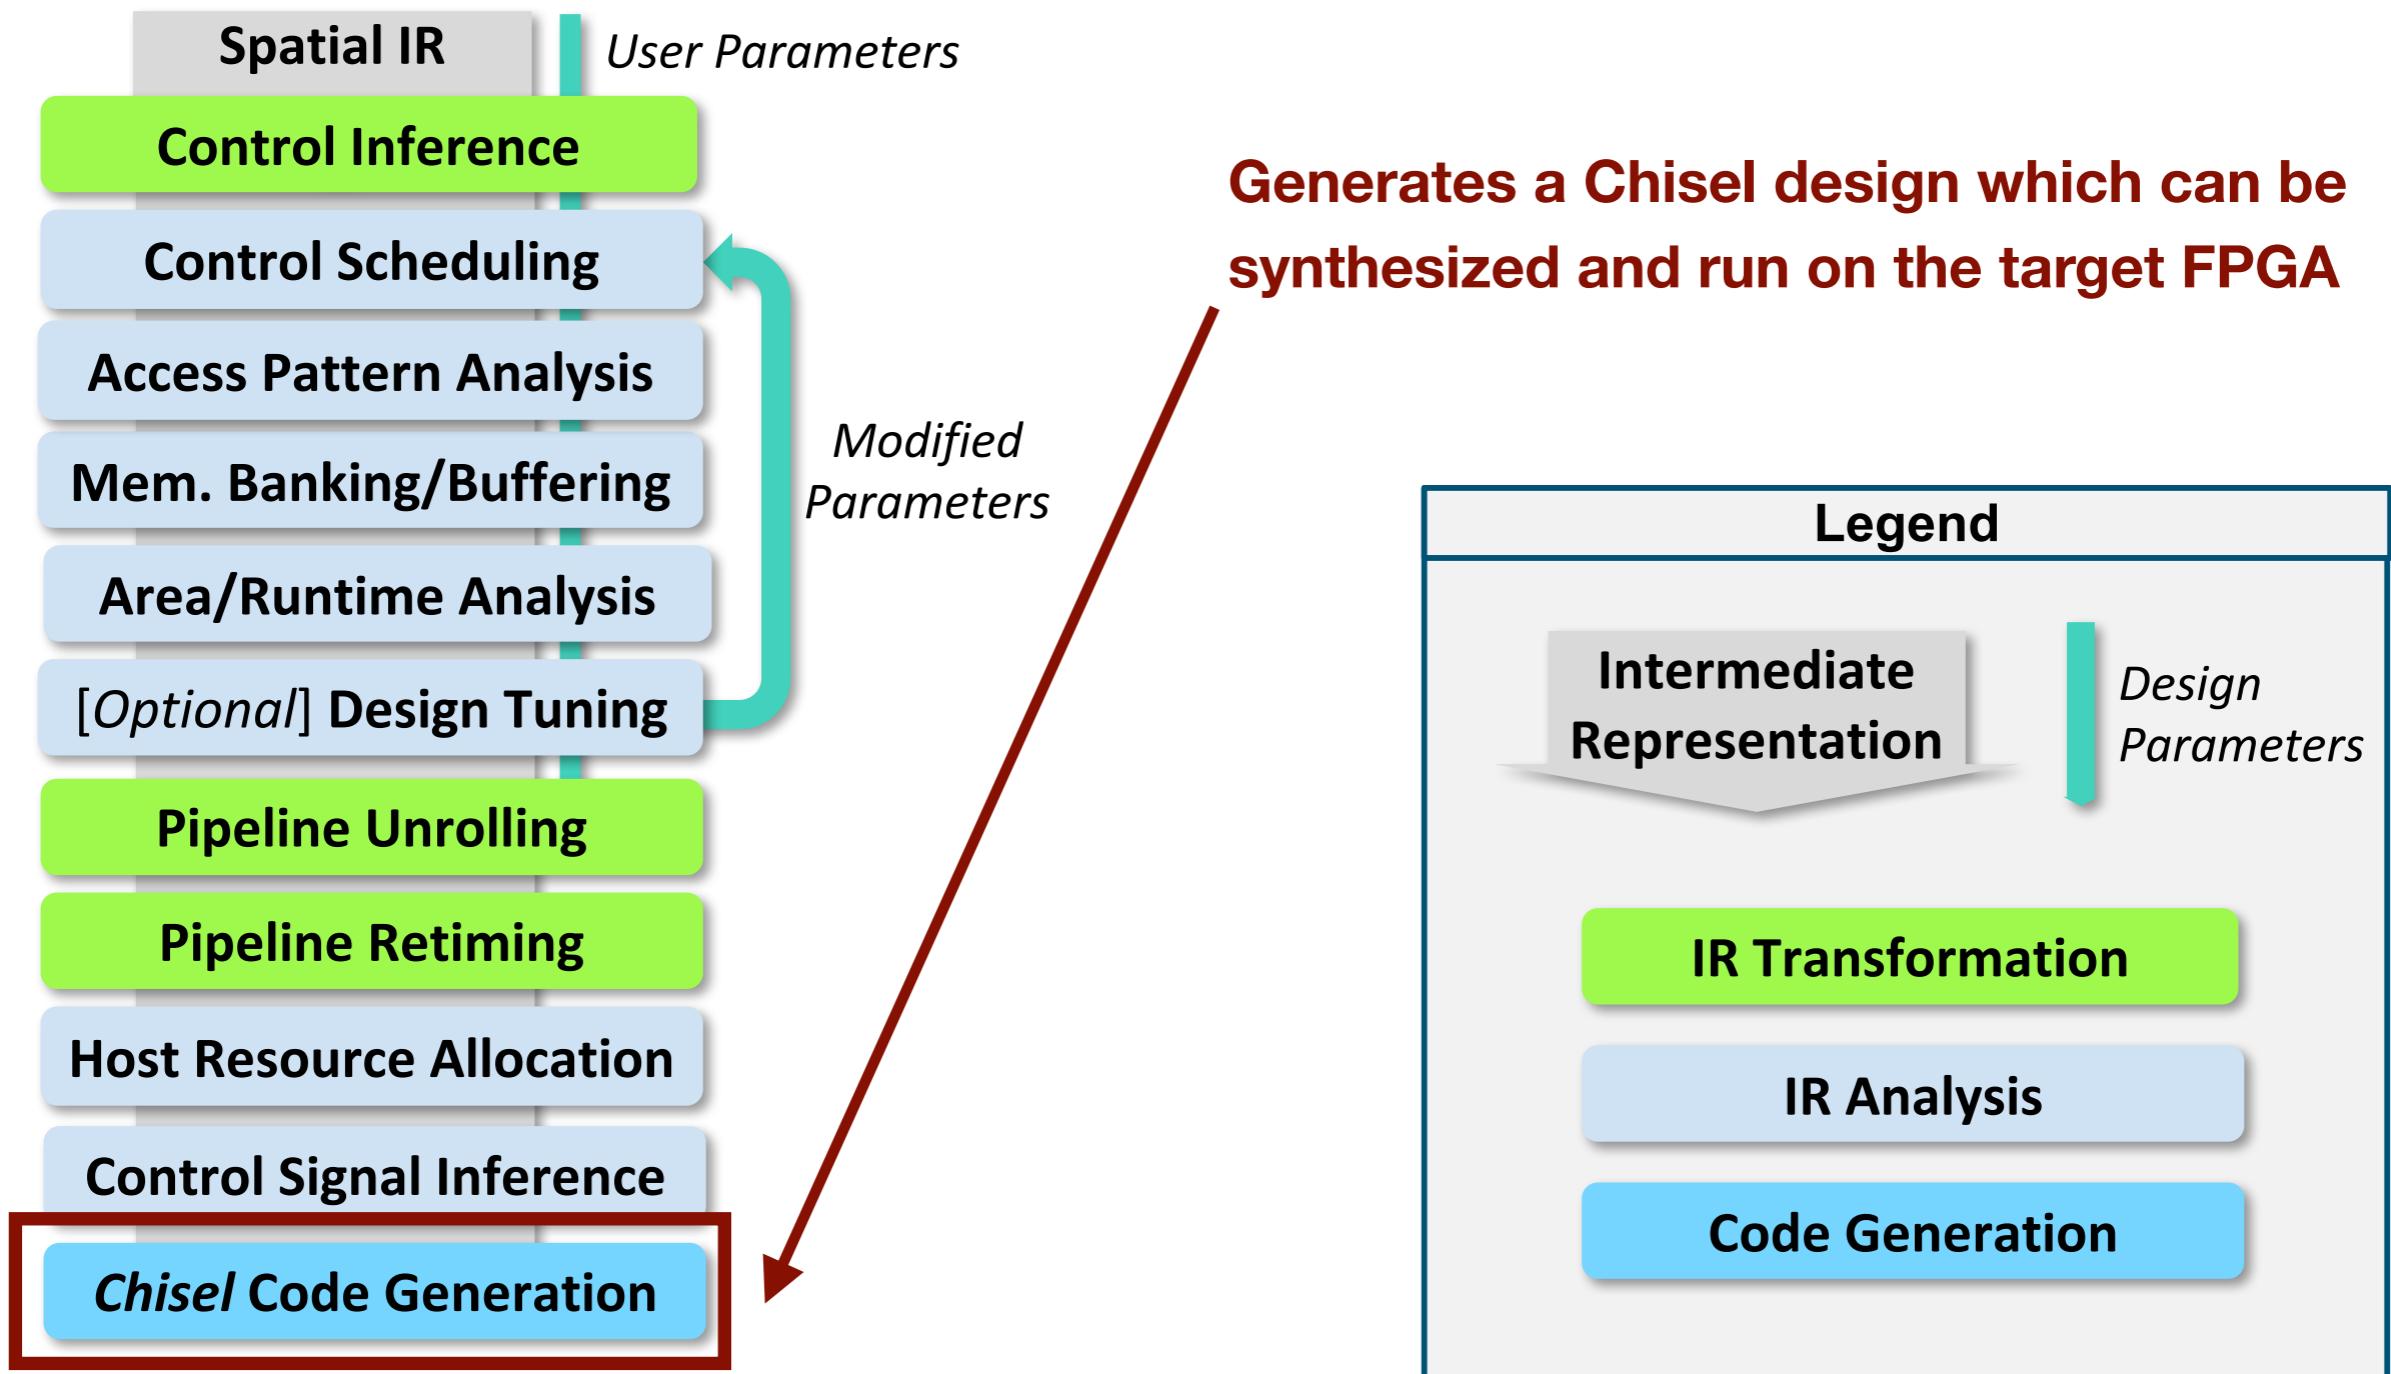
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



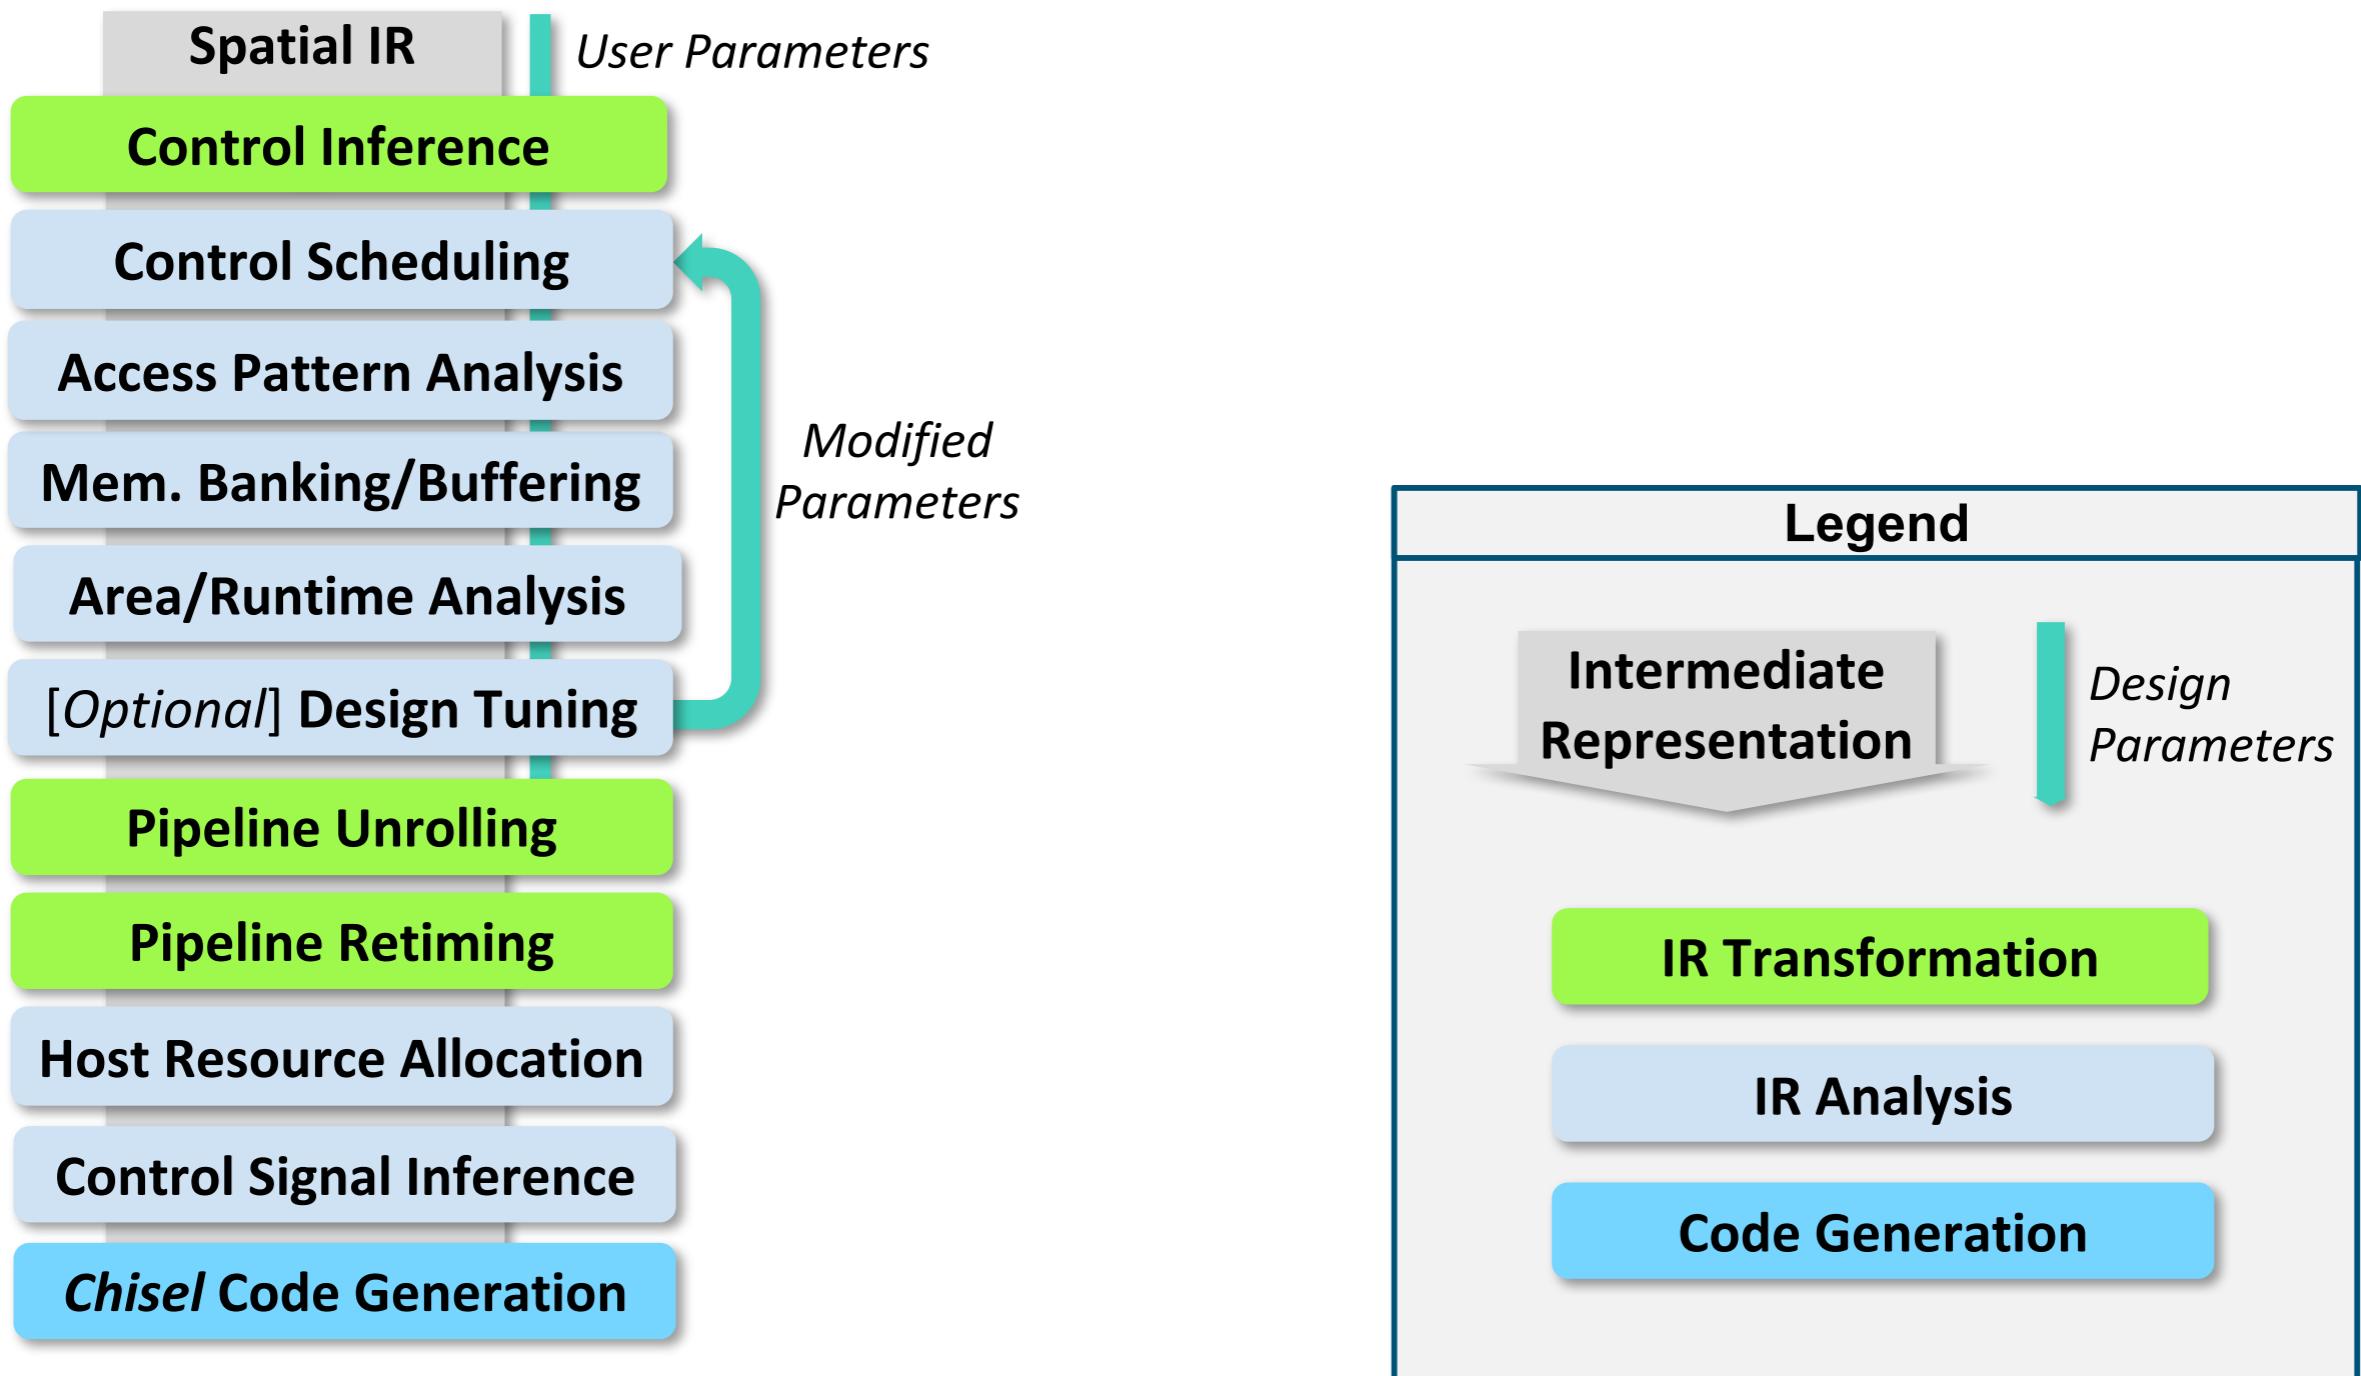
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



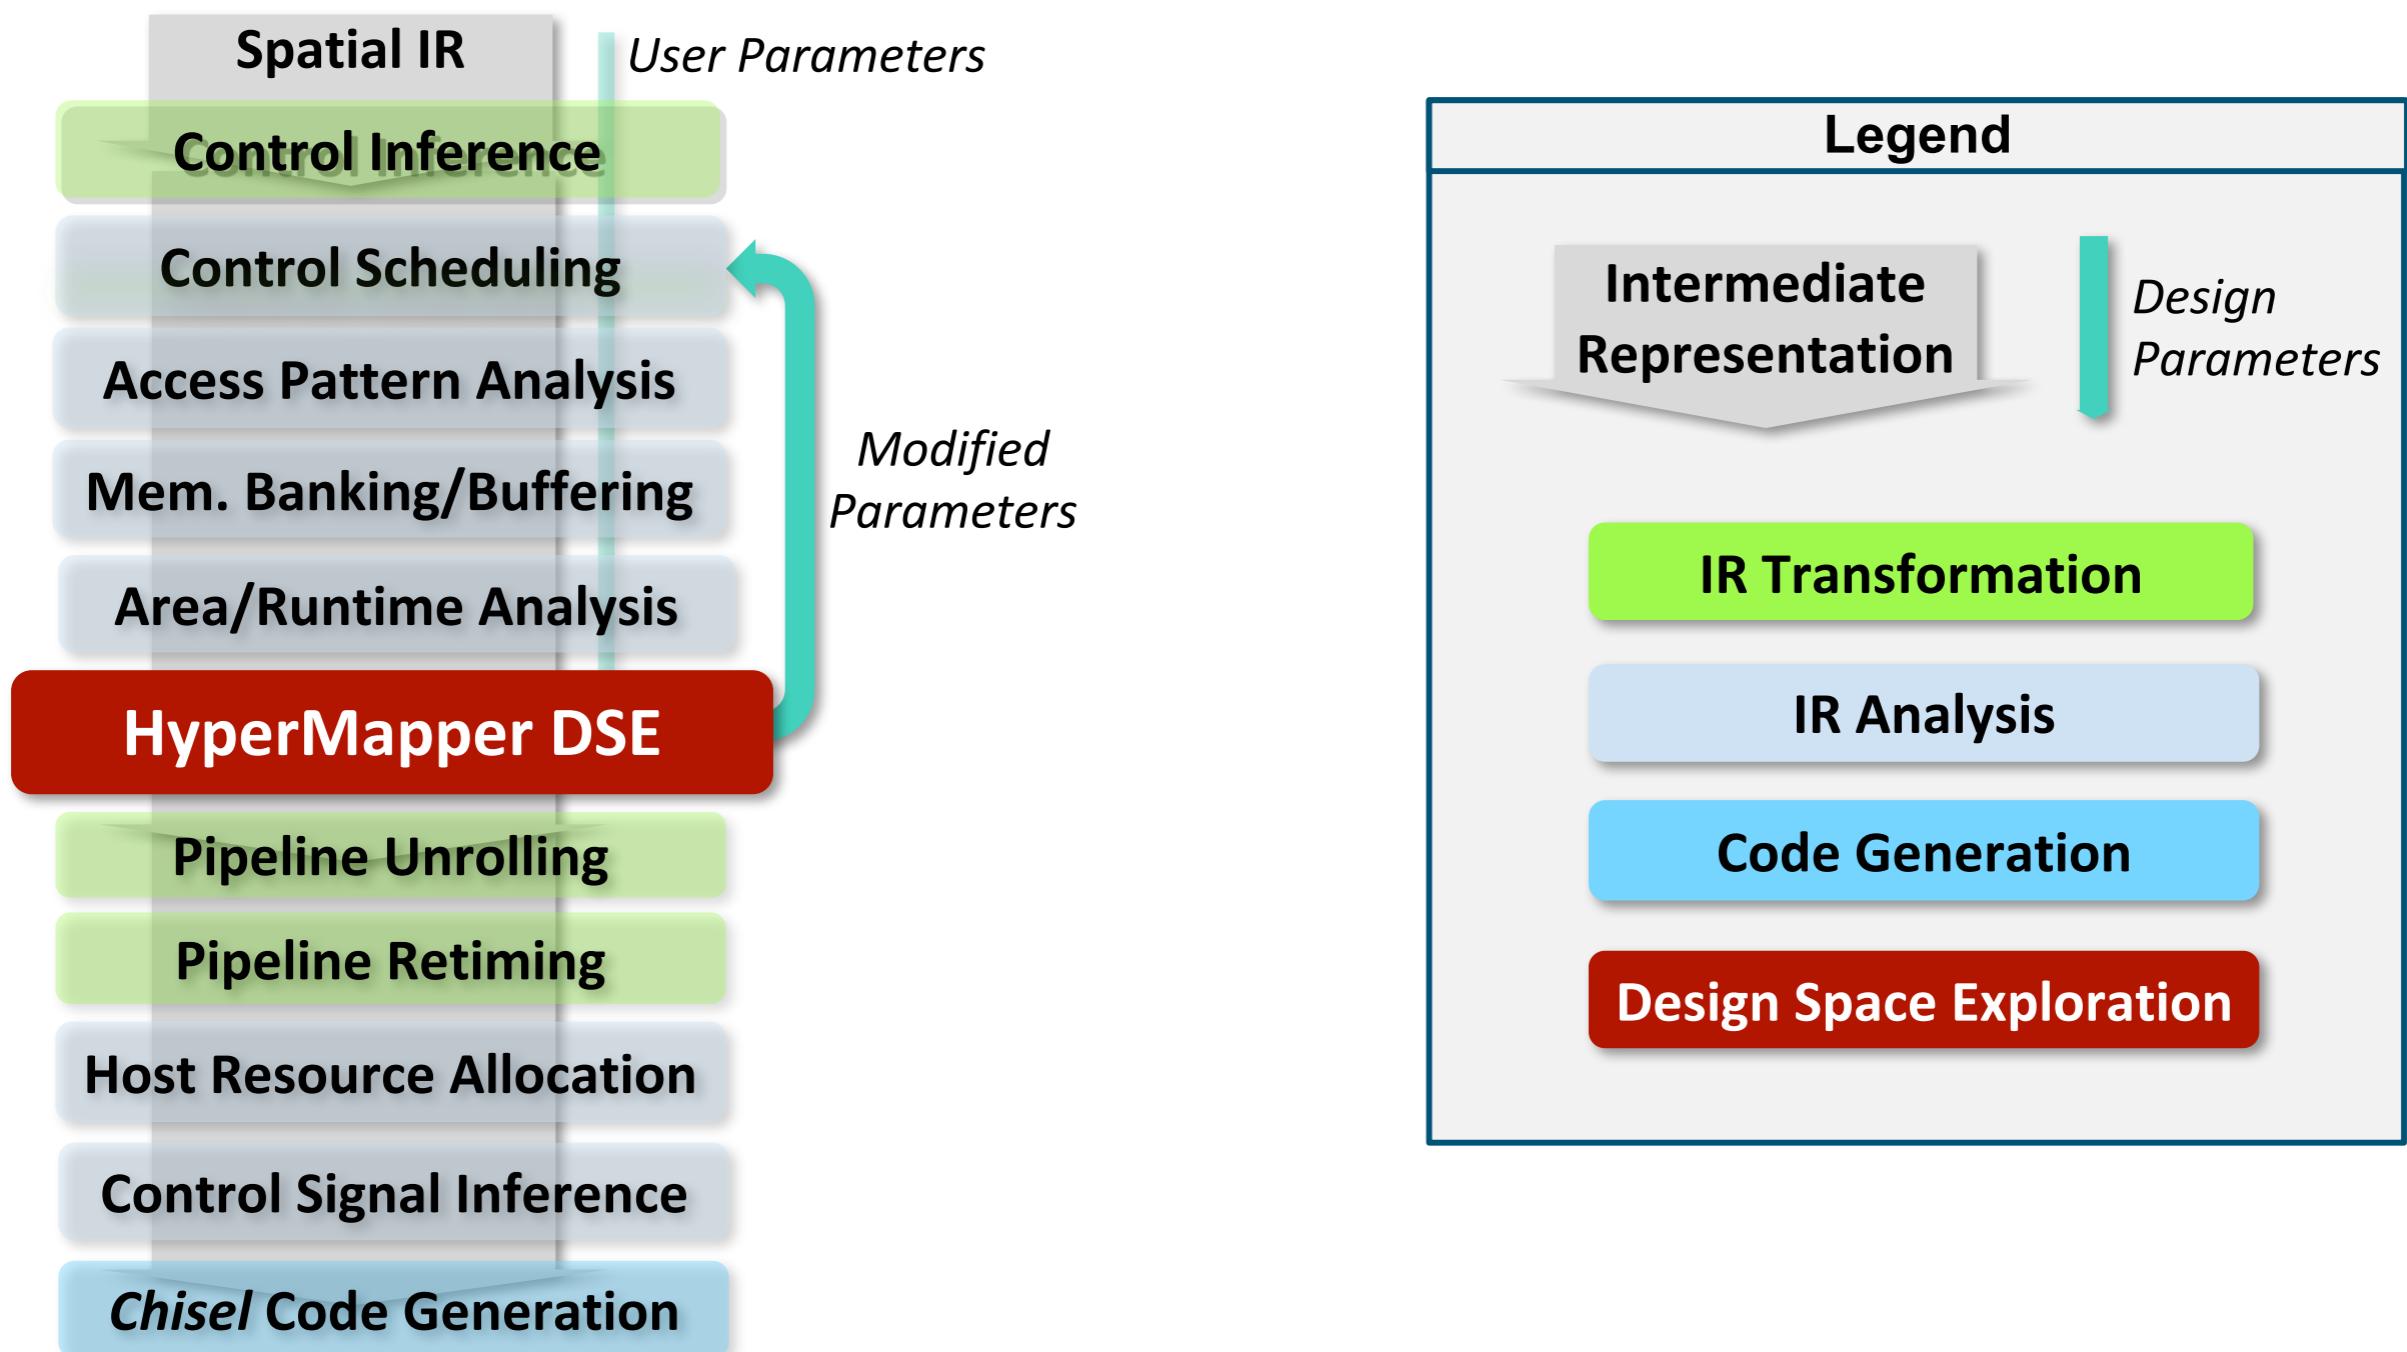
- Goal of Spatial [Koeplinger, et al.]: design of application accelerators
- On reconfigurable architectures FPGAs and CGRAs
- Spatial compiler lowers user programs into synthesizable Chisel [Bachrach, et al.]

[Koeplinger, et al.] "Spatial: a language and compiler for application accelerators." PLDI 2018

[Bachrach, et al.] "Chisel: constructing hardware in a Scala embedded language." DAC 2012.

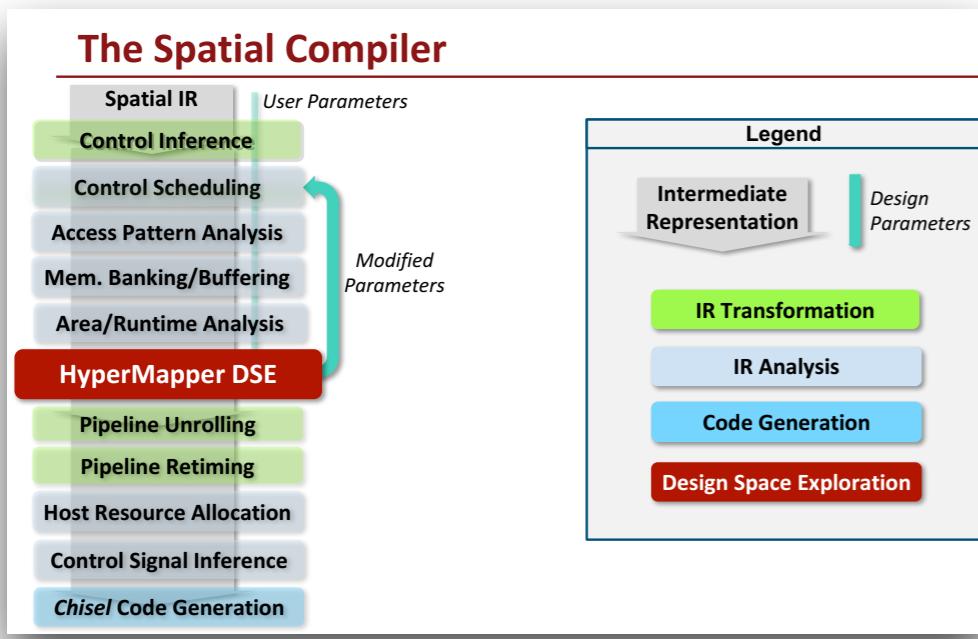
# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



# Application 2: FPGA Design Space Exploration (Spatial Compiler)

Luigi Nardi, Ph.D. - Lund/Stanford



Benchmark	Variables	Space Size
BlackScholes	4	$7.68 \times 10^4$
K-Means	6	$1.04 \times 10^6$
OuterProduct	5	$1.66 \times 10^7$
DotProduct	5	$1.18 \times 10^8$
GEMM	7	$2.62 \times 10^8$
TPC-H Q6	5	$3.54 \times 10^9$
GDA	9	$2.40 \times 10^{11}$

## Input

Compiler automatically provides params:

- Tile size (ordinal)
- Inner and outer loop pipelining (ordinal)
- Meta-pipe (categorical)
- Unrolling factor (ordinal)
- Memory banking (ordinal)
- Parallelism (categorical)

## Output

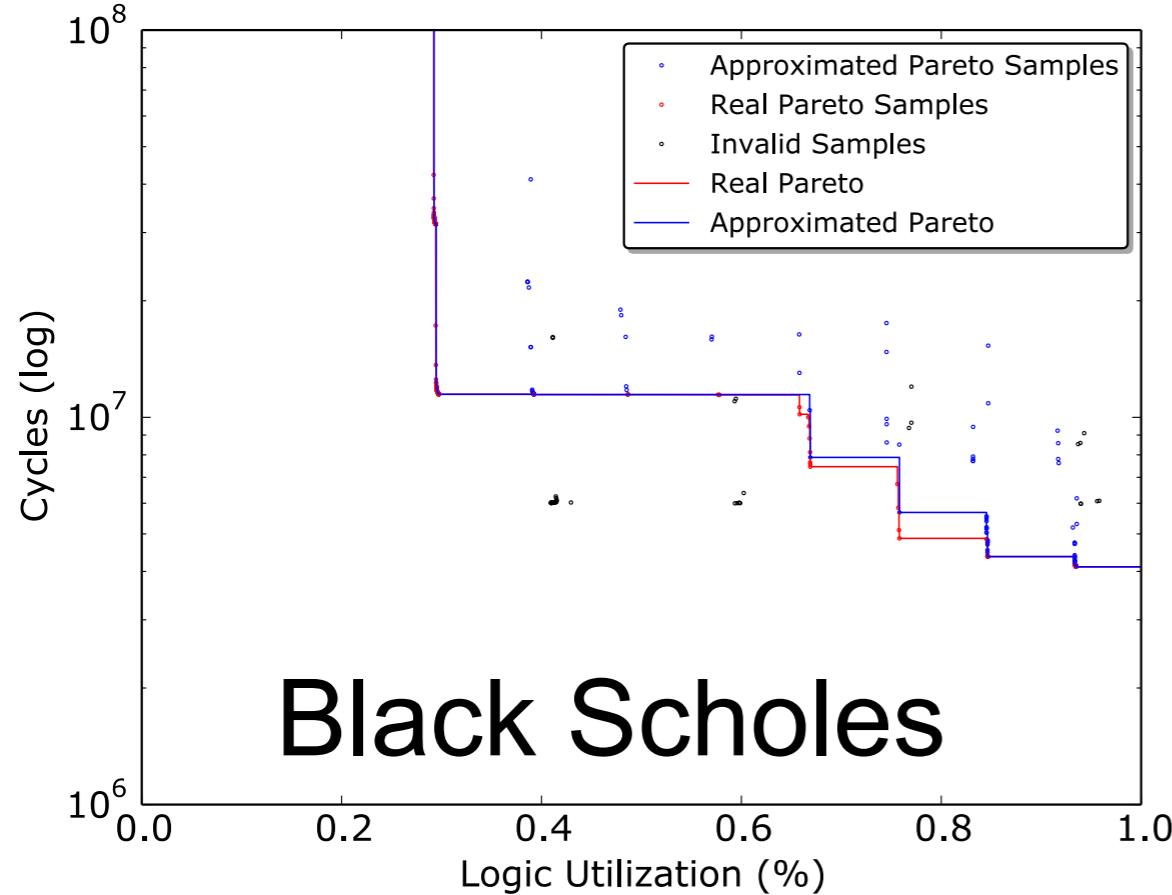
### Two objectives:

1. Minimize clock cycles (runtime)
2. Minimize FPGA logic utilization
  - For fitting multiple applications
  - Proxy for energy consumption

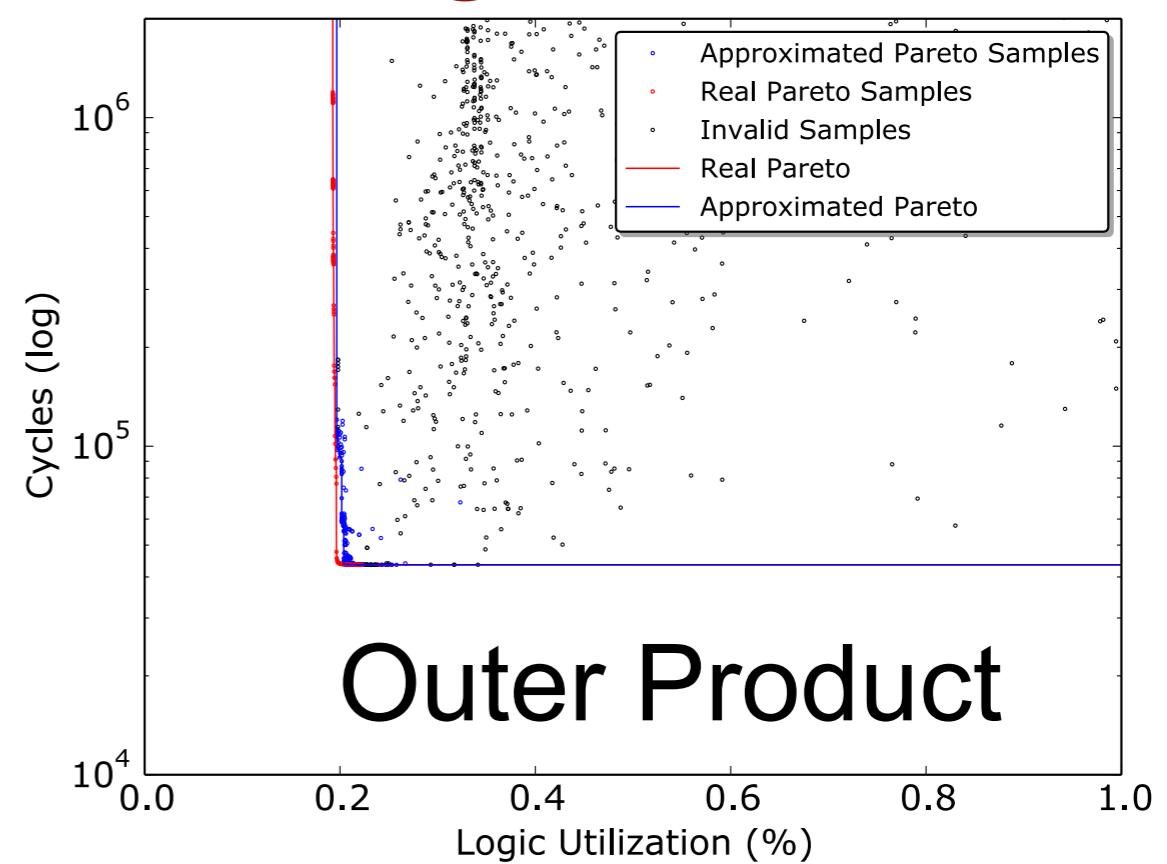
### One constraint:

design must fit in the chip

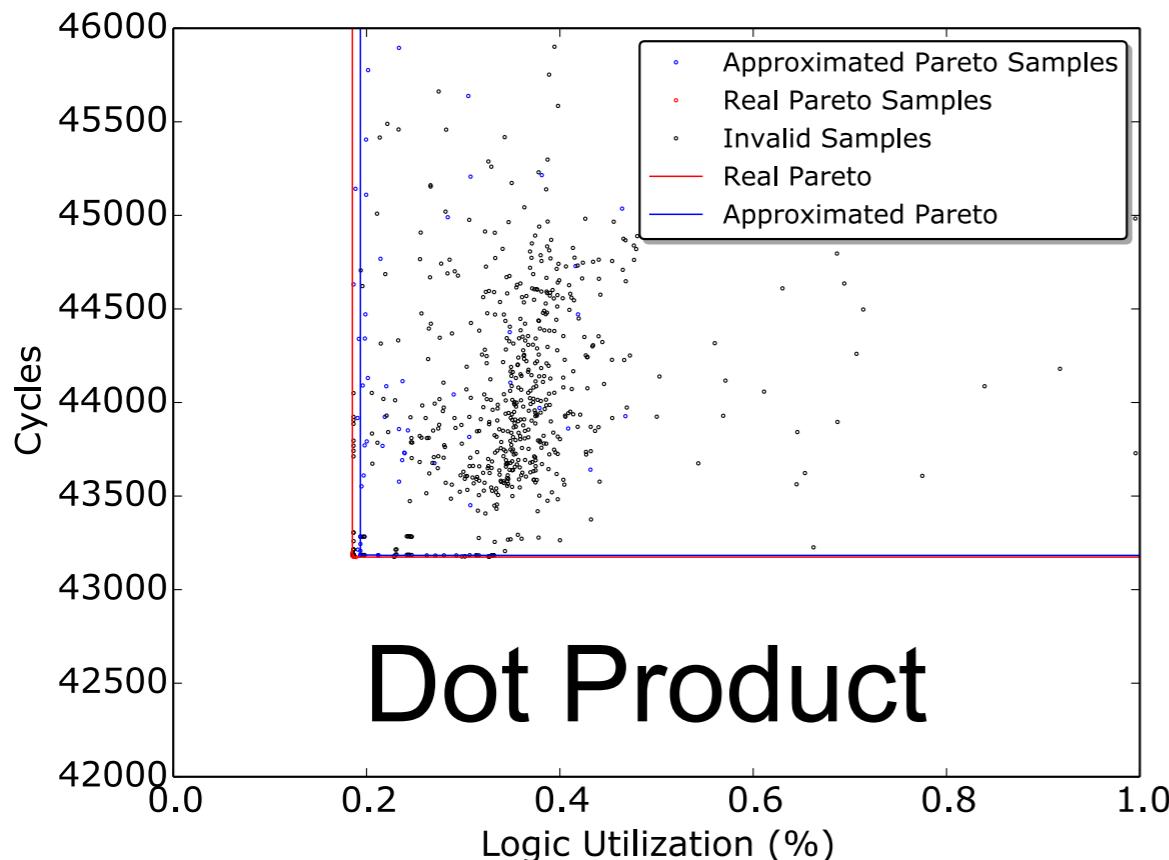
# Small benchmarks - Cycles/Logic Utilization



Black Scholes



Outer Product

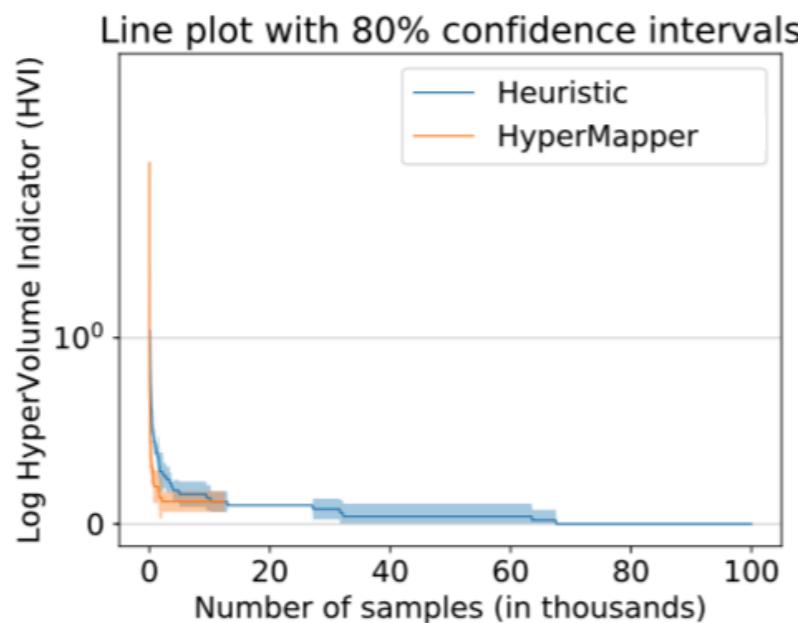


Dot Product

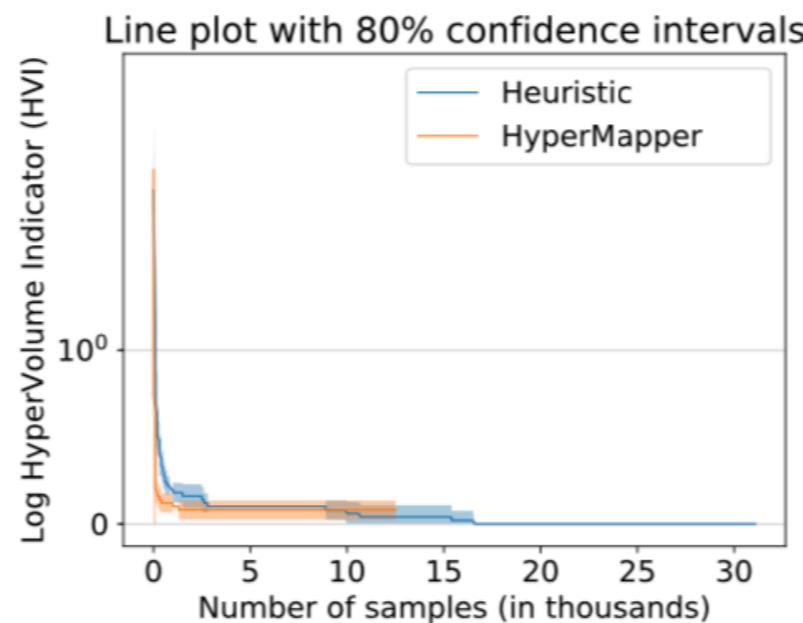
- 3 small benchmarks: compute real Pareto
- Real Pareto is exhaustive search
- Exhaustive: 6 to 12 hours on 16 cores
- Approximated Pareto is HyperMapper

# HVI Cycles/Logic Utilization

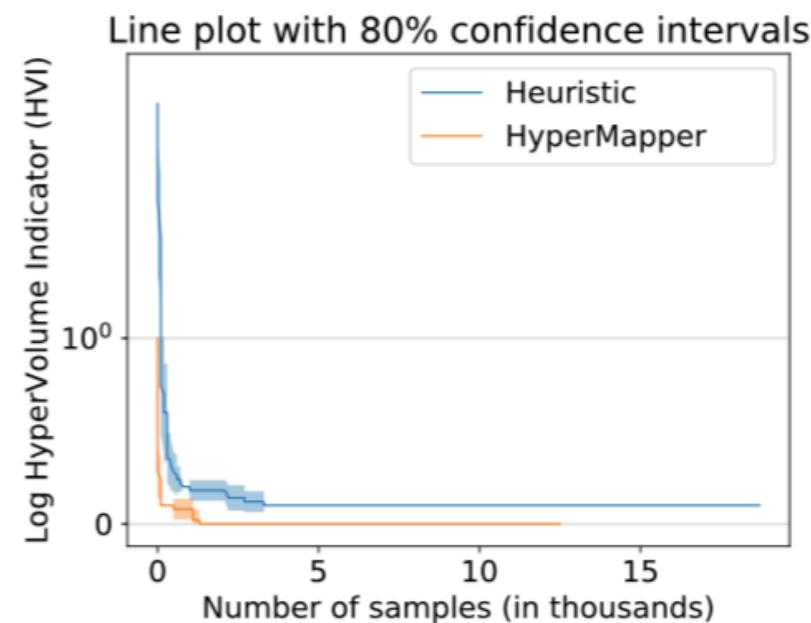
- Hypervolume indicator (HVI): scalar to compare Paretos
- Baseline is a pruning (based on heuristics) and 100K random sampling



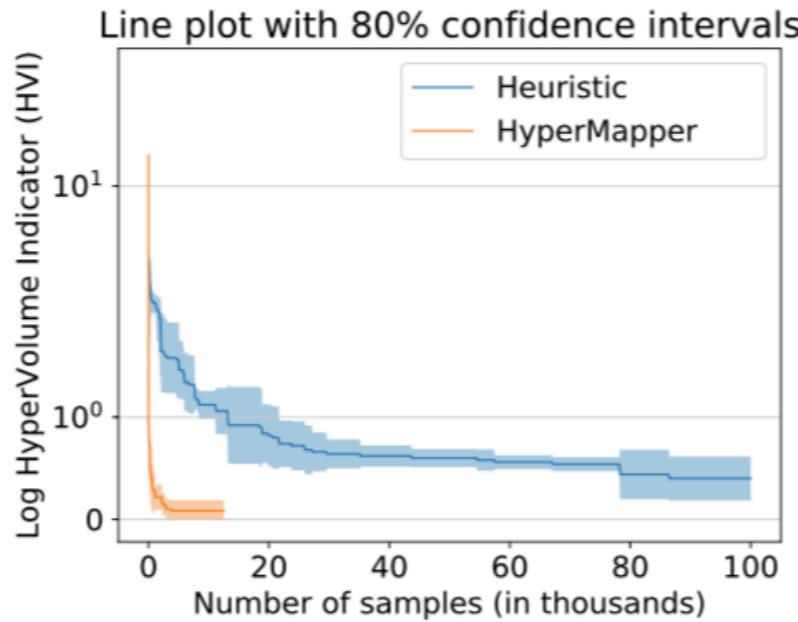
GEMM



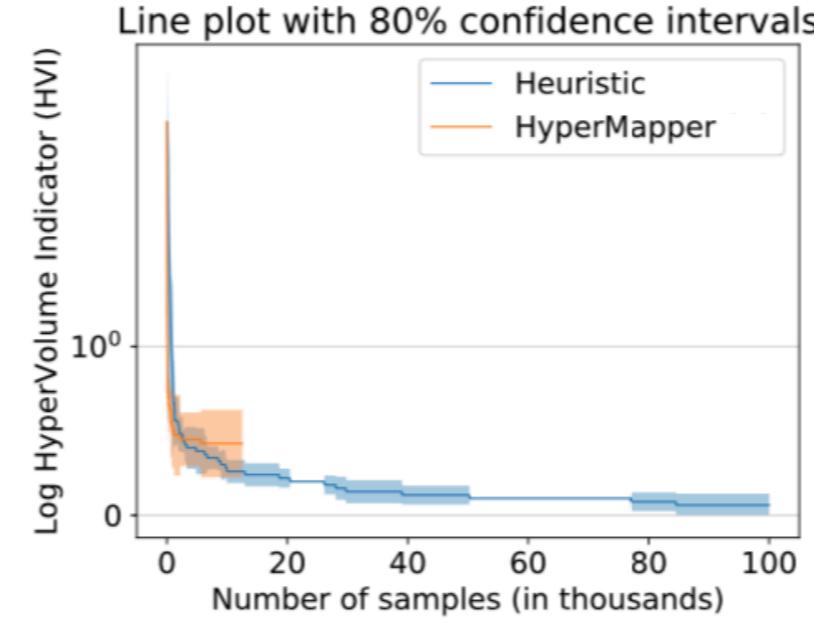
OuterProduct



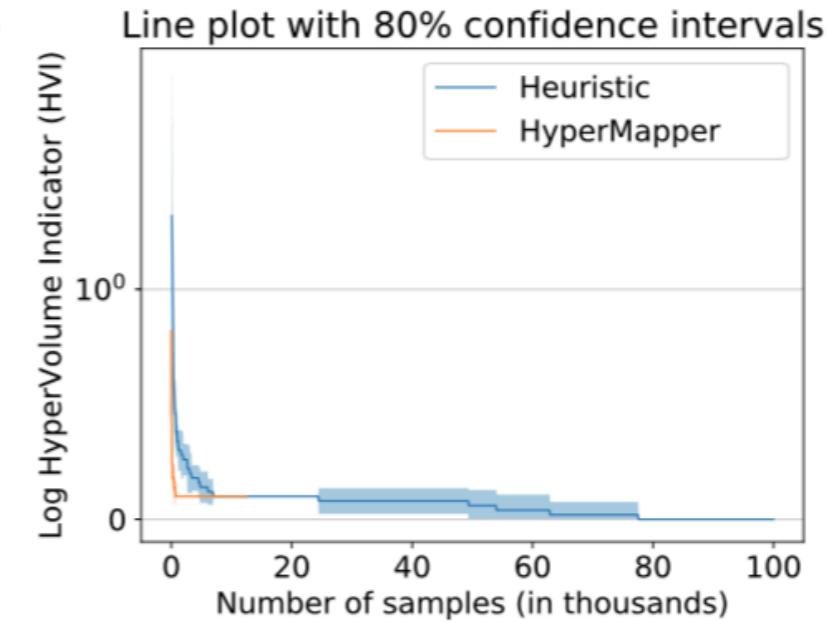
K-Means



GDA



T-PCH Q6



DotProduct

## Application 3:

# Simultaneous Localization And Mapping Optimization

Build a coherent world representation and localize the camera in real-time

I will focus on two SLAM algorithms from SLAMBench:

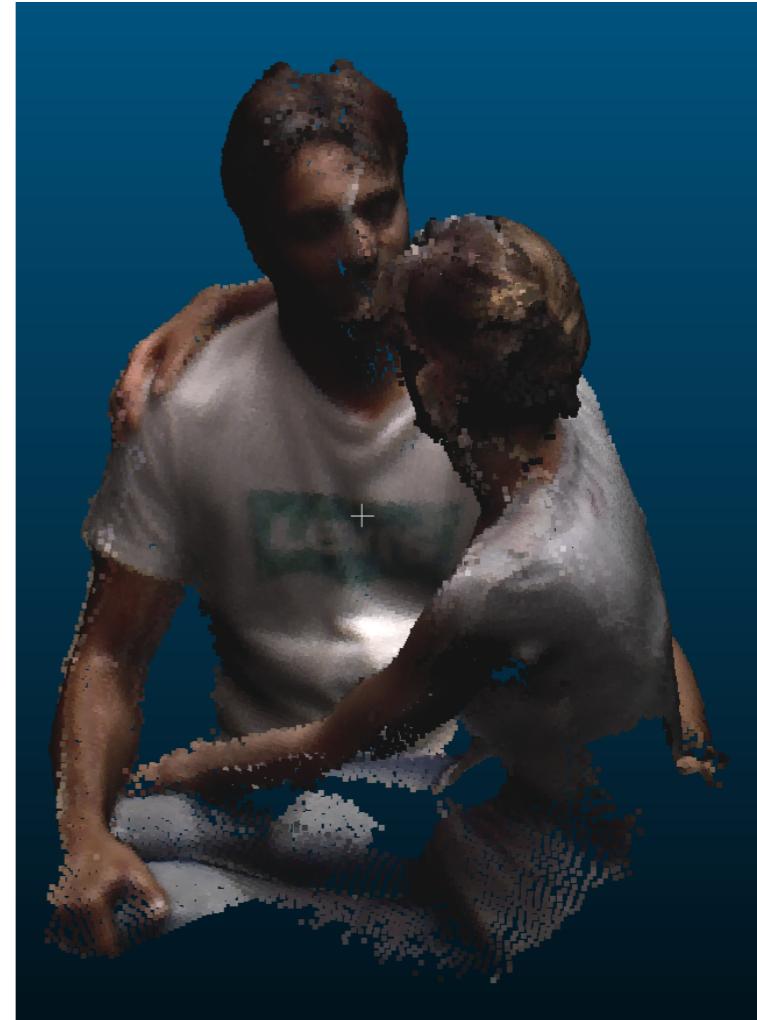
- KinectFusion [[Newcombe et al. ISMAR 2011](#)]
- ElasticFusion [[Whelan et al. RSS 2015](#)]



**Jesse Clayton (NVIDIA)**  
**3D reconstruction**

Pictures featured using ElasticFusion

- Applications, e.g.:
- Robotics
  - Autonomous driving
  - 3D printing
  - Augmented reality
  - Telepresence

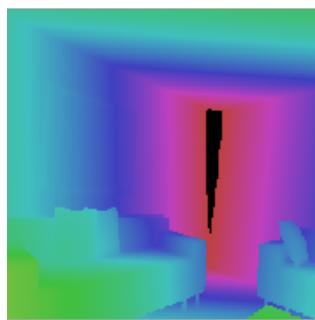


**Daniele and Daniela**  
**3D reconstruction**

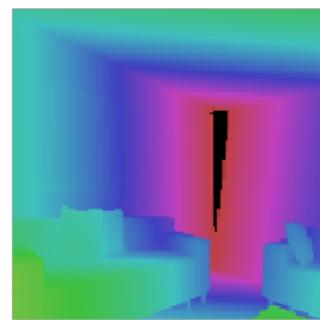
# Application 3: SLAM (KinectFusion) algorithmic features

Features	Ranges
Volume resolution	64x64x64, 128x128x128, 256x256x256, 512x512x512
$\mu$ distance	0 .. 0.5
Pyramid level iterations (3 levels)	0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11
<b>Image resolution (image ratio)</b>	1, 2, 4, 8
Tracking rate	1, 2, 3, 4, 5
ICP threshold	10 <sup>-6</sup> .. 10 <sup>2</sup>
Integration rate	1 .. 30

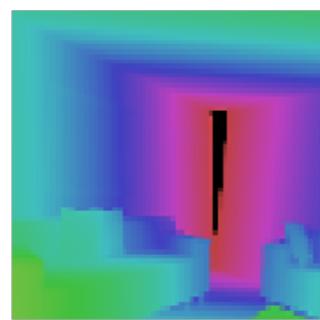
Image resolution (image ratio)



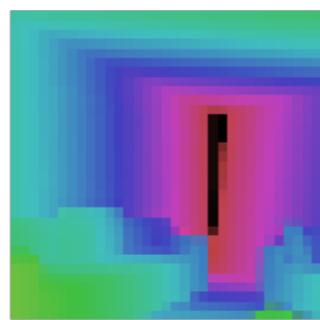
640x480



320x240



160x120

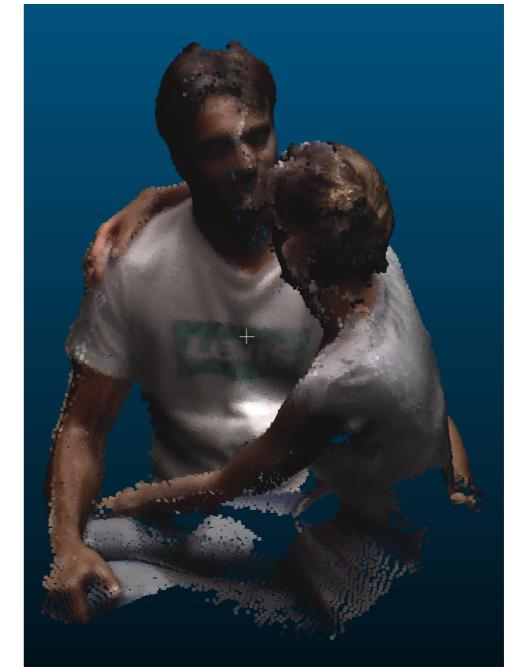
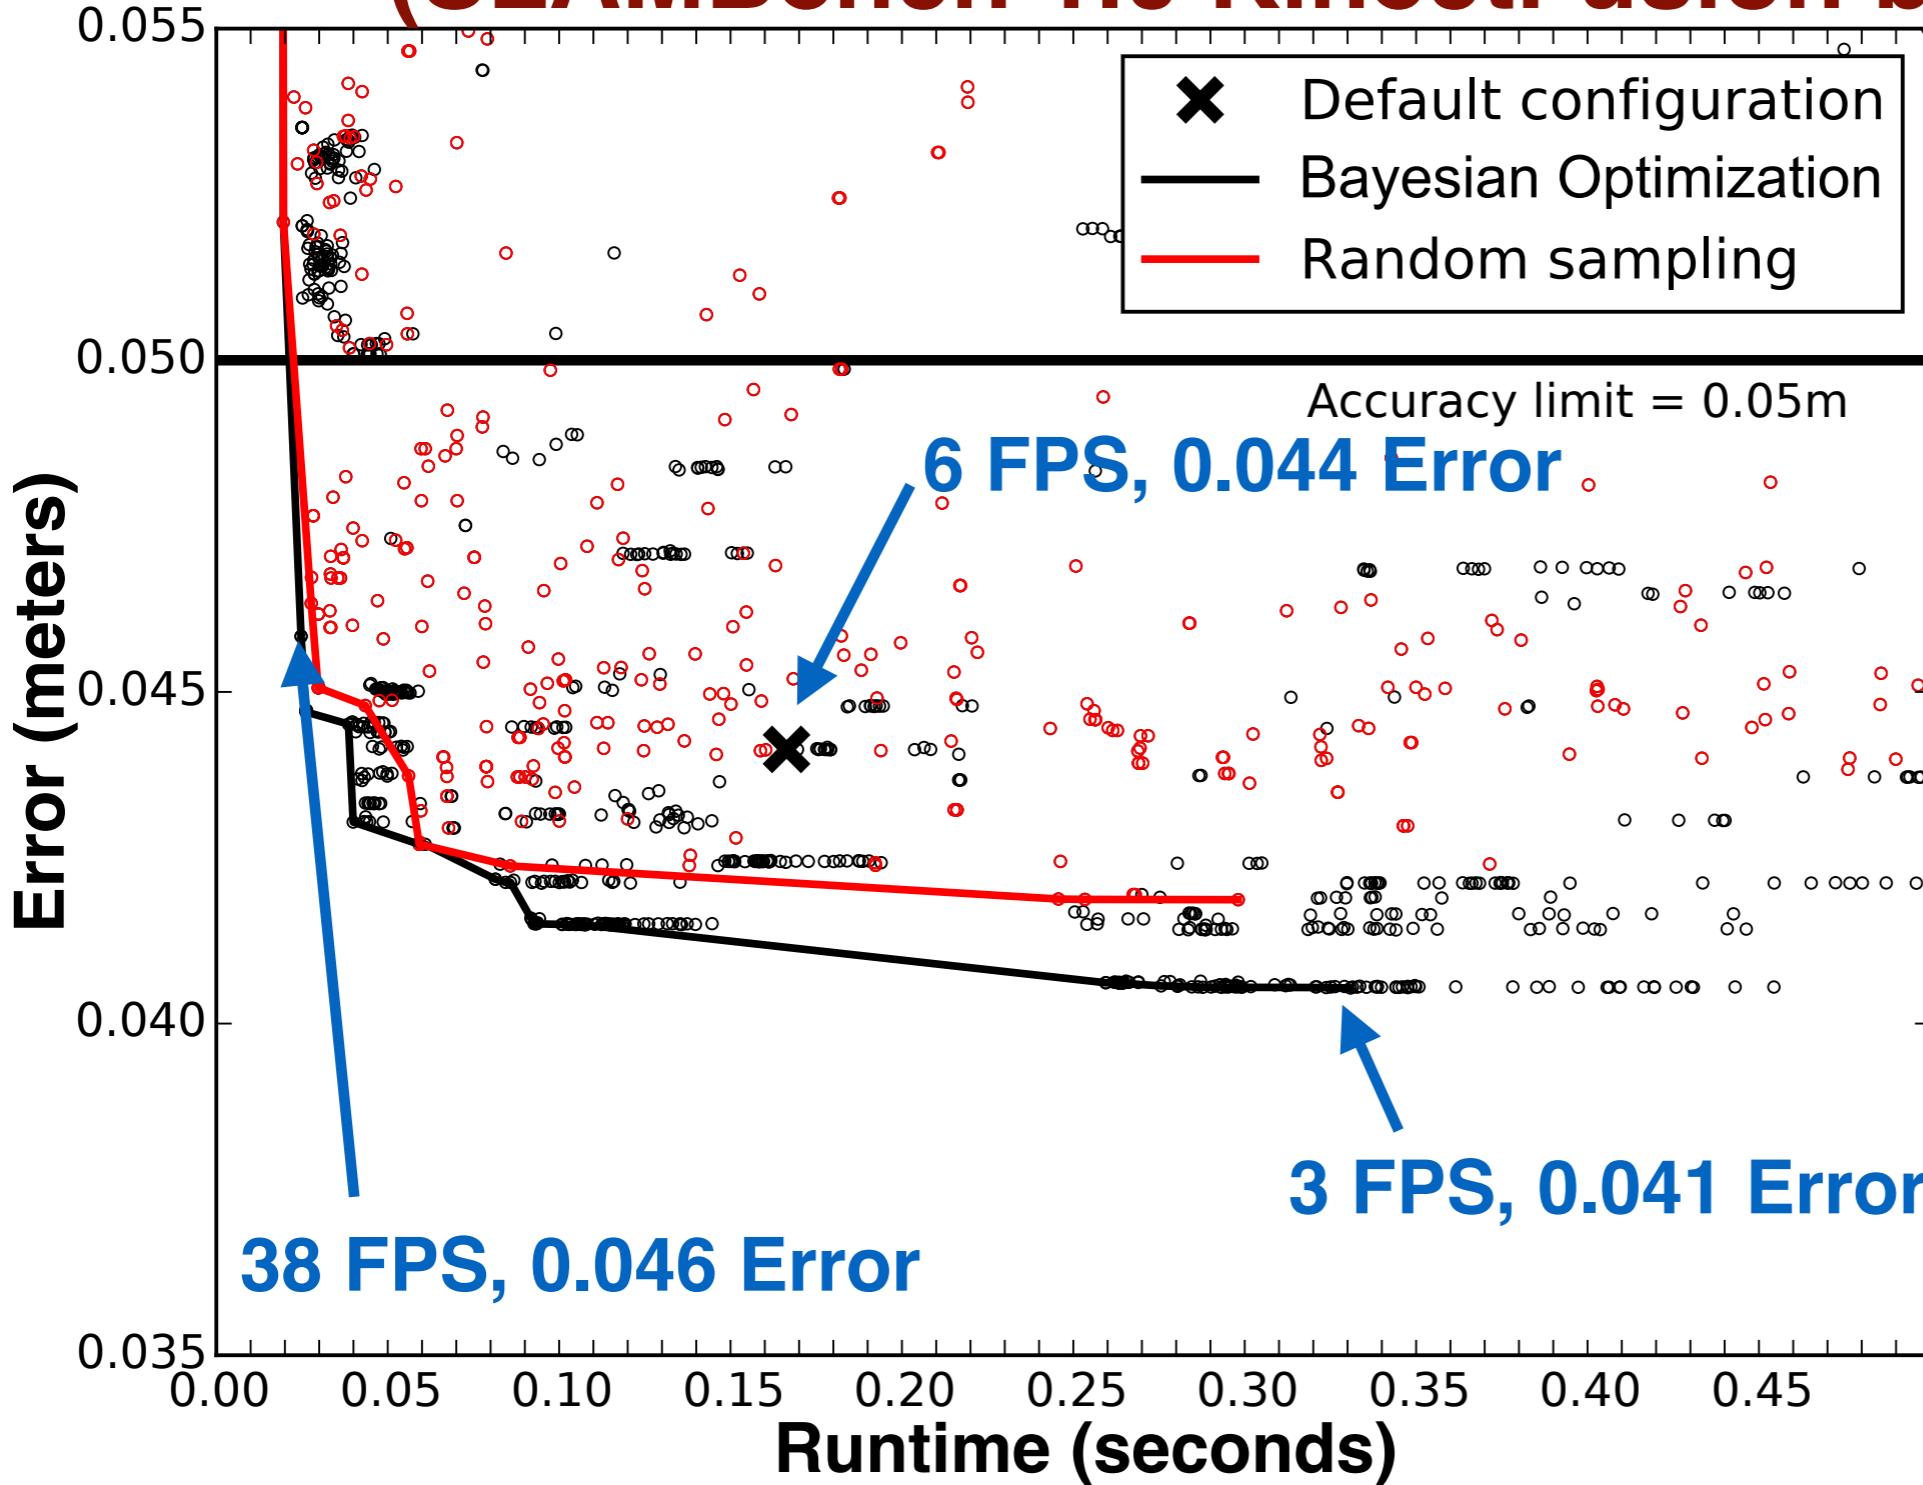


80x60

**Three objectives:**

1. Error
2. Runtime
3. Energy consumption

# Example 3: Computer Vision Use Case (SLAMBench 1.0 KinectFusion benchmark)

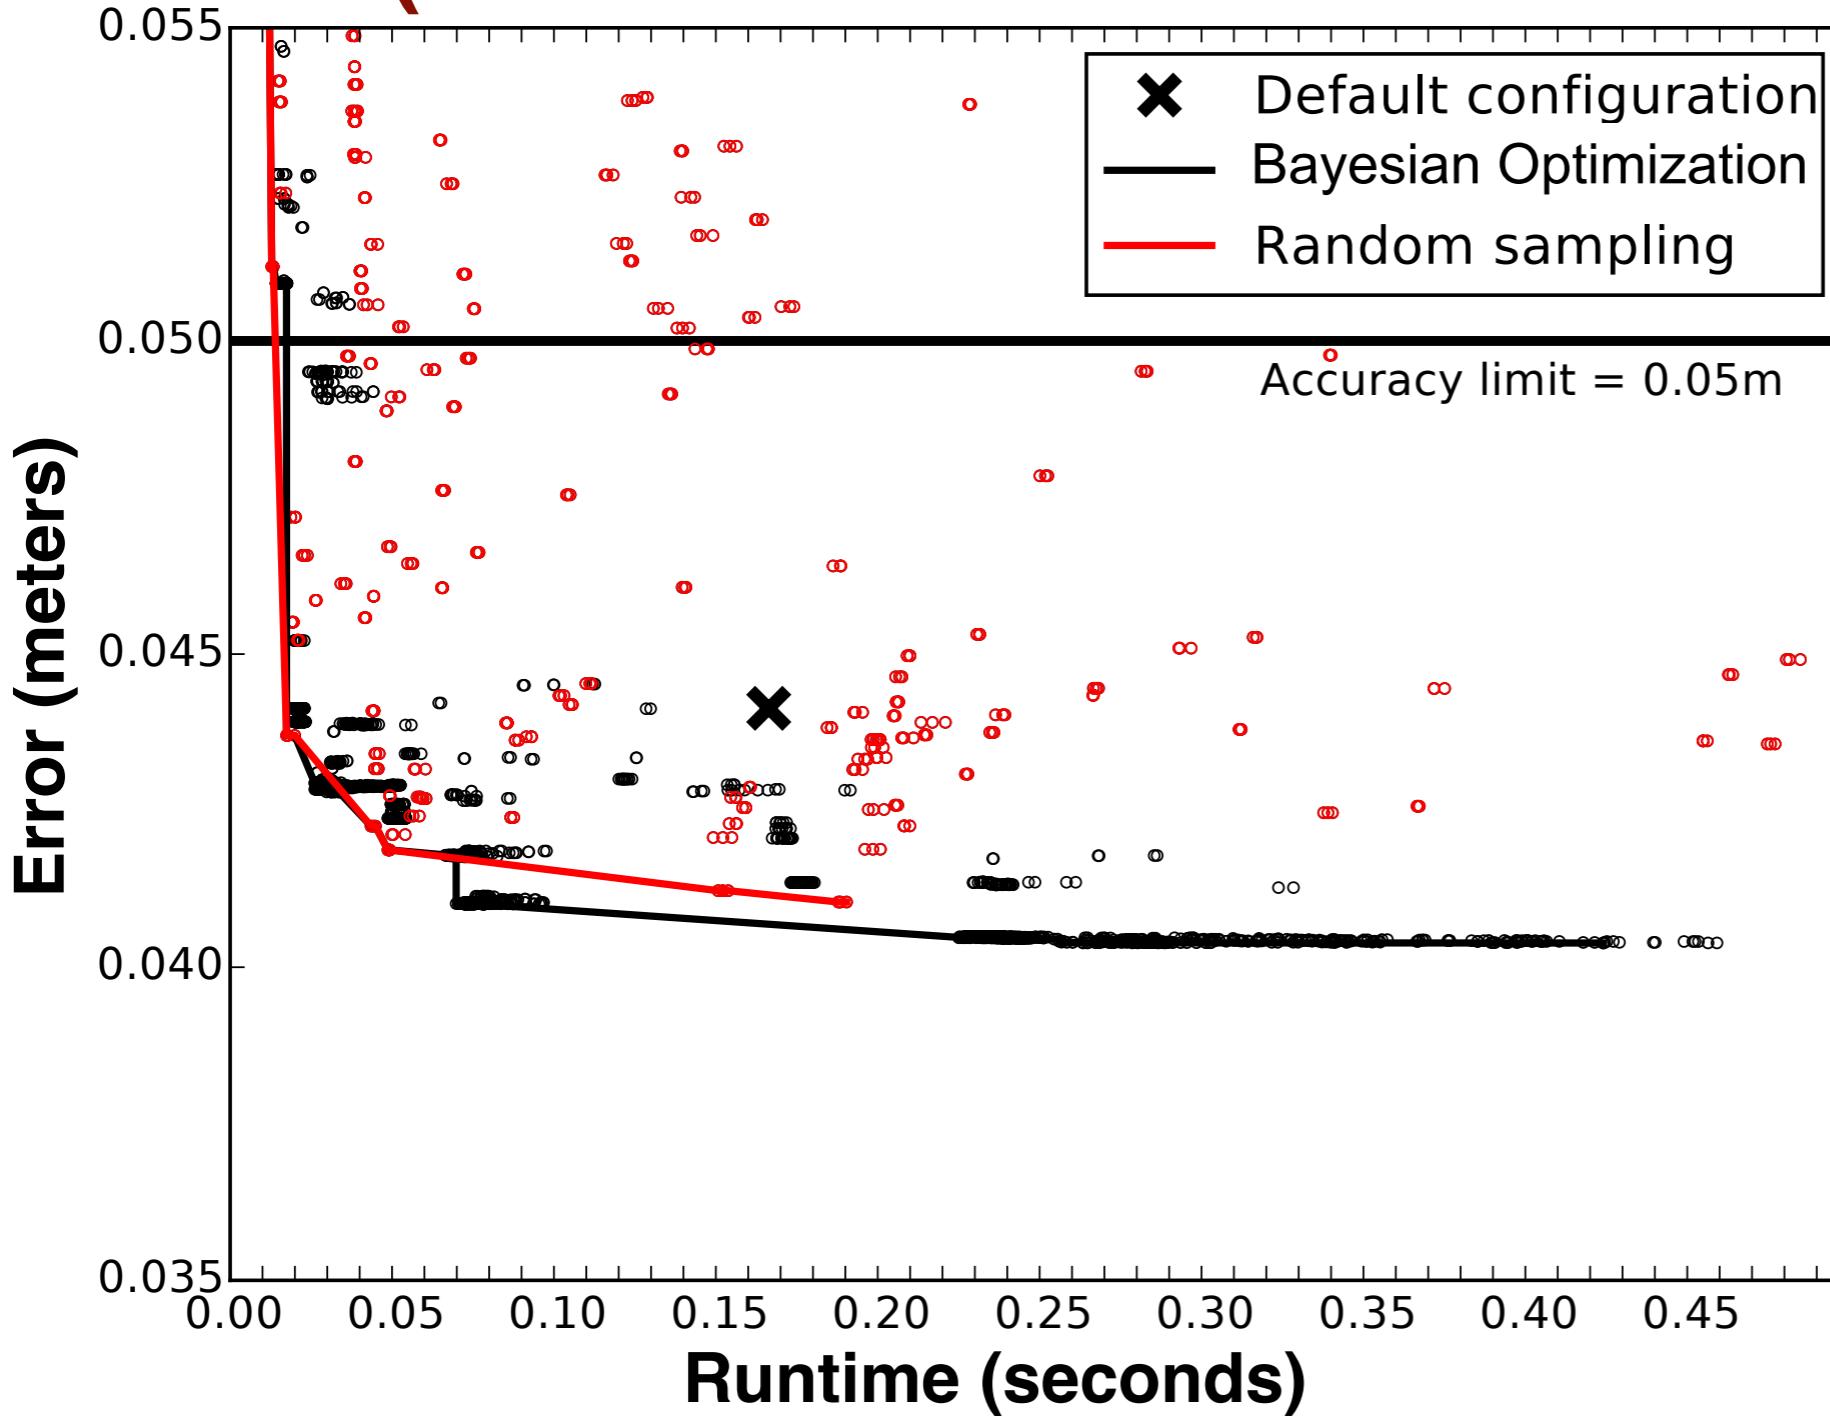


3D reconstruction



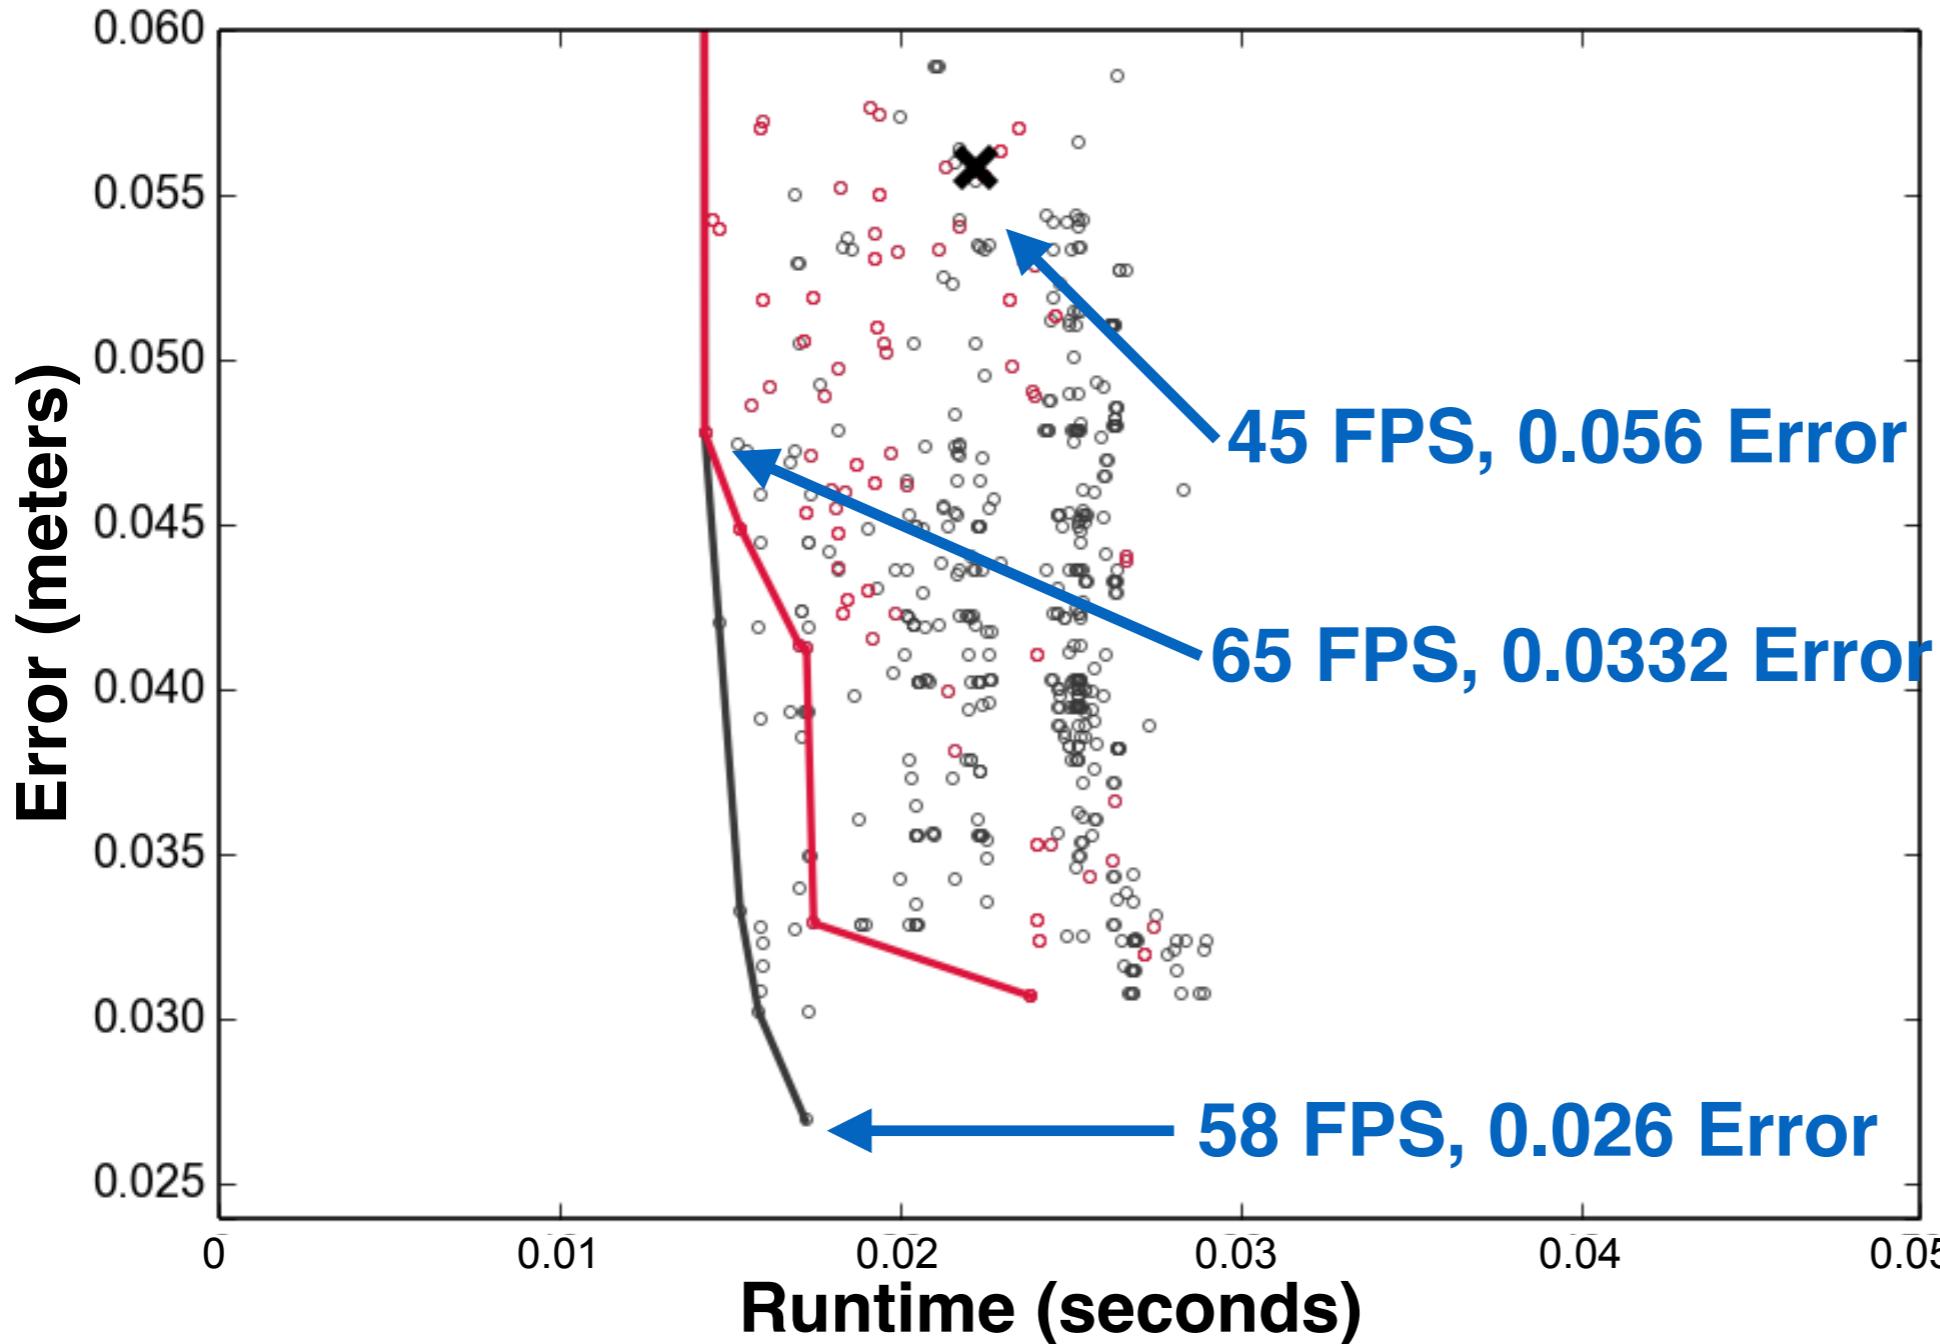
Machine	Type	CPU	CPU name	CPU cores	GPU	GPU name
Hardkernel ODROID-XU3	Embedded	ARM A15 + A7	Exynos 5422	4 + 4	ARM	Mali-T628

# Example 3: Computer Vision Use Case (SLAMBench 1.0 KinectFusion benchmark)



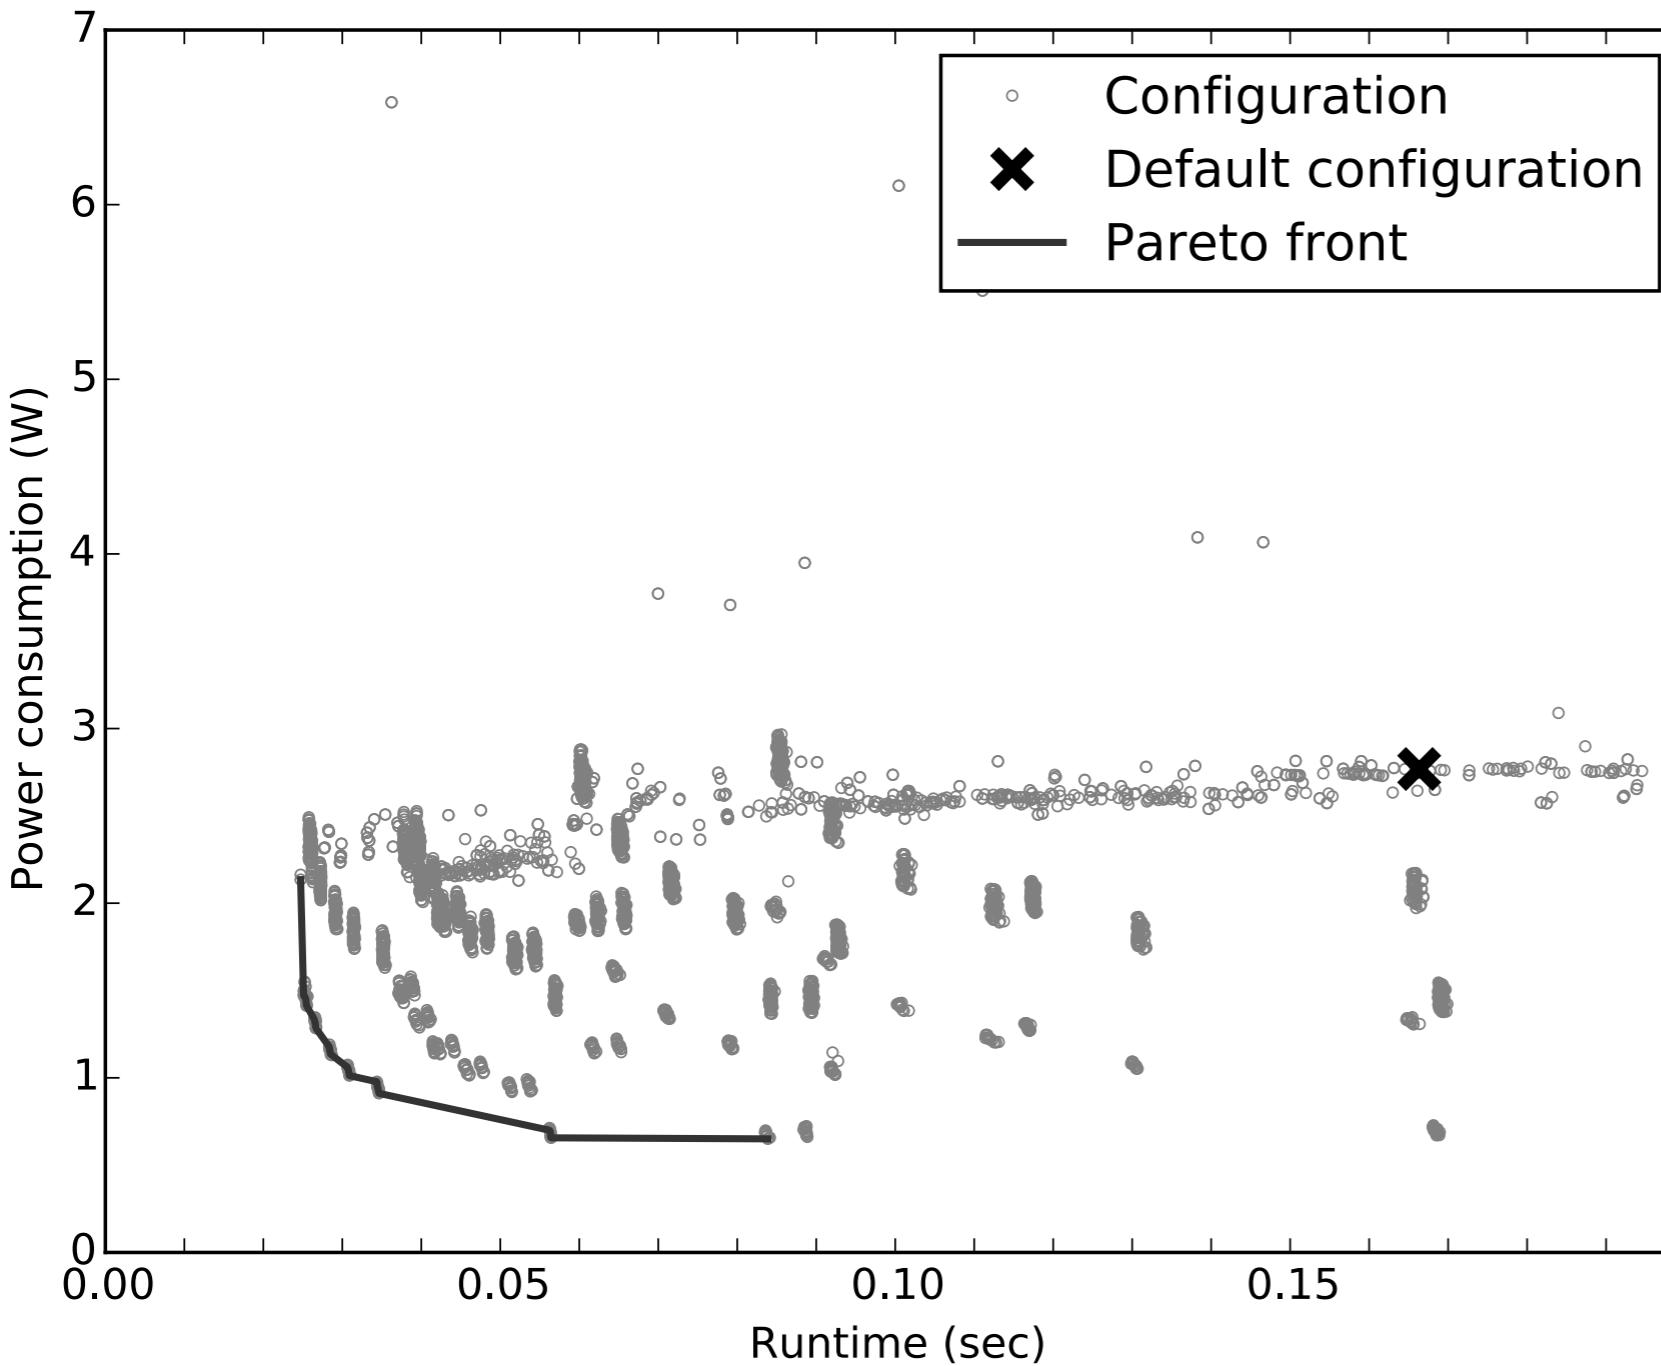
Machine	Type	CPU	CPU name	CPU cores	GPU	GPU name
ASUS T200TA	Detachable laptop	Intel Silvermont	Atom Z3795	4	Intel	HD Graphics

# Example 3: Computer Vision Use Case (SLAMBench 2.0 ElasticFusion benchmark)



Machine	Type	CPU	CPU name	CPU cores	GPU	GPU name
NVIDIA/Intel	Desktop	Intel Ivy Bridge	E5-1620	8	NVIDIA	GTX 780 Ti

# Example 3: Computer Vision Use Case (SLAMBench 1.0 KinectFusion benchmark)



Machine	Type	CPU	CPU name	CPU cores	GPU	GPU name
Hardkernel ODROID-XU3	Embedded	ARM A15 + A7	Exynos 5422	4 + 4	ARM	Mali-T628

# Conclusion and Future Work

- Goal: HyperMapper is a multi-objective BBO framework
- Demonstrated on:  
hardware design, DBMS and CV applications
- Future:
  - Dynamic Black-box optimization
  - Use of prior knowledge in the multi-objective posterior
  - Asynchronous batch optimization
  - Transfer learning
  - Applications:
    - Microsoft PostgreSQL/Azure full scenario
    - Multi-fidelity in SLAM
    - Halide computer vision language

# Research Team and Collaborators

- Work in collaboration with many people from academia and industry
- **We are hiring!** WASP-AI is funding positions for my research group:
  - 2 Ph.D. students (<https://lu.varbi.com/en/what:job/jobID:280143/>)
  - 2 Postdocs (<https://lu.varbi.com/en/what:job/jobID:280619/>)



Luigi Nardi, Ph.D.  
WASP-AI Assistant Professor  
in Machine Learning



Artur Luis  
Ph.D. Student  
UFMG

# Info on HyperMapper

- Join HyperMapper on **Slack**: [hypermapper.slack.com](https://hypermapper.slack.com)
- **Repo**: <https://github.com/luinardi/hypermapper>
- **Wiki**: <https://github.com/luinardi/hypermapper/wiki>
- HyperMapper **SaaS**: [www.hypermapper.com](http://www.hypermapper.com)

## Adopters



**teradata.**



**UC San Diego**



Microsoft - Database Management Systems

Teradata - Database Management Systems

Stanford University - Hardware Design

UCSD - Spector Benchmarks (FPGAs)

UT Austin - Capri (Approximate Computing)

ICL - Computer Vision and Robotics

Etc.

# Backup slides

# HyperMapper SaaS

The screenshot shows the 'New experiment' creation interface for HyperMapper SaaS. The interface is divided into several sections:

- Experiment name:** A text input field where the user has typed "wei".
- Optimization iterations:** A text input field.
- Optimization objectives:** A large text input field with placeholder text "Enter values separated by commas".
- Input parameters:** A section for defining parameters. It includes fields for Name (text input), Type (dropdown menu set to "real"), Min (text input), Max (text input), Default (text input), and two buttons (- and +). There is also a "Create" button in a blue box and a "JSON" button.

# Demo HyperMapper

<https://github.com/luinardi/hypermapper/wiki/Quick-Start-Guide>